Cortex AgentiX
Documentation -

Detailed

1. Learn about Cortex AgentiX 

1.1. Get Started with Cortex AgentiX 

1.2. Understand Cortex AgentiX licenses  📝🟢

note

**You can purchase additional compute units.
*You can extend incident retention by
purchasing an add-on. For more information,
see Data retention policy.
Intel feed quotas are based on the selected
Fetches Indicators field in the integration
instance settings, not the enabled status.
Disabling an integration instance does not
affect the intel feed...

1.3. Agentic AI in Cortex AgentiX 

💡🔵

tip

The system Help Center agent delivers fast,
context-aware assistance to answer your
questions. You can ask natural language
questions, such as "How do I create a
dashboard?" or "Where can I review my data
retention policies?" and the agent retrieves
concise, relevant information from the
documentation. If a question remains
unresolved, the agent...

1.3.1. Manage Cortex AgentiX agents  📝🟢

note

When users create custom agents, the agents
have the same or fewer permissions as the user.
For example, if a user has sufficient permissions
to isolate an endpoint, the user's agent can also
isolate an endpoint.
All users have access to all system agents and
all public agents, but plan execution is limited
by the permissions of the individual user. A...

1.3.2. Cortex AgentiX use cases  📝🟢

note

If a team member in the channel sees the
summary and attempts to ask the agent, "Give
me more details on the first SLA issue," the
team member receives an access denied
message because the active session is only
available to you, the initiator.

1.3.3. Agentic Assistant security 

1.4. Key components and concepts  📝🟢

note

Agents are bound by the same rules and robust
permissions as a human user. In addition, you
can mark actions that make real-world changes
in production systems as sensitive, requiring a
quick manual review and confirmation, ensuring
peace of mind before critical system changes
are made.

1.5. Data retention policy 

1.6. Supported ciphers 

1.7. Supported web browsers  📝🟢

note

Cortex AgentiX is optimized for best
performance using the HTTPv2 protocol. If you
are using a non-supported browser version or a
network middleware that uses a protocol other
than HTTPv2, you may experience slowness in
loading and response times between the
browser and the Cortex AgentiX service.
Consult with your browser or middleware
vendor to see...

2. Onboard Cortex AgentiX 

2.1. Plan and prepare your deployment 

2.2. Onboarding checklist 

2.3. Activate Cortex AgentiX 

☑️🟣

prer…

The Cortex AgentiX activation email.
A Customer Support Portal (CSP) account.
You need to set up your CSP account. For more
information, see How to Create Your CSP User
Account.
When you create a CSP account, you can set
up two-factor authentication (2FA) to log into
the CSP by using an Email, Okta Verify, or
Google Authenticator (non-FedRAMP...

📝🟢

note

Activation can take about an hour and does not
require you to remain on the activation page.
Cortex AgentiX sends a notification to your
email when the process is complete.

📝🟢

note

If you want to change your tenant's name, the
subdomain, or activate a development tenant
(subject to license), on the right-hand side, click
the ellipsis.
You can only change the subdomain once, and
it cannot be undone.
After deleting the subdomain, you can reuse it
after 7 days.

2.3.1. Cortex AgentiX supported regions 

2.3.2. Enable access to Palo Alto Networks

resources



📝🟢

note

Some of the IP addresses required for access
are registered in the United States. As a result,
some GeoIP databases do not correctly
pinpoint the location from which IP addresses
are used. All customer data is stored in your
deployment region, regardless of the IP address
registration, and data transmission is restricted
through any infrastructure to...

📝🟢

note

<agentix-tenant> refers to the chosen
subdomain of your Cortex AgentiX tenant, and
<region> is the region in which your tenant is
deployed.
The port is 443.

2.4. Install and configure content 

📝🟢

note

Cortex AgentiX includes AI system actions,
used by system agents. Some system actions
are available out-of-the-box with pre-installed
content packs, such as Common Scripts and
other system AI actions require you to
download and install the relevant content pack.
To find and install additional content packs that
include actions, go to Marketplace and...

📝🟢

note

Not all content packs are available to
download on the Data Sources page. If you
require content packs such as Phishing and
Malware, you need to download the pack from
Marketplace and configure the integration.

2.5. Post-deployment steps 

2.5.1. Post-deployment checklist 

2.5.2. Set up users and roles 

2.5.3. Configure server settings 

📝🟢

note

Keyboard shortcuts, timezone, and timestamp
format are not set universally and only apply to
the user who sets them.

📝🟢

note

If the Password Protection (for downloaded
files) setting under Settings → Configuration →
General → Server Settings is enabled, enter the
password 'suspicious' to download the file.

📝🟢

note

In MSSP environments, this option is not
available on the parent tenant.

📝🟢

note

Requires the TIM add-on.

🧭➡️

navi…

Settings → Configurations → General → Server
Settings

2.5.4. Configure security settings 

2.5.5. Set up an engine  🧭➡️

navi…

Settings → Configurations → Data Broker →
Engines → Create New Engine

2.5.6. Log forwarding 

2.5.6.1. Forward logs and data from Cortex
AgentiX to external services 

2.5.6.1.1. Configure external applications for
forwarding 

📝🟢

note

There are two options for configuring external
applications. To configure relevant external
applications before you begin creating
forwarding notifications, follow the steps in the
following topics using the menu path Settings
→ Configurations → Integrations → External
Applications. You can also configure an
external application as part of the...

2.5.6.1.1.1. Integrate a syslog receiver  📝🟢

note

Up to TLS 1.3 is supported.
Verify the self-signed CA includes your public
key.

2.5.6.1.1.2. Integrate Slack for outbound
notifications  📝🟢
note

You are directed to the Slack browser to install
Cortex AgentiX. You can only use this link to
install Cortex AgentiX on Slack. Attempting to
install from Slack Marketplace will redirect you
to Cortex AgentiX documentation.

2.5.6.1.2. Configure notification forwarding 

☑️🟣

prer…

Before you can select an external service for
notification forwarding, you must integrate the
external service with Cortex AgentiX. For more
information, see Configure external
applications for forwarding. No prior
configuration is required to send data to an
email distribution list.

📝🟢

note

Forwarding destinations: Only issues and cases
can be forwarded to Slack, Splunk, Amazon
SQS, Amazon S3, or Webhook.
Notification forwarding by domain: To
configure notification forwarding for issues by
domain, select Issues and filter the Issues table
by Issue Domain.
Alert vs. issue format:By default, new
configurations use the issue format, but you...

📝🟢

note

Not all data and log types can be sent to all
external services. For more information, see
Forward logs and data from Cortex AgentiX to
external services.

📝🟢

note

The Grouping Timeframe defines the time
frame, in minutes, of how often Cortex AgentiX
sends notifications. Every 20 issues or 20 events
aggregated within this time frame are sent
together in one notification, sorted according
to severity. To send a notification when one
issue or event is generated, set the time frame
to 0.

2.5.6.1.3. Monitor administrative activity 

2.5.6.2. Data and log notification formats 

📝🟢

note

Issues can be forwarded to email, syslog
servers, and Slack in the alert format, if you
prefer. The alert format can be selected when
you configure your forwarding notification.

2.5.6.2.1. Issue notification format 

2.5.6.2.2. Management Audit log notification
format 

3. Cortex AgentiX configuration 

3.1. Configure Cortex AgentiX 

3.2. Content configuration 

3.2.1. Configure content using Data Sources 

3.2.1.1. Manage instances 

3.2.1.2. Add a new data source or instance 

📝🟢

note

If a data source contains multiple integrations,
the integration configured as the default
integration will used by the Data Onboarder.
The default integration of the content pack is
indicated in each content pack's
documentation. The other integrations are
available for configuration in the Data Sources
& Integrations page after installing the content...

📝🟢

note

If you are adding a new instance to an existing
data source, these options are View only.
You can adjust the view-only options on the
relevant page in the system, for example
Correlations, Playbooks, or Scripts.
If you are adding a new instance to an existing
data source, these options are View only.
You can adjust the view-only options on the
relevant...

3.2.1.3. Configure integration permissions  🧭➡️

navi…

Settings → Configurations → Data Collection
→ Integration Permissions

3.2.2. Marketplace 

3.2.2.1. Cortex Marketplace  🧭➡️

navi…

Settings → Configurations → Marketplace) or
at Cortex Developer Docs Marketplace

3.2.2.2. Content packs  📝🟢

note

Cortex AgentiX includes a built-in default mail
sender. You also have the option of installing a
different mail sender content pack, such as
Microsoft Exchange Online.

3.2.2.3. Content Pack Support Types  📝🟢

note

Palo Alto Networks is not liable for and does
not warrant or support any content pack
produced by a third-party publisher.
Palo Alto Networks does not support content
packs that do not have official available
documentation.

3.2.2.4. Manage content packs 

📝🟢

note

Optional content packs that are already
installed are treated like they are required
content packs to preserve content integrity.

📝🟢

note

Cortex AgentiX includes a built-in default mail
sender. You also have the option of installing a
different mail sender content pack, such as
Microsoft Exchange Online.

📝🟢

note

In addition to content packs that you install
from Marketplace, related content packs are
automatically downloaded when you adopt
playbooks or edit tasks that require content
items such as scripts or integrations.

📝🟢

note

Third-party product integrations are developed
and tested against a specific product version.
For products that are on-prem or cloud-based
with specific API versions, the version
developed and tested against will be included
in the integration's documentation. Newer
versions of the product are not always
immediately tested, and it is expected that...

⚠️🟡

caut…

If you want to downgrade, any content that
depends on the content pack including any
customizations may be deleted if it does not
exist in the target content pack version.

⚠️🟡

caut…

If another content pack is dependent on the
content pack you want to delete, it may break
the other content pack. You can reinstall the
content pack, but you cannot restore detached
and customized content.

💡🔵

tip

You can also find content packs that require
updates by going to Settings → Data Sources &
Integrations and filtering by Pack Version =
Update Available. If you click on an integration
in the filtered list, there is a link to the content
pack in Marketplace for updates.

3.2.2.5. Marketplace FAQs 

3.2.2.6. Content changes when upgrading
Cortex AgentiX versions 

3.2.2.7. Content pack contributions 

3.2.3. Integrations 

📝🟢

note

Some integrations can be downloaded directly
without having to initially download a content
pack from Marketplace. For more information,
see Configure content using Data Sources.
In addition to content packs that you install
from Marketplace, related content packs are
automatically downloaded when you adopt
playbooks or edit tasks that require content...

📝🟢

note

If the integration was installed from a content
pack, you need to duplicate the integration
before editing.

📝🟢

note

Integration commands are only available when
the integration instance is enabled. Some
commands depend on a successful connection
between Cortex AgentiX and third-party
integrations.

💡🔵

tip

In the Playground, you can clear the context
data, if needed, which deletes everything in the
Playground context data, but does not affect
the actual issue or case. To clear the context,
run !DeleteContext all=yes' from the CLI or
click Clear Context Data while viewing the
context data.

🧭➡️

navi…

Settings → Configurations → Marketplace

3.2.3.1. Manage API keys 

👀🟠

imp…

Before configuring, ensure that you review the
Manage user scope topic.

👀🟠

imp…

By default, Enable Scope Based Access
Control is disabled in Settings →
Configurations → General → Server Settings,
and granular scoping is not enforced. Before
enabling SBAC, we recommend that an
administrator or a user with Access
Management permissions first ensures that the
users, user groups, and API Keys defined in
Cortex AgentiX are...

👀🟠

imp…

You will not be able to view the API key again
after you complete this step. Ensure that you
copy the API key before closing the notification.

3.2.3.2. Integration use cases 

3.2.3.3. Add an integration instance 

📝🟢

note

This procedure describes how to add an
integration instance from the Data Sources &
Integration page. Some integration instances
can also be configured on the Data Sources
page. For more information, see Add a new data
source or instance.

📝🟢

note

In some integrations, a classifier, an issue type,
and mapper fields are included.

3.2.3.4. Fetch issues from an integration instance 

📝🟢

note

In some integrations, the Issue Fetch interval is
called Feed Fetch Interval.
If the integration instance does not have the
Issue Fetch Interval field, you need to add this
field by editing the integration settings. If the
integration is from a content pack, you need to
create a copy of the integration. Any future
updates to this integration will not...

📝🟢

note

Some content packs include preconfigured
correlation rules, but you should review them to
see if they suit your use case and duplicate
them if required. Go to Threat Management →
Detection Rules → Correlations, search for the
relevant rule, right-click, and select Preview
Rule. For example, the ServiceNow v2 Alerts
(automatically generated)...

3.2.3.4.1. Map fields to issue types 

📝🟢

note

In the Classification & Mapping page, the
mapping does not indicate for which issue
types they are configured. Therefore, when
creating a mapper, it is best practice to add to
the mapper name, the issue types the mapper is
for. For example, Mail Listener - Phishing.

📝🟢

note

When mapping a list, we recommend you map
to a multi select field. Short text fields do not
support lists. If you do need to map a list to a
short text field, add a transformer in the
relevant playbook task, to split the data back
into a list.

🧭➡️

navi…

Settings → Configurations → Object Setup →
Issues → Classification & Mapping

3.2.3.4.2. Classify events using a classifier for
issue types 

3.2.3.5. Forward Requests to Long-Running
Integrations 

📝🟢

note

For more information about how to set up the
integration, see Manage external dynamic lists.

📝🟢

note

When running on the tenant, you can only use
long-running integrations provided by Cortex
AgentiX, you cannot create custom ones.
Custom long-running integrations are
supported only on engines at this time.
Configuring custom certificates or private API
Keys in the long-running integration instance is
supported only on engines, not on the Cortex...

📝🟢

note

For the TAXII Server and TAXII2 Server
integrations, the xdr string is automatically
replaced by crtx. For the Microsoft Teams
integration, you can use the microsoft-teams-
create-messaging-endpoint command to get
the correct messaging endpoint based on the
server URL, the server version, and the instance
configurations. For more information, see...

3.2.3.6. Manage credentials  🧭➡️

navi…

Settings → Configurations → Integrations →
Credentials → New Credential

3.2.4. Troubleshoot Integrations  📝🟢

note

If an integration instance consistently
encounters API rate limits when running on the
tenant, consider configuring the instance to run
on an engine to change the source of the
outbound traffic.

3.2.5. Measuring data freshness  📝🟢

note

The max_delay metric is taken from the
maximum bucket value with a restricted limit;
therefore, metrics show whole numbers.
The median and ninetieth_percentile metrics are
statistical calculations that give an
approximation of the real value; therefore,
metrics show decimal numbers.
Time slots with a zero log count or zero byte
count display records...

3.2.6. Overview of data ingestion metrics  ☑️🟣

prer…

For Cortex AgentiX to monitor data ingestion
health and create health issues, you must
enable Cortex - Analytics. Go to Configurations
→ Cortex - Analytics. For more information, see
Enable the Analytics Engine and Identity
Analytics.Enable the Analytics Engine and
Identity Analytics

3.2.7. Verify collector connectivity  📝🟢

note

If an issue is triggered, the investigation options
in the right-click menu of the Health Issues
pages are context-specific. Make sure that you
specify the relevant issue category.

3.2.8. Creating correlation rules to monitor data
ingestion health 

3.3. Configure the Cortex Agentic Assistant 

3.3.1. Agentic Assistant components and
concepts  📝🟢
note

Agents are bound by the same rules and robust
permissions as a human user. In addition, you
can mark actions that make real-world changes
in production systems as sensitive, requiring a
quick manual review and confirmation, ensuring
peace of mind before critical system changes
are made.

3.3.2. Agents Hub 

📝🟢

note

To manage agents in the Agents Hub, you must
have the proper permissions. For more
information, see Agentic Assistant role-based
access control.

3.3.2.1. Manage actions 

📝🟢

note

To manage actions in the Agents Hub, you must
have the correct permissions. For more
information, see Agentic Assistant role-based
access control.

💡🔵

tip

System actions may rely on content packs that
need to be installed and configured.

3.3.2.2. Register actions 

📝🟢

note

To register actions, ensure you have the correct
permissions. For more information, see Agentic
Assistant role-based access control.

📝🟢

note

Mandatory arguments cannot be deselected.

📝🟢

note

The execution of system or custom actions that
are based on integration commands can be
restricted using integration permissions.

👀🟠

imp…

A single content item can be registered as
different actions, with each action using
different inputs and outputs from the same
script. Only register the same script, command,
or AI prompt as a new action if it is required for
your use case, as providing an agent with many
actions with overlapping abilities can reduce
the ability of the agent to choose...

3.3.2.3. Manage agents 

📝🟢

note

In some cases, an agent may include actions
with content items that are not relevant for all
licenses. If that occurs, the grey circle appears,
but you are not able to install the related
content.

💡🔵

tip

System agents can include actions that require
additional content packs to be installed and
configured. To view all actions assigned to a
system agent, including actions not available
due to missing content, click on the system
agent in the Agents Hub. There may be actions
assigned to a system agent that are not
relevant to your organization. For...

3.3.2.4. Build agents 

📝🟢

note

To build agents in the Agents Hub, you must
have view/edit permissions. For more
information, see Agentic Assistant role-based
access control.

📝🟢

note

If you clear the checkbox for a role, all actions
associated with that role are also cleared. The
exception is if another role is also selected,
which is associated with the same actions.
If you clear the checkbox for an action, all roles
associated with that action are cleared. For
example, if you select the Investigator role, and
Send Mail and Tavily...

3.3.2.5. Expand agent capabilities with MCP
integrations 

📝🟢

note

For agentic actions to be created from tools on
an MCP server, the tools must have input and
output descriptions on the MCP server. These
descriptions are required for Cortex AgentiX to
understand the tools' capabilities and the
correct use cases.

📝🟢

note

If you are using OAuth-based authentication,
the Test button returns an error containing the
command to run in the playground in order to
test the connection.

📝🟢

note

If an MCP tool is removed from the MCP server,
the action will be unavailable due to missing
content. If the tool is restored on the MCP
server, the action is automatically reenabled.
If you have a development and a production
tenant, you can push actions created from MCP
tools from development to production.

3.3.3. Agentic Assistant role-based access
control  📝🟢
note

Agents are limited by the individual permissions
of the user. For example, if users do not have
sufficient permissions to isolate an endpoint,
they cannot use an agent to isolate an endpoint.
The execution of system or custom actions that
are based on integration commands can be
restricted to specific roles using integration
permissions.

3.4. Cortex MCP server  3.4.1. Cortex MCP server overview 

📝🟢

note

The Cortex MCP Server empowers you to
integrate AI into your security workflows using
natural language. When using LLM-based
suggestions, always review and approve actions
suggested by the AI before they're executed.
We recommend deploying the Cortex MCP
server in a secure environment where access is
limited to authorized users.

3.4.1.1. Install the Cortex MCP server 

☑️🟣

prer…

If you are running the Cortex MCP server in a
Poetry virtual environment, you must have
Python 3.13 or higher.
If you plan to run the Cortex MCP server in a
Docker container, you must have Docker
installed.

📝🟢

note

The MCP Server uses public APIs to
communicate and is limited by the license
quotas available in your tenant. This is
particularly relevant when running XQL queries.
For more information on running XQL query
APIs, see Run XQL query APIs.

📝🟢

note

It is critical to avoid assigning excessive
permissions when creating an API key for the
Cortex MCP Server. Since the key has both read
and write capabilities, overly broad permissions
can lead to unintended actions and potentially
compromise your environment. Ensure the key
follows the principle of least privilege and is
granted only the minimum...

📝🟢

note

When using Docker, we recommend using an
.env file to set the Cortex API credentials as
environment variables. While the credentials
can be provided in the MCP client configuration
settings, the .env file provides safer handling of
API credentials and makes your configuration
easily reproducible.

📝🟢

note

By default, stdio (standard input/output) is
used. You can also configure Streamable HTTP,
to send requests directly to the tenant instead
of through the MCP client. Streamable HTTP
can be useful for testing in the browser without
a MCP client and to bypass limits that may be in
place for your MCP client. For Docker, you can
include the Streamable HTTP...

👀🟠

imp…

To configure the Cortex MCP Server, you need
the Cortex API URL, Cortex API key, and Cortex
API key ID. You will not be able to view the API
key again after you complete this step. Ensure
that you copy the API key before closing the
notification.

3.4.1.2. Configure the MCP client  🧭➡️

navi…

Settings → Developer → Edit Config

3.4.1.3. Use the Cortex MCP server  📝🟢

note

The built-in tools retrieve information, but do
not write to the tenant. You can create your
own tools that include write actions. The
examples below include both.

3.4.1.4. Create custom Cortex MCP server tools  📝🟢

note

Any new or updated components provided by
Cortex are automatically downloaded into the
remote_components folder. During each
update, the folder is fully replaced and all
existing contents are recreated. Do not add
custom tools to this directory, as it is managed
entirely by Cortex and is overwritten at every
update.

3.5. Remote repository management 

3.5.1. Cortex AgentiX development tenant  👀🟠

imp…

Development tenants are not intended for
performance checks; they cannot access
production data, and they are connected to a
limited number of endpoints. As a result, all
development tenants have fewer resources than
the production tenant, including data ingestion
capacity and performance, and compute
capabilities. In a development tenant, extreme...

3.5.2. Set up a remote repository 

🧭➡️

navi…

If the content repository option is disabled for
the production or development tenant (under
Settings → Configurations → General →
Remote Repository Settings, toggle the
Content repository slider to off), the tenant
becomes standalone and does not push or pull
content

If you are changing your remote repository
settings, back up existing content to your local
computer by navigating to Settings →
Configurations → General → Server Settings →
Custom Content and click Export all custom
content

3.5.2.1. Set up a built-in remote repository 

📝🟢

note

Once enabled, development tenants have a red
banner on the top left showing DEV.

📝🟢

note

The first tenant that is enabled pushes its
content to the remote repository first. For
example, these instructions describe enabling
the production tenant first, so the remote
repository will initially contain production
tenant content. You can enable a development
tenant first if you want the remote repository to
initially contain the content from...

🧭➡️

navi…

Settings → Configurations → General →
Remote Repository Settings and toggle the
Content repository slider to enable the remote
repository

3.5.2.2. Set up a Private Remote Repository 

📝🟢

note

Once enabled, the development push tenant
has a red banner on the top left showing DEV.

📝🟢

note

For repository vendors that use tokens, enter
the token type in the username field and the
token in the password field. Verify details with
your vendor.
If your private Git remote repository uses
personal access tokens instead of usernames
and passwords, enter the token type in the
username field and the access token in the
password field. For example,...

📝🟢

note

You can't add an engine that has been added to
a Load-Balancing Group.

📝🟢

note

If your private Git remote repository uses
personal access tokens instead of usernames
and passwords, enter the access token in the
password field and leave the username field
blank.
For repository vendors that use tokens, the
token type is entered in the username field and
the token is entered in the password field.
Verify details with your vendor.
If...

📝🟢

note

The first tenant that is enabled pushes its
content to the remote repository first. For
example, these instructions describe enabling
the production tenant first, so the remote
repository will initially contain production
tenant content. You can enable a development
tenant first if you want the remote repository to
initially contain the content from...

📝🟢

note

For repository vendors that use tokens, enter
the token type in the username field and the
token in the password field. Verify details with
your vendor.
If your private Git remote repository uses
personal access tokens instead of usernames
and passwords, enter the token type in the
username field and the access token in the
password field. For example,...

📝🟢

note

You can't add an engine that has been added to
a Load-Balancing Group.

📝🟢

note

If your private Git remote repository uses
personal access tokens instead of usernames
and passwords, enter the access token in the
password field and leave the username field
blank.
For repository vendors that use tokens, the
token type is entered in the username field and
the token is entered in the password field.
Verify details with your vendor.
If...

💡🔵

tip

Due to security concerns, there is a closed
allow list of approved URLs for private
repositories. If you want to use a URL that is
excluded from the allow list, use an engine
(engine groups are not supported).

3.5.3. Push and pull content 

☑️🟣

prer…

When pulling playbooks or scripts from a
remote repository into a production tenant, you
must define how ownership and access are
assigned for the incoming object's new content.
Select Settings → Configurations → General →
Remote Repository Settings.
Under Access to new content, choose how to
assign ownership for pulled content:
Keep the...

📝🟢

note

For more information, see Manage access to
playbooks and scripts, including the the
Remote Repositories (Push/Pull) section.

📝🟢

note

In a production (pull) tenant, manual editing of
synchronized content is blocked:
Users cannot be granted Editor permissions.
Any existing Editor permissions (set before the
tenant was configured for remote repository
synchronization) are effectively treated as
Viewer access. These capabilities are removed
from Cortex AgentiX and blocked by the
system...

3.5.4. Remote repository troubleshooting 

3.6. Users and roles management 

3.6.1. Users and roles in Cortex AgentiX  📝🟢

note

You can only set dataset access permission for
a user role in the tenant. These settings are
disabled in Cortex Gateway

3.6.2. Roles management 

📝🟢

note

The user who activated the Cortex product is
assigned the Account Admin role.
Only users with the Account Admin role can
add or remove another Account Admin user
role.

📝🟢

note

By default, users do not have roles assigned. If
no direct or user group role has been assigned,
users don't have permission to view or edit data
in the Cortex tenant.

3.6.2.1. Manage roles in the Cortex AgentiX
tenant 

💡🔵

tip

For analysts, we recommend limiting specific
permissions, such as:
Removing the ability to install, delete,
Marketplace/Data Sources, which should be
reserved for engineers and administrators. We
recommend setting these permissions for
analysts to None or View.
Removing access to API keys. Under
CONFIGURATIONS , set the Public API access
to None or...

💡🔵

tip

We recommend making a copy of out-of-the-
box roles and editing the copies, rather than
creating new roles, to avoid missing any
important permissions.

3.6.2.2. Role permissions by component 

💡🔵

tip

While Cortex AgentiX provides predefined, out-
of-the-box roles, it is highly recommended to
make copies of these roles and edit them rather
than creating new roles from scratch, ensuring
you do not miss any critical underlying
permission dependencies.

🧭➡️

navi…

You can manage role permissions in Cortex
AgentiX under Settings → Configurations →
Access Management → Roles

3.6.2.2.1. Dashboards and Reports permissions 

3.6.2.2.1.1. Dashboards permissions  📝🟢

note

By default, all new dashboards are restricted
(private) and visible only to the creator (the
owner). An administrator (or role with
equivalent access) can configure how
dashboards are shared across the tenant under
Settings → Configurations → Access
Management → Objects. These settings control
whether dashboard owners are allowed to
share their...

3.6.2.2.1.2. Reports permissions 

3.6.2.2.1.3. Command Center Dashboard

permissions 

3.6.2.2.2. Cases and Issues permissions 

📝🟢

note

Some roles require specific permissions. For
example, a Security Engineer may require
View/Edit for Playbooks, but a SOC Tier-1
Analyst does not.

⚠️🟡

caut…

To set Cases & Issues to View or View/Edit, you
must first set the Scripts and Playbooks
permissions to Enabled.
When SBAC is set to Restrictive mode, users
who don't have all the required tags shouldn't
be able to read or edit the parent case (fields
or context). For more information on setting
restrictive mode, see Configure server settings.
If users...

3.6.2.2.3. Investigation and Response permissions 

3.6.2.2.3.1. Query Library permissions 

📝🟢

note

Users must have at least View permission in
Query Center to access the Query Library from
Query Center.
If you want to execute and schedule queries
from the Query Library, you also need View/Edit
permission in the Query Center.

🧭➡️

navi…

Users can primarily access Query Library from
Investigation & Response → Search → XQL
Search, and select the Query Library tab

3.6.2.2.3.2. Query Center permissions  ⚠️🟡

caut…

Access to the Query Center is strictly governed
by Scope-Based Access Control (SBAC). Even if
users have View/Edit permissions, they will only
see data returned from the endpoint groups or
log sources defined in their specific role scope.

3.6.2.2.3.3. Playbook permissions 

📝🟢

note

Playbooks can only be Enabled when Scripts
are Enabled first.

📝🟢

note

Playbooks can only be Disabled after Cases
and Issues (under Cases & Issues) are set to
None first.

👀🟠

imp…

Playbooks are a prerequisite for the entire
Investigation & Response workspace. You
cannot set Cases and Issues to View/Edit unless
Playbooks is Enabled.

3.6.2.2.3.4. Script permissions 

📝🟢

note

Users can also view and edit Lists (under
Settings → Configurations → Object Setup →
Lists) provided they have Cases & Issues
permissions.

📝🟢

note

Scripts can only be Disabled after first setting
Playbooks to Disabled and Playground to None.

⚠️🟡

caut…

Scripts are a prerequisite for Playbooks. You
cannot set Playbooks to Enabled unless Scripts
is Enabled.
By enabling the Scripts component and
selecting Create scripts that will run with super
user, users gain unrestricted access to sensitive
system resources. Because these scripts bypass
standard security controls, this permission must
be strictly...

3.6.2.2.3.5. Playground permissions  🧭➡️

navi…

Controls Playground Investigation & Response
→ Automation → Playground, which is a testing
environment for commands and scripts that is
not connected to a live (active) investigation

3.6.2.2.3.6. Automation Exclusion Center

permissions  🧭➡️
navi…

Controls access to the Automation Exclusion
Center (Settings → Configurations →
Automation → Automation Exclusion Center),
which prevents a command or script from a
remediation action

3.6.2.2.4. Threat Management permissions 

3.6.2.2.4.1. Detection Rules permissions  ⚠️🟡

caut…

Users must have View/Edit access to the Query
Center to create or edit Correlation rules.

3.6.2.2.4.2. Threat Intelligence permissions 

🔒🟣

cont…

This feature requires the Threat Intelligence
Management (TIM) license.

🧭➡️

navi…

Located under Threat Management → Threat
Intelligence → Indicators, these permissions
govern how your organization interacts with
indicators (IPs, URLs, Domains, Hashes) and
intelligence feeds

3.6.2.2.5. Marketplace permissions  ⚠️🟡

caut…

Granting View/Edit access to the Marketplace
allows users to install new content packs. As
content packs often contain Python scripts and
automated playbooks, this permission
effectively allows users to introduce new
executable code into the tenant. This should be
restricted to Security Engineers and
Administrators.

3.6.2.2.6. Exceptions Configuration permissions 

3.6.2.2.7. Configuration permissions 

3.6.2.2.7.1. Auditing permissions  ⚠️🟡

caut…

The Auditing module is intentionally restricted
to View-only access. Audit logs are immutable
records designed to maintain compliance and
forensic integrity. They cannot be modified or
deleted by any user.

3.6.2.2.7.2. Alert Notifications permissions 

⚠️🟡

caut…

Configuring alert notifications requires
underlying infrastructure. For Slack or Email,
users need access to the Integrations
permissions.

🧭➡️

navi…

Configurations: Manage rules through Settings
→ Configuration → General → Notifications:
Includes main notification rules and user
notifications

External forwarding: Configure through
Settings → Configuration → Integrations →
External Applications

3.6.2.2.7.3. General Configuration permissions 

3.6.2.2.7.4. Access management permissions 

⚠️🟡

caut…

SSO Configuration Risk: Granting View/Edit
access allows users to modify the tenant's
Single Sign-On (SSO) and authentication
settings. Misconfigurations can cause tenant-
wide lockouts. Ensure only authorized identity
or infrastructure administrators hold this
permission.
Auditing is Mandatory: It is highly
recommended that any user managing access...

🧭➡️

navi…

Set permissions for Users, Roles, User Groups,
and Authentication Settings under Access
Management (Settings → Configurations →
Access Management

3.6.2.2.7.5. Data Sources permissions  🧭➡️

navi…

Data Sources is managed under Settings →
Configurations → Data Collection → Data
Sources & Integrations

3.6.2.2.7.6. Integrations - instance permissions 

📝🟢

note

There are two integration permissions:
Integrations (under Data Collection): Configure
data collection integrations
Integration Permissions (under Integrations):
Configure command permission for integrations.

📝🟢

note

Users can view the Credentials page, but
cannot edit it without the Credentials View/Edit
permission.

3.6.2.2.7.7. Data Management permissions 

⚠️🟡

caut…

This permission only offers None or View/Edit
access. While SOC Tier-2, Tier-3, and Threat
Hunters may need to view data schemas for
advanced investigations, granting them
View/Edit access gives them full control to
modify datasets, and change ingestion
pipelines. Grant this permission with caution
and ensure proper change management
protocols are in...

🧭➡️

navi…

Controls access to dataset configuration, data
transformation rules, and data lifecycle
management in Cortex AgentiX through
Settings → Configurations → Data Management

3.6.2.2.7.8. Public API 

3.6.2.2.7.9. Threat Intelligence permission - API
configuration  🧭➡️
navi…

Controls access to the configuration page for
external threat intelligence API keys (Virus
Total) on Settings → Configurations →
Integrations → Threat Intelligence

3.6.2.2.7.10. Long-running HTTP Integrations
configuration  📝🟢
note

This permission is used by administrators to set
up and maintain the EDL service infrastructure.

3.6.2.2.7.11. Integrations Permissions 

3.6.2.2.7.12. Credentials permissions 

📝🟢

note

To view the Credentials page, users require the
following view permissions:
Integrations
Data Sources
Without these permissions, users can't view the
Credentials page.

🧭➡️

navi…

Users can access Credentials on the
Credentials page by going to Settings →
Configurations → Integrations → Credentials

3.6.2.2.7.13. Case Properties permissions 

⚠️🟡

caut…

System-default statuses cannot be deleted,
and domain names cannot be changed after
creation.

🧭➡️

navi…

Controls access to the following tabs from
Cases (Settings → Configurations → Object
Setup → Cases

3.6.2.2.7.14. Exclusion List permissions 

📝🟢

note

When managing Indicators (located under
Threat Management → Threat Intelligence →
Indicators), users who lack View/Edit
permissions for the Exclusion List will find that
exclusion-related features (such as the
Exclusion reason field and Do not add to
exclusion list checkbox) are automatically
hidden by the system.

🧭➡️

navi…

Controls access to the indicator exclusion list
configuration under Settings → Configurations
→ Object Setup → Indicators → Exclusion List

3.6.2.2.7.15. Fields and Types permissions  🧭➡️

navi…

Controls access to custom fields and indicator
types within Object Setup Settings →
Configurations → Object Setup

3.6.2.2.7.16. Layout permissions  🧭➡️

navi…

Controls access within Object Setup (Settings
→ Configurations → Object Setup) to the
following

3.6.2.2.7.17. Sync Profile permissions  🧭➡️

navi…

Controls access to case mirroring profiles
configuration in Settings → Configurations →
Object Setup → Issues → Sync Profiles

3.6.2.2.8. Help permissions  ⚠️🟡

caut…

The tenant must have an active support
contract, and the user must have CSP portal
access.

3.6.2.2.9. Cortex Agentic Assistant permissions 

3.6.2.2.9.1. AI Prompts 

📝🟢

note

You need to add Playbooks View/Edit
permission to edit playbooks.

⚠️🟡

caut…

To effectively embed AI Prompts directly into
playbook workflows, users must be granted the
Manage prompts in playbook editor checkbox,
and they must also hold View/Edit permissions
for the Playbooks module.

🧭➡️

navi…

Controls access to the AI Prompts Library
(Investigation & Response → Automation → AI
Prompts), where users create and manage
reusable prompt templates (including system
instructions and few-shot examples) used to
guide the LLM's behavior

3.6.2.2.9.2. Cortex Agentic Assistant Agents  📝🟢

note

If Manage Agents is checked, the user can only
edit their own agents.

3.6.3. User group management 

📝🟢

note

If users have been created in the CSP, but you
want them to access the tenant through SSO
only, skip this field and add only SAML group
mapping after SSO is set up, otherwise, users
can access the tenant through both the CSP
and SSO.
If you have not yet created any users, skip this
field and add them later. See Set up
authentication .

📝🟢

note

When using Microsoft Entra ID for SSO, the
SAML group mapping needs to be provided
using the group object ID (GUID) and not the
group name.

👀🟠

imp…

The Cloud Identity Engine (CIE) is used
exclusively for directory group management; it
is not utilized for individual user account
management, provisioning, or authentication
workflows. Consequently, disabling, suspending,
or removing user objects directly within CIE
does not automatically disable, restrict, or
delete those corresponding users inside...

👀🟠

imp…

By default, Enable Scope Based Access
Control is disabled in Settings →
Configurations → General → Server Settings,
and granular scoping is not enforced. Before
enabling SBAC, we recommend that an
administrator or a user with Access
Management permissions first ensures that the
users, user groups, and API Keys defined in
Cortex AgentiX are...

3.6.4. User management 

📝🟢

note

If a user does not have a direct role or user
group assigned, the user is revoked and is not
saved in the Cortex Gateway.

📝🟢

note

If a user does not have a direct role or user
group assigned, the user is revoked and is not
saved in the Cortex Gateway.

3.6.4.1. Manage users in the Cortex AgentiX
tenant 

📝🟢

note

To remove users that were added to your CSP
account, you need to do this in the CSP and not
in the tenant or Cortex Gateway.

📝🟢

note

To update the permissions attributable to each
role, change them in the Roles tab or Cortex
Gateway.
If users have been created in the CSP, but you
want them to access the tenant through SSO
only, you should not assign a direct role. If you
sign a direct role, users can access the tenant
through both the CSP and SSO.

📝🟢

note

You cannot deactivate a user who has an
Account Admin role. If you want to deactivate
users from all tenants, deactivate them in
Cortex Gateway.
The user will be deactivated in the tenant, but
may still be active in other tenants. If you want
to deactivate the tenant for multiple tenants,
deactivate the user in Cortex Gateway.

📝🟢

note

When you remove a role, the role associated
with the API keys is deleted.
If more than one role was associated with the
API key, a yellow warning symbol appears next
to the API key in the API key table. When you
hover over the symbol, a message indicates that
some of the roles associated with the API key
have been deleted.
If all roles associated with...

3.6.4.2. Manage user scope 

☑️🟣

prer…

Configuring user scopes in Cortex AgentiX
Access Management requires View/Edit RBAC
permissions for Access Management (under
Configurations). Account Admin and Instance
Administrator roles are granted this permission
by default. For more information, see
Predefined user roles in Set up users and
roles.Set up users and roles
By default, Enable Scope...

📝🟢

note

Make sure to assign the required default
granular scoping for users. This depends on the
structure and divisions within your organization
and the particular purpose of each
organizational unit to which scoped users
belong.

👀🟠

imp…

Some areas and features in Cortex AgentiX do
not respect SBAC. In these cases, use RBAC
permissions to restrict access.

3.6.5. Manage access to objects 

📝🟢

note

Only the Owner or an Administrator can delete
a custom object.

👀🟠

imp…

While Per-object access controls the visibility
of the object (such as a dashboard or saved
query), the underlying data remains governed
by Scope-Based Access Control (SBAC). A user
must have the appropriate SBAC permissions to
view the data available through an object.

3.6.5.1. Manage access to custom dashboards 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for for all custom
objects, including dashboards:
Enable "Owners can Share objects they
created": Grants owners...

📝🟢

note

Only Account Admins and Instance
Administrators have the authority to change the
owner of an object.

📝🟢

note

If you're designated as an Editor, you can
always duplicate the widget and make your
changes on the copy.

3.6.5.2. Manage access to report templates 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for all custom
objects, including report templates:
Enable "Owners can Share objects they
created": Grants...

📝🟢

note

Only Account Admins and Instance
Administrators have the authority to change the
owner of an object.

3.6.5.3. Manage access to playbooks and scripts 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for all custom
objects, including report templates:
Enable "Owners can Share objects they
created": Grants...

📝🟢

note

For more information, see Push and pull content.

📝🟢

note

In a production (pull) tenant, be aware of the
following:
Users cannot be granted Editor permissions.
Any existing Editor permissions (set before the
tenant was configured for remote repository
synchronization) are effectively treated as
Viewer access.

📝🟢

note

Playbooks can only be Enabled when Scripts
are Enabled first.

📝🟢

note

To detach a playbook or script, a user must
have both the Create and the Edit Public sub-
permissions assigned to their role.
Reattaching an existing automation to a
Marketplace content pack can be performed
by the playbook or script Owner or an
Administrator.

3.6.5.4. Manage access to saved queries 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for saved queries
in the Query Library, including the options that
appear to users when clicking the three dot,...

📝🟢

note

If the role of a user is set to Edit Public Queries
but not Create Queries, they can update
existing public queries, but the Save as drop-
down menu will be hidden, preventing them
from creating new Query Library entries.

📝🟢

note

When the tenant-level setting Owners and
editors can change the general access is
unselected, the drop-down is disabled and only
an administrator can configure this option.

3.6.6. Set up authentication 

💡🔵

tip

You should have at least one user in the
Customer Support Portal for backup, in case of
any authentication issues with your IdP provider.

3.6.6.1. Authenticate users through the
Customer Support Portal 

📝🟢

note

You must be assigned the Super User role in the
Customer Support Portal to add users in the
Customer Support Portal.
The first Super User who logs into Cortex
Gateway is automatically assigned the Account
Admin role and has access to the tenant. The
user who activates the Cortex AgentiX tenant
will also be assigned the Account Admin role (if
there is...

📝🟢

note

After users are added to the Customer Support
Portal and they accept the invitation, you can
manage them in Cortex Gateway or the Cortex
AgentiX tenant.

3.6.6.2. Authenticate users using SSO 

📝🟢

note

To set up SSO authentication in the tenant, you
must be assigned an Instance Administrator or
Account Admin role.
SAML 2.0 users must log in to Cortex AgentiX
using the FQDN (full URL) of the tenant. To
allow login directly from the IdP to , you must
set the relay state on the IdP to the FQDN of
the tenant.
If you have multiple tenants, you must set up...

📝🟢

note

The first SSO cannot be deleted, it can only be
deactivated by toggling SSO Enabled to off.
The Domain parameter is predefined for the
first SSO.
If you add additional SSO providers, you must
provide the email Domain in the SSO Integration
settings for all providers except the first.
Cortex AgentiX uses this domain to determine
to which identity...

📝🟢

note

Cortex AgentiX requires the IdP to send the
group membership as part of the SAML token.
Some IdPs send values in a format that include
a comma, which is not compatible with Cortex
AgentiX. In that case, you must configure your
IdP to send a single value without a comma for
each group membership. For example, if your
IdP sends the Group DN (a...

👀🟠

imp…

PROGRAMMATIC CONTRAINT:
There is no public API endpoint available to
provision or de-provision users
programmatically within Cortex AgentiX. All
target accounts must be initialized or explicitly
managed using the native interactive Single
Sign-On (SSO) or Customer Support Portal
(CSP) interface workflows defined in this guide.
To review the list of...

3.6.6.3. Set up Okta as the Identity Provider
Using SAML 2.0  📝🟢
note

When using SAML 2.0, users are required to
authenticate by logging in directly at the tenant
URL. They cannot log in via Cortex Gateway.

3.6.6.4. Set up Microsoft Entra ID as the Identity
Provider Using SAML 2.0 

📝🟢

note

The default attributes shown on the main single
sign-on page in Microsoft Entra ID are not the
values you need. You must click Edit next to
Attributes and Claims to view and copy the
actual values.

📝🟢

note

The attribute claim name must exactly match
the value sent by your IdP. In some cases, this
may be the full attribute name/namespace,
depending on the configuration of our IdP

📝🟢

note

When using SAML 2.0, users are required to
authenticate by logging in directly at the tenant
URL. They cannot log in via Cortex Gateway.

👀🟠

imp…

When copying the Single Sign-On URL value,
remove idp/saml and leave the trailing /.
For example, if the Single Sign-On URL is
https://clientname.panproduct.region.paloaltone
tworks.com/idp/saml, just copy
https://clientname.panproduct.region.paloaltone
tworks.com/.

3.7. Engines 

3.7.1. What is an engine? 

📝🟢

note

You cannot share a multiple-engine installation
with a single-engine installation.

📝🟢

note

When you add an engine to a load-balancing
group, you cannot use that engine separately.
The engine does not appear in the engines
menu when configuring an integration instance,
but you can choose the load-balancing group.

3.7.2. Engine requirements 

📝🟢

note

The Cron package is required to install engines
on a Linux machine.

📝🟢

note

CentOS 8.x reached End of Life (EOL) on
December 31, 2021, and is no longer supported
as an operating system.
CentOS 7.x reached End of Life (EOL) on June
30, 2024, and is no longer supported as an
operating system.

📝🟢

note

If you have configured a range of Approved IP
Ranges under Allowed Sessions on the Security
Settings page, the engine must communicate
through one of the approved IPs.

3.7.3. Install an engine 

📝🟢

note

If you are using DEB, RPM, or Zip installation,
install Docker or Podman.
Natively running Python or PowerShell
integrations/scripts on Windows or Linux is not
supported on Cortex AgentiX engines.

📝🟢

note

When upgrading an engine that was installed
using the Shell installation, you can use the
Upgrade Engine feature in the Engines page. For
Amazon Linux 2 type engines, you need to
upgrade these engine types using a zip-type
engine and not use the Upgrade Engine feature.
If you use the shell installer, Docker/Podman is
automatically installed. We...

📝🟢

note

Use DEB and RPM installation when the shell
installation is not available. You need to
manually install Docker or Podman and any
dependencies.

📝🟢

note

If you receive a permissions denied error, it is
likely that you do not have permission to access
the /tmp directory.
If the installer fails to start due to a permissions
issue, even if running as root, add one of the
following two arguments when running the
installer:
--target <path> - Extracts the installer files
into the specified custom...

👀🟠

imp…

For DEB/RPM engines, Python (including 3.x)
and the containerization platform
(Docker/Podman) must be installed and
configured. For Docker or Podman to work
correctly on an engine, IPv4 forwarding must be
enabled.

💡🔵

tip

For Linux systems, we recommend using the
shell installer. If using Amazon Linux 2, use the
zip installer (see step 4).

3.7.3.1. Docker 

📝🟢

note

This section is relevant when installing an engine.

📝🟢

note

Docker images can be downloaded together
with their relevant content packs for offline
installation.

3.7.3.1.1. Install Docker 

📝🟢

note

For Red Hat's Docker distribution, you need
Mirantis Container Runtime (formerly Docker
Engine - Enterprise) to run specific Docker-
dependent integrations and scripts. For more
information, see Install Docker distribution for
Red Hat on an engine server.
To use the Mirantis Container Runtime (formerly
Docker Engine - Enterprise) follow the
deployment...

3.7.3.1.1.1. Install Docker distribution for Red Hat
on an engine server

 📝🟢

note

If running RHEL v8 or higher, the engine installs
Podman packages and configures the operating
system to enable Podman in rootless mode.

3.7.3.1.1.2. Docker image security 

3.7.3.1.1.3. Docker FAQs 

3.7.3.1.1.4. Troubleshoot Docker issues 

3.7.3.1.1.5. Configure Docker pull rate limit 

3.7.3.1.1.6. Change the Docker installation folder 

3.7.3.1.2. Docker hardening guide 

📝🟢

note

The settings below can also be applied to
Podman, with the exception of limiting available
memory, limiting available CPU, and limiting
PIDS.

📝🟢

note

If you do not want to apply Docker memory
limitations, you should explicitly set the
advanced parameter: limit.docker.memory to
false.

💡🔵

tip

We recommend reviewing Docker network
hardening below before changing any
parameters in the configuration file.
To securely run Docker containers, we
recommend using the latest Docker version.

3.7.3.2. Podman 

📝🟢

note

When upgrading an engine, the engine keeps
the previously used container management
type (regardless of distribution version).

3.7.3.2.1. Change the container storage directory 

📝🟢

note

Alternatively, the runroot setting may be set to
some temporary directory that is accessible by
the user demisto. If you choose to set the
runroot, it must be a directory that is mounted
as tmpfs (temporary filesystem), unlike the
graphroot.

⛔🔴

war…

Do not use NAS storage or a temporary (tmpfs)
directory for the graphroot setting. The
graphroot needs to be a local, non-temporary
directory for Podman to work. For more
information, see
https://en.wikipedia.org/wiki/Network-
attached_storage.

💡🔵

tip

We recommend reserving 150 GB for container
storage, either in the /home partition or a
different storage directory that you have set
using the graphroot key.

3.7.3.2.2. Install Podman 

📝🟢

note

This procedure is intended for RHEL 8 or later. It
may not work for other operating system types.
Do not use NAS storage for the $HOME
directory. The directory needs to be a local
directory for Podman to work.

📝🟢

note

If you edit the file with the root user, make sure
to set the demisto user as file owner by running
chown demisto:demisto
/home/demisto/.config/containers/registries.con
f

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

3.7.3.2.3. Migrate From Docker to Podman 

3.7.3.2.4. Troubleshoot Podman 

📝🟢

note

This step removes all Podman images, including
any custom images you may have created.

📝🟢

note

$(id -u demisto) is used to get the demisto user
ID, which is part of the directory name. For
example, /tmp/podman-run-993
Not all the directories above may be present.

📝🟢

note

When you run podman container cleanup --rm -
a, you might see a message such as running or
paused containers cannot be moved without
force. The message can be safely ignored, as it
only pertains to current running containers,
which are not removed.

📝🟢

note

To change to the demisto OS user, run the
following command:
sudo su - -s /bin/bash demisto

📝🟢

note

You must switch to user demisto before running
the "system migrate" (running it as root will
have no effect).

📝🟢

note

The "system migrate" must be done by the user
demisto.

👀🟠

imp…

The runroot must be located under the tmpfs
file system type. This is required to clean
Podman's run state on reboot and for
performance reasons.

⚠️🟡

caut…

Unlike the runroot, the graphroot must NOT be
located under the tmpfs file system type. Using
tmpfs for the graphroot might corrupt
container images, causing command executions
to fail. It also degrades performance by forcing
Podman to needlessly re-pull images.

3.7.4. Manage engines  🧭➡️

navi…

You can manage your engines and load-
balancing groups by going to Settings →
Configurations → Engines

3.7.5. Upgrade an engine 

3.7.6. Remove an engine 

3.7.7. Configure engines 

📝🟢

note

A day is defined as 24 hours and may not
exactly correspond to calendar days due to
daylight savings, leap seconds, etc.

📝🟢

note

In addition, to support engine upgrades from
the UI, edit the /usr/local/demisto/upgrade.conf
file on the engine to include the SERVER_URLS
setting with the new tenant's address. Include
only the host, without https:// or any additional
path at the end of the host name. For example:
SERVER_URLS="api-
example.us.paloaltonetworks.com"

📝🟢

note

If the engine loses communication for longer
than this time, it will disconnect and you need
to restart the service.

3.7.7.1. Configure the engine to use a web proxy 

📝🟢

note

You need to configure Docker to use a proxy.
When using a BlueCoat proxy, ensure you
encode the values correctly.

📝🟢

note

In an environment with a single engine, go to
/usr/local/demisto/upgrade.conf. In an
environment with multiple engines, go to
/usr/local/demisto/<engine-
name>/upgrade.conf, replacing <engine-name>
with the name of the engine.
Note that the key is in the upgrade.conf file and
must be https_proxy, even if your proxy address
starts with http://.

3.7.7.2. Configure the engine to call the server
without using a proxy

 3.7.7.2.1. Use NGINX as a reverse proxy 

3.7.7.3. Configure an engine to use custom
certificates 

3.7.8. Use an engine in an integration  📝🟢

note

Long-running integrations should not run on
load-balancing groups.

3.7.9. Run a script using an engine  🧭➡️

navi…

Investigation & Response → Automation →
Scripts

3.7.10. Troubleshoot engines 

📝🟢

note

Ensure that pop-ups are not blocked by your
browser.

📝🟢

note

If the installer fails to start due to a permissions
issue, even if running as root, add one of the
following two arguments when running the
installer:
--target <path> - Extracts the installer files
into the specified custom path.
--keep - Extracts the installer files into the
current working directory (without cleaning at
the end).
If using installer...

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

📝🟢

note

Any changes made to variables in the
upgrade_engine.sh file are reset after each
upgrade. We recommend instead using the
upgrade.conf file to set variables.

📝🟢

note

You can ignore the following error: Cannot
create folder '/var/lib/demisto'

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

📝🟢

note

If the installer fails to start due to a permissions
issue, even if running as root, add one of the
following two arguments when running the
installer:
--target <path> - Extracts the installer files
into the specified custom path.
--keep - Extracts the installer files into the
current working directory (without cleaning at
the end).
If using installer...

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

🧭➡️

navi…

When troubleshooting engines, access the logs
from Settings → Configurations → Engines and
select the engine from which you want to
download the logs

3.7.11. Troubleshoot integrations running on
engines 

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

📝🟢

note

Docker CE installations typically run Docker,
while Docker EE installations typically run
dockerroot.

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

3.8. Cases and issues configuration 

3.8.1. Case and issue lifecycle 

3.8.2. Customize your cases 

3.8.2.1. External integrations 

3.8.2.2. Create a case domain 

📝🟢

note

Adding custom domains requires a View/Edit
RBAC permission for Case Properties (under
Object Setup).
Once created, a custom case domain cannot be
deleted or renamed.

⛔🔴

war…

Before you add a custom domain, please review
the built-in options. For more information, see
Case and issue domains.
We recommend using the built-in domains
where possible. Custom domains might not be
supported by all content. In addition, custom
domains affect Cortex AgentiX's ability to
learn, correctly identify, and score future cases.
Smart...

🧭➡️

navi…

You can see all domains under Configurations
→ Object Setup → Cases → Domains

Settings → Configurations → Object Setup →
Cases → Domains

Settings → Configurations → Access
Management → User Groups

Settings → Configurations → Access
Management → Users

To enable access to the domain for an API key,
got to Settings → Configurations →
Integrations → API Keys

3.8.2.3. Create a starring configuration  💡🔵

tip

Right-click an issue field to add it as match
criteria.

3.8.2.4. Set up case scoring 

📝🟢

note

For scoped users, a small lock icon indicates
that you don't have permissions to edit a rule.

💡🔵

tip

Right-click an issue field to add it as match
criteria.

3.8.2.5. Create custom case statuses and
resolution reasons



📝🟢

note

Before you create a custom status, please
review the built-in options. For more
information, see Resolution reasons for cases
and issues.
We recommend using the built-in statuses and
resolution reasons where possible. Custom
statuses and resolution reasons might not be
supported by all content, and status syncing
can take time.
In addition, custom...

📝🟢

note

After creation, custom statuses and resolution
reasons cannot be deleted or modified.

🧭➡️

navi…

Configurations → Object Setup → Cases →
Properties

3.8.3. Customize issue fields and layouts 

3.8.3.1. Issue fields 

⛔🔴

war…

Deleting an issue field or uninstalling a content
pack containing an issue field may affect
detection and other capabilities based on the
deleted field. For example, correlation, layouts,
case scoring, starring rules, and playbook
triggers.

3.8.3.1.1. Issue field types 

📝🟢

note

You cannot set filters, starring rules,
automation rules, layout rules, or issue
exclusions based on the values in custom
timestamp fields.

📝🟢

note

When grid field is shown in the Issues table, if
there are values in the field, they do not display
in the table. Instead, the column shows Data
Available.

📝🟢

note

When an HTML field is shown in the Issues
table, if there is a value in the field, it does not
display in the Issues table. Instead, the column
shows Data Available.
The following HTML tags are not permitted:
blockquote, del, dd, div, dl, dt, fieldset, form, h1,
h2, h3, h4, h5, h6, hr, iframe, ins, li, math,
noscript, ol, pre, p, script, style, table,...

📝🟢

note

When a Markdown field is shown in the Issues
table, if there is a value in the field, it does not
display in the table. Instead, the column shows
Data Available.

3.8.3.1.2. Create custom issue fields 

🧭➡️

navi…

Settings → Configurations → Object Setup →
Issues → Fields → New Field

3.8.3.1.2.1. Create a grid field for an issue 

3.8.3.1.2.2. Issue field triggered scripts 

📝🟢

note

Fields that can hold a list, such as multi-select
custom fields, return the delta in an array as a
new argument. For example, if a multi-select
field value has changed from ["a"] to ["a", "b"],
the new argument of the script gets a value of
["b"].

📝🟢

note

Issue field-triggered scripts must have the field-
change-triggered tag to appear in the list.
Issue field trigger scripts are not supported for
all system fields. The following fields are not
supported for issue field trigger scripts and
may result in failing to populate the issue
layout:
Cases: Case ID, Cases IDs
Asset Fields: Asset IDs, Asset Names,...

📝🟢

note

When creating the script, add the field-change-
triggered tag in the script settings.

📝🟢

note

Verify that User can add rows is selected.

⚠️🟡

caut…

This feature assumes fair and intended usage
of field-triggered scripts. In cases of excessive
or abusive usage, execution may be restricted
or disabled. If script execution is restricted or
disabled, fields are still updated, but without
the results of the assigned script.

3.8.3.1.2.3. Issue timer fields 

🧭➡️

navi…

Settings → Configurations → General → Server
Settings → Issues and Enable Timer Field

3.8.3.1.2.3.1. Configure issue timer fields 

3.8.3.1.2.3.2. Configure a playbook to run timers 

3.8.3.1.2.3.3. Automate changes to issue fields
using timer scripts 

3.8.3.1.2.3.4. Use issue timer field commands

manually in the CLI 

3.8.3.2. Issue layouts 

3.8.3.2.1. Create custom issue layouts 

📝🟢

note

The script that runs when an action button is
clicked accepts only mandatory arguments
through the pop up window and does not
provide an option for any non-mandatory
arguments to be filled in when the button is
clicked. We recommend using a wrapper script
to collect and validate arguments in scenarios
where there can be a combination of
mandatory and...

💡🔵

tip

Limit the number of issue fields to 50 in each
section. You can create additional sections as
needed.

3.8.3.2.2. Add a custom widget to an issue layout  🧭➡️

navi…

Investigation & Response → Automation →
Scripts and upload the following script: Expand
script commonfields: id: ee3b9604-324b-4ab5-
8164-15ddf6e428ab version: 49 name:
IndicatorWidgetBar script: |- # Constants HIGH
= 3 SUSPICIOUS = 2 LOW = 1 NONE = 0
indicators = [] scores = {HIGH: 0, SUSPICIOUS: 0,
LOW: 0, NONE: 0} incident_id = demisto

Investigation & Response → Automation →
Scripts and upload the following script

Settings → Configurations → Object Setup →
Issues → Layouts

3.8.3.2.3. Create rules for issue layouts  🧭➡️

navi…

Settings → Configurations → Object Setup →
Issues → Layout Rules → New Rule

3.8.4. Customize case fields and layouts 

3.8.4.1. Case fields 

⛔🔴

war…

Deleting a case field or uninstalling a content
pack containing a case field may affect
capabilities based on the deleted field and
layouts.

3.8.4.1.1. Case field types 

📝🟢

note

You cannot set filters, starring rules,
automation rules, layout rules, or issue
exclusions based on the values in custom
timestamp fields.

📝🟢

note

When grid field is shown in the Cases table, if
there are values in the field, they do not display
in the Cases table. Instead, the column shows
Data Available.

📝🟢

note

When an HTML field is shown in the Issues
table, if there is a value in the field, it does not
display in the Cases table. Instead, the column
shows Data Available.

📝🟢

note

When a Markdown field is shown in the Cases
table, if there is a value in the field, it does not
display in the Cases table. Instead, the column
shows Data Available.

📝🟢

note

If you make changes to case fields you can
update the context data by running a playbook,
script, or command. For more information, see
Update case fields.
To update dynamic custom case fields, such as
SLA and Timer fields, see Update case timer
and SLA fields.

3.8.4.1.2. Create custom case fields 

🧭➡️

navi…

Settings → Configurations → Object Setup →
Cases → Fields → New Field

3.8.4.1.2.1. Create a grid field for a case 

3.8.4.1.2.2. Update case fields  📝🟢

note

Ensure that you have the required RBAC
permission to write scripts.

3.8.4.1.3. Create case timers and SLAs 

☑️🟣

prer…

Before you can create a case SLA, you must
first create a timer field. A timer field can be
associated with a single case SLA.

📝🟢

note

If you edit this filter after creation, the timer
and associated SLA will be removed from any
case that no longer qualifies, even if the timer is
already running.

🧭➡️

navi…

Settings → Configurations → Object Setup →
Cases and open the Fields tab

3.8.4.1.4. Update case timer and SLA fields 

3.8.4.2. Case layouts 

3.8.4.2.1. Create custom case layouts 

📝🟢

note

The script that runs when an action button is
clicked accepts only mandatory arguments
through the pop up window and does not
provide an option for any non-mandatory
arguments to be filled in when the button is
clicked. We recommend using a wrapper script
to collect and validate arguments in scenarios
where there can be a combination of
mandatory and...

💡🔵

tip

Limit the number of case fields to 50 in each
section. You can create additional sections as
needed.

3.8.4.2.2. Create rules for case layouts 

📝🟢

note

Layout rules support SBAC (scoped based
access control). The following parameters are
considered for editing access.
If Scope-Based Access Control (SBAC) is
enabled and Endpoint Scoping Mode is set to
restrictive mode, you can edit a rule if you are
scoped to all tags in the rule.
If Scope-Based Access Control (SBAC) is
enabled and Endpoint Scoping...

🧭➡️

navi…

Settings → Configurations → Object Setup →
Cases → Layout Rules → New Rule

3.8.5. Issue syncing 

☑️🟣

prer…

You must set up an integration before you can
sync issues. For more information, see Set up an
integration for mirroring issues.

☑️🟣

prer…

You must set up an integration before you can
sync issues. For more information, see Set up an
integration for mirroring issues.

☑️🟣

prer…

You must set up an integration before you can
sync issues. For more information, see Set up an
integration for mirroring issues.

📝🟢

note

Multiple tickets can be linked to an issue with
outbound syncing. Issues with inbound syncing
can be linked to a single ticket only.

📝🟢

note

Using issue fields as variables is not currently
supported.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

📝🟢

note

Using issue fields as variables is not currently
supported.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

📝🟢

note

The inbound syncing flow runs every two
minutes, and the outbound syncing flow runs
every five minutes.
In a bi-directional set-up, if the same field is
updated in both tickets, the most recently
updated value is used.
In the external ticket, the logged history shows
updates to the ticket. The user name that is
logged with the history reflects the user...

📝🟢

note

After an issue is resolved, ticket syncing
remains active for up-to seven days. Therefore,
you still update, change, or reopen the issue or
external ticket and the tickets will continue to
sync.

📝🟢

note

If you change the selected inbound sync
profile, the original sync profile is immediately
overwritten.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

⛔🔴

war…

If you leave this field blank, all configured
instances will be used.

⛔🔴

war…

If you leave this field blank, all configured
instances will be used.

💡🔵

tip

You can find a sync profile ID under Settings →
Configurations → Object Setup → Issues →
Sync Profiles. By default the ID field is not
displayed in the table. Click the three dot menu
and add it to the table layout.

3.8.5.1. Create a sync profile 

📝🟢

note

If an issue is using bi-directional syncing, you
need to provide both an Inbound and an
outbound sync profile.

📝🟢

note

Blank fields are skipped.
You must define exact values.
Custom status values are not currently
supported.
Support is currently limited to a specific set of
fields.

3.8.6. Create issue exceptions 

☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete issue exception
approvers and to turn off and on approver
functionality.
You must have Exception Management Admin
permission to create issue exception rules,
delete or disable exception rules, and view the
All Exception Rules page.

📝🟢

note

Issue exceptions apply specifically to issues,
not individual findings.
Issue exceptions do not automatically update
standard dashboards, reporting widgets, or
compliance reporting profiles. To view the
aggregate impact of exceptions, you must
manually filter the issues table.
Exceptions do not alter Attack Path policies.
Exceptions do not bypass any...

3.8.6.1. Configure the issue exception approval
workflow



☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete approvers.

☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete approvers.

☑️🟣

prer…

You must have the Exception Approver Admin
permission to disable or re-enable the issue
exception approval workflow.

🧭➡️

navi…

Settings → Configurations → Server Settings
and scroll to the Exception Management section

Settings → Configuration → General → Server
Settings and scroll to the Exception
Management section

3.8.6.2. Create an issue exception rule  📝🟢

note

You can create a maximum of 10 issue
exceptions per day.

3.8.6.3. Create an exception rule from an issue  📝🟢

note

You can create a maximum of 10 issue
exception rules per day.

3.8.6.4. View issue Exception Rules  🧭➡️

navi…

Settings → Exceptions Configuration →
Exception Rules

3.8.6.5. Disable issue exception rules 

📝🟢

note

You cannot re-enable an issue exception rule
after it has been disabled.

🧭➡️

navi…

Settings → Exceptions → Configuration →
Exception Rules

3.8.6.6. View excepted issues 

3.8.7. What's a correlation rule? 

🧭➡️

navi…

After you configure your correlation rules, you
can manage them in Threat Management →
Detection Rules → Correlations, and view and
analyze the generated issues in Cases and the
Issues Table

3.8.7.1. Correlation rule details 

📝🟢

note

Only an administrator can create and view
queries built with an unknown dataset that
currently does not exist in Cortex AgentiX .

📝🟢

note

For more information on troubleshooting server
errors in scheduled correlation rules,
Troubleshoot server errors in scheduled
correlation rules.Troubleshoot server errors in
scheduled correlation rules

📝🟢

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

📝🟢

note

Only an administrator can create and view
queries built with an unknown dataset that
currently does not exist in Cortex AgentiX.

🧭➡️

navi…

If you are assigned a role that enables
Investigation → Rules privileges, you can
manage all user-defined Correlation Rules from
Threat Management → Detection Rules →
Correlations

3.8.7.2. Create a correlation rule 

☑️🟣

prer…

To enable pivots from issues to a third-party
source system, ensure that you know the
dataset field name that contains the URL of the
source system. Some vendors already provide
this URL as part of their API, and you can find it
in the third-party product dataset. If there is no
URL, you cannot enable this feature.

📝🟢

note

When creating a Real Time Correlation Rule,
you can only generate an issue as the resulting
action for the Correlation Rule. All other
options are disabled.
To ensure your Correlation rules generate
issues efficiently and do not overcrowd your
Issues table, Cortex AgentiX automatically
disables Correlation rules that reach 5000 or
more hits over a...

📝🟢

note

When you open the New Correlation Rule
editor from XQL Search, this XQL Search field is
already populated with the XQL query that you
defined.
An administrator can create and view queries
built with an unknown dataset that currently
does not exist in Cortex AgentiX . All other
users can only create and view queries built
with an existing dataset.

📝🟢

note

The XQL features for call, top, and wildcards in
datasets (dataset in (<dataset prefix>_*)) are
currently not supported in Correlation Rules. If
you add them to the XQL definition, you will not
be able to Create or Save the Correlation Rule.
The XQL features for transaction in datasets
(dataset in (<dataset prefix>_*)) are currently
not supported in...

📝🟢

note

The incident that contains the issues will also
be assigned to the selected domain. For more
information, see Case and issue domains.
If you want to generate a health issue, choose
the Health domain. For more information about
health issues, see About health issues.

📝🟢

note

If a generated issue has severity Medium or
above, a case is automatically opened.
Low severity issues generated by correlation
rules are not grouped into cases.

📝🟢

note

There is no validation or auto complete for
these parameters and the values can be null or
empty. In these scenarios, Cortex AgentiX does
not display the null or empty values, but adds
the text NULL or EMPTY in the descriptions.

📝🟢

note

Cortex AgentiX groups into cases based on a
grouping logic that evaluates relationships,
context, and shared artifacts across issues.
However, not all fields influence case grouping.
For information about the specific fields that
influence grouping, see Optimize case grouping
in correlations.Optimize case grouping in
correlations

📝🟢

note

When more than one XDM field is listed, Cortex
AgentiX looks for the field value according to
the order of the fields listed.

📝🟢

note

When more than one XDM field is listed, Cortex
AgentiX looks for the field value according to
the order of the fields listed.

📝🟢

note

When Use the Cortex AgentiX default incident
enrichment is not selected and you have not
mapped any issue fields, the issue is dispatched
into a new incident.

👀🟠

imp…

The maximum size of a lookup dataset is 50 MB.
If the data exceeds this limit, the add to lookup
action fails.

💡🔵

tip

When creating your correlation rule, you can
use predefined values for different fields in the
editor, such as Alert Name, Alert Description,
and Drill-Down Query. For more information,
see Field replacement syntax in correlation
rules.Field replacement syntax in correlation
rules

🧭➡️

navi…

You can create a new correlation rule from
either the Threat Management → Detection
Rules → Correlation Rules page or when
building a query in XQL Search

3.8.7.3. Manage correlation rules  📝🟢

note

Execute Rule is not available for real time
correlation rules.

3.9. Automations 

3.9.1. Automation in Cortex AgentiX  📝🟢

note

You can build end-to-end automation
workflows from within the playbook editor,
including creating automation rules,
configuring integration instances, and creating
and editing tasks. For more information, see
Playbooks.

3.9.2. Quick Actions  📝🟢

note

Quick Actions appear as War Room entries, but
do not appear in the Work Plan.

3.9.3. Automation Exclusion Center 

📝🟢

note

By default, all users have read and edit
permissions to lists. When creating a list of
critical assets, we recommend limiting the read
and edit permissions to specific roles.

3.9.3.1. Manage automation exclusion policies 

📝🟢

note

For the IAM User Hard Remediation and User
Soft Remediation policies, we recommend
including username, email, and ID for each user
you want to exclude. Example: username1,
user@example.com, userID112.

📝🟢

note

You can also right click on a policy from the
main Automation Exclusion Center page to
disable or enable the policy.
If you click on a list name in the Exclude
column, that list opens in the Lists page.

3.9.4. Playbooks 

☑️🟣

prer…

To work with playbooks, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Scripts: This component (under Investigation &
Response → Automations) must be set to
Enabled first. It is the foundational permission
for all automation; if Scripts are not enabled,
you cannot...

3.9.4.1. Playbooks overview 

3.9.4.2. Access to playbooks 

3.9.4.3. Playbook development checklist 

3.9.4.4. Plan your playbook 

3.9.4.5. Manage playbooks  📝🟢

note

The library by default shows only playbooks
that are not adopted. Click the Show Adopted
checkbox to show the adopted playbooks,
indicated by an Adopted mark.
The library shows the most updated playbook
version. Adopting an older version than shown
should be done through Marketplace.
Adopting a playbook does not make it run.
Some content packs include...

3.9.4.6. Build your playbook 

👀🟠

imp…

The ability to create, edit, or share custom
playbooks is governed by access management.
If certain options are unavailable, contact your
administrator. For more information, see
Manage access to playbooks and scripts.

3.9.4.6.1. Task 1. Choose from existing playbooks
or create your own



📝🟢

note

When a task inside of a sub-playbook is not
configured, the alert is propagated to the main
playbook. If multiple sub-playbooks are nested,
and any of the sub-playbooks have non-
configured components, the alert appears in
the main playbook as well as in the sub-
playbooks. Alerts also appear for the individual
non-configured tasks within the...

📝🟢

note

You must have the required permissions to
create playbooks to view this button. For more
information, see Manage access to playbooks
and scripts.

💡🔵

tip

If there are additional relevant playbooks in
Marketplace that are not in your Org repository,
you can click Explore them now to see them in
the Playbook Catalog and choose to adopt.

💡🔵

tip

To open multiple playbooks at the same time,
edit the first playbook and then click New next
to the playbook name to create a new tab. You
can either create a new playbook, or add an
existing one.
You can view recently modified or deleted
playbooks by clicking version history for all
playbooks .

🧭➡️

navi…

Open the Investigation & Response →
Automation → Playbooks page to find an
existing playbook, customize it, or create a
playbook

3.9.4.6.2. Task 2. Configure playbook settings 

📝🟢

note

If the playbook has inputs and outputs, the
Playbook Starts task will show back and forth
arrows. Clicking them opens the Playbook
Settings pane Inputs/Outputs tab.

📝🟢

note

For more information on Automation Rules, see
Create an automation rule.

📝🟢

note

This rule will trigger the playbook to run if no
other Automation Rule triggers the playbook
first. You can view and edit the order the rules
run in the Automation Rules page.

📝🟢

note

If you do not add any fields, the group will be
deleted when you click Save.

3.9.4.6.3. Task 3. Add objects from the Task
Library



3.9.4.6.3.1. Add commands and scripts 

📝🟢

note

If you want to add a script that is not yet
adopted, Cortex AgentiX automatically installs
the content pack containing the script. If the
script requires an integration instance, you are
prompted to configure one.

📝🟢

note

The output value is dynamic and is derived
from the context at the time that the task is
processed. As a result, parallel tasks that are
based on the same output may return
inconsistent results.

💡🔵

tip

If your needs require a custom script, use the
Agentic Assistant with the Automation Engineer
agent to leverage the Cortex AgentiX built-in
LLM to quickly and efficiently generate
functional Python scripts from natural
language prompts. For more information, see
Create a script.

3.9.4.6.3.2. Add sub-playbooks 

📝🟢

note

To modify the settings or task configurations
for system playbooks (out-of-the-box
playbooks), your role must have the Edit Public
Playbooks permission enabled. If this permission
is not enabled, system playbooks will remain
read-only. For more information, see Manage
access to playbooks and scripts.

💡🔵

tip

Balance between the number of iterations and
the interval so you do not overload the server.

3.9.4.6.3.3. Add AI prompt tasks 

📝🟢

note

Selecting Local AI Prompt creates a local
version of the task within your playbook. Any
future updates made to the Prompts Library
will not sync to this specific task.
Choose Local AI Prompt if you need to
customize the logic for a specific workflow and
want to ensure its behavior remains unchanged,
even if the Prompts Library is updated with...

💡🔵

tip

If an AI prompt task returns a 504 Error with the
status DEADLINE_EXCEEDED, it is likely due to
the task requiring more time to process than
the default setting allows. To resolve this, open
the Task Details pane of the AI task and select
the Advanced tab, then increase the Execution
timeout (seconds) based on how complex your
prompt is and how long...

3.9.4.6.3.4. Add manual tasks and blank tasks 

3.9.4.6.3.4.1. Create a standard task 

3.9.4.6.3.4.2. Create a conditional task 

3.9.4.6.3.4.3. Create a communication task 

3.9.4.6.3.5. Create a section header 

3.9.4.6.3.6. Configure script error handling in a
playbook 

📝🟢

note

Use the GetErrorsFromEntry script (part of the
Common Scripts Pack) to check whether the
given entry returns an error and returns an error
message. For example, when using the script in
a playbook, you can fetch the error message
from a given task, such as a runtime error. You
can then add a step in the playbook flow to
send those error messages to the...

📝🟢

note

You can set up script error handling only when
running a script in a Standard task or a
Conditional task. For more information about
error handling settings for these tasks, see
Create a standard task or Create a conditional
task.
Built-in, Manual, and Ask Conditional tasks have
On Error settings for number of retries and
retry interval, but not Error...

3.9.4.6.4. Task 4. Add custom playbook features 

3.9.4.6.5. Task 5. Test and debug the playbook 

3.9.4.6.6. Task 6. Manage playbook content 

3.9.4.7. Customize your playbook 

3.9.4.7.1. Configure a sub-playbook loop 

📝🟢

note

Consider the following when adding a loop:
The maximum number of loops (default is 100).
A high number of loops or a high wait time
combined with a large number of issues may
affect performance.
Periodically check looping conditions to ensure
they are still valid for the data set.
If you want a sub playbook task to loop over an
array passed into its...

💡🔵

tip

Balance between the number of iterations and
the interval so you do not overload the server.

3.9.4.7.2. Filter and transform data 

⚠️🟡

caut…

You can change the context data root to filter,
but it is not recommended to select a different
root, as it affects the filter results. The drop-
down list displays the filter root for backward
compatibility.

🧭➡️

navi…

Investigation & Response → Automation →
Scripts → New Script

3.9.4.7.2.1. Filter considerations, categories, and
built-in filters



📝🟢

note

Is defined considers false and empty strings
and lists to be defined values. If you don't want
those to be included as defined, use Is not
empty.

📝🟢

note

Not defined considers false and empty strings
and lists to be defined values. If you don't want
those to be included as defined, use Is empty.

3.9.4.7.2.2. Transformer considerations,
categories, and built-in transformers

 📝🟢

note

To make regex case non-sensitive, use the (?i)
prefix (for example (?i)yourRegexText.

3.9.4.7.3. Extend context 

3.9.4.7.4. Extract indicators 

📝🟢

note

This configuration may delay playbook
execution (issue creation).
While indicator creation is asynchronous,
indicator extraction and enrichment are run
synchronously. Data is placed into the issue
context and is available via the context for
subsequent tasks.

📝🟢

note

When using out of band, the extracted
indicators do not appear in the context. If you
want the extracted indicators to appear select
inline.

3.9.4.7.5. Update issue fields with playbook tasks  📝🟢

note

The setIssue script includes all available fields;
use the scroll bar to see all the fields.
The name field has a limit of 600 characters. If
there are more than 600 characters, you can
shorten the name field to under 600 characters
and then include the full information in a long
text field such as the description field.
There are many fields already...

3.9.4.8. Test your playbook 

📝🟢

note

The debugger does not support using
parentIncidentFields.

📝🟢

note

Using an existing issue in the debugger does not
affect the original issue or change the original
context data.

3.9.4.8.1. Troubleshoot playbook performance 

3.9.4.9. Manage playbook content 

📝🟢

note

Opening a playbook does not immediately lock
it. A playbook is considered to be in edit mode
only when a user makes a modification that
causes the Save button to become available.

👀🟠

imp…

Manually unlocking a playbook may cause
version conflicts. This turns off concurrent
editing protection, and if the original user is still
editing, their changes might be lost or
overwritten.

🧭➡️

navi…

Investigation & Response → Automations →
Playbooks

3.9.4.10. Best practices 

📝🟢

note

If you reattach a detached playbook, any
customizations you have made to the playbook
will be overwritten when the playbook updates
to the current version.

📝🟢

note

Retries are not supported for data collection
tasks that have errors sending emails (indicated
by a server timeout). This is because retries only
work on automation execution failures, not on
email delivery issues.

3.9.5. AI Prompts 

📝🟢

note

In MSSP environments, this option is not
available on the main tenant.

🧭➡️

navi…

On the AI Prompts page (Investigation →
Response → Automation → AI Prompts), you
can view, edit, and manually create prompts

3.9.5.1. AI prompts role-based access control 

3.9.5.2. Use existing prompts  🧭➡️

navi…

Investigation & Response → Automation → AI
Prompts and in the Prompts Library search for
the prompt you want to use

3.9.5.3. Create a prompt  🧭➡️

navi…

Investigation & Response → Automation → AI
Prompts and click + New Prompt

3.9.5.4. Write effective prompts 

📝🟢

note

An AI prompt cannot create new context keys;
it can only refer to existing ones. The only key
an AI prompt can create is the output key.
Therefore, if a variable like ${PostKey} is not
already mapped to existing context data, it
cannot be used in the prompt.

📝🟢

note

The LLM's response is always returned in a text
format, even if you explicitly ask the LLM to
return a structured format such as JSON.
Therefore, the subsequent task in your
playbook will have to parse the response to
extract the relevant data.

3.9.6. Agentic Response (Preview) 

3.9.7. Create an automation rule 

📝🟢

note

You can also define the conditions that trigger
a specific playbook in the playbook editor. For
more information, see Task 2. Configure
playbook settings

📝🟢

note

Quick Actions, by default, run using all
available integration instances that contain the
command. When selecting a Quick Action for
an automation rule, you can instead choose one
specific integration instance to use.

📝🟢

note

Only users with both the Edit Public Playbooks
permission, located under the Playbooks
component in the role configuration, and the
Interact with agents permission, located under
the Cortex Agentic Assistant → Agents
component in the role configuration, can create
automation rules that trigger agents.

📝🟢

note

Only public agents can be selected. If you
create a custom agent, it must be set as a
public agent to use it in an automation rule.

📝🟢

note

To avoid making real world changes in third-
party systems, we recommend expanding the
Test Data field and clicking Edit to manually
change the sample data. When you run the test,
third-party actions are executed. For example,
if the prompt leads the agent to create a Jira
issue as part of the test, an actual Jira issue will
be created.

📝🟢

note

When automation rules trigger agents, there is
a limited number of parallel agentic chats. If
this limit is reached, there may be delays in
execution, but all rules are still applied and
agents are still triggered.
If the agent tries to execute an action that is
marked as sensitive and requires manual
approval, the pending task shows in the
Resolution...

3.9.8. Scripts 

☑️🟣

prer…

To work with scripts, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Scripts: This component (under Investigation &
Response → Automations) must be set to
Enabled first. It is the foundational permission
for all automation; if Scripts are not enabled,
you cannot...

3.9.8.1. Access to scripts 

3.9.8.2. Use existing scripts  🧭➡️

navi…

Investigation & Response → Automation →
Scripts and in the Scripts Library search for the
script you want to use

3.9.8.3. Create a script 

📝🟢

note

You can enable/disable a script in the Settings
without having to duplicate the script.
You can view recently modified or deleted
scripts by clicking the version history for all
scripts .

👀🟠

imp…

If you choose Python, from the Agentic
Assistant you can use the Automation Engineer
agent.

💡🔵

tip

If there is an error, you can copy the error
message from the test result into the Agentic
Assistant prompt and ask the Automation
Engineer agent to correct the error.

3.9.8.4. Use the Automation Engineer agent to
accelerate script development and deployment 

📝🟢

note

The Automation Engineer agent is available for
users with script editing permissions.

📝🟢

note

If you make manual edits in the script and don't
save the changes and you then modify the
script with the Automation Engineer agent, you
are prompted to confirm overwriting the
manual edits.

💡🔵

tip

Use detailed prompts, for example, Get failed
logins from last 24 hours and return as a table.
Clearly define argument names, types, and
default values.
Mention if you expect the output in a specific
format, such as a table, JSON, or plain text.

💡🔵

tip

If there is an error, you can copy the error
message from the test result into the Agentic
Assistant prompt and ask the Automation
Engineer agent to correct the error.

3.9.8.5. Change the Docker image in an
integration or script  📝🟢
note

You can access publicly available Docker
images from the Cortex AgentiX tenant even if
there is no external connection to the Demisto
registry, for example, if due to firewall
constraints, your engine cannot access the
Demisto registry.

3.9.9. Context data 

3.9.9.1. Issue context data 

3.9.9.2. Case context data 

3.9.9.3. Search context data 

3.9.9.4. Add context data to an issue 

3.9.9.5. Add context data to a case  📝🟢

note

If you run the command in the issue War Room,
the data is added to the following places:
The case context data.
The issue context data under the case tab.
If you run the command in the Case War Room,
the data is added to the case context data only.

3.9.9.6. Delete context data from a case 

3.9.9.7. Use context data in a playbook  ⚠️🟡

caut…

Users with Trigger Playbook permissions on a
given issue may still be able to modify the
parent case via commands and scripts, even
without full access to the case.

3.10. Data management 

3.10.1. Dataset management 

🔒🟣

cont…

Managing Roles requires an Account Admin or
Instance Administrator role.

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

Only datasets that are already handled as part
of the GB license are supported for this license.
In addition, the retention configuration is only
available in Cortex AgentiX, as opposed to the
public APIs.

📝🟢

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.
Datasets include dataset permission
enforcements in the Cortex Query
Language(XQL), Query Center, and XQL
Widgets. For example, to view or access any of
the endpoints and host_inventory datasets, you
need role-based access control (RBAC)...

📝🟢

note

You'll only be able to save the dataset view if
the query contains no errors; otherwise, the
Save button is disabled.

👀🟠

imp…

Cortex AgentiX enforces retention on all log-
type datasets, excluding metrics and users.

👀🟠

imp…

This column is updated once a day. Therefore,
if the dataset was created or updated by the
target or lookup flows, it's possible that the
Last Updated value is a day behind when the
queries or reports were run as it was before this
column was updated.

3.10.1.1. What are datasets? 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

CEF and LEEF logs are often wrapped in a
Syslog envelope (RFC 3164/5424). Cortex
AgentiX automatically detects and strips the
Syslog header to extract the reporting device IP
and hostname before parsing the inner CEEF or
LEEF payload.

📝🟢

note

For more information on tasks specific to
lookup datasets, see Lookup datasets.

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query and can appear with a slight delay.

👀🟠

imp…

By default, forensic datasets are not included
in XQL query results, unless the dataset query
is explicitly defined to use a forensic dataset.

3.10.1.2. Lookup datasets 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

A lookup dataset can only be deleted if there
are no other dependencies. For example, if a
Correlation Rule is based on a lookup dataset,
you wouldn't be able to delete the lookup
dataset until you removed the dataset from the
XQL query of the Correlation Rule.

👀🟠

imp…

When you create or add data to a lookup
dataset using the target stage, the _time field
won't be included by default unless you
explicitly add it with the fields stage.

3.10.1.2.1. Import a lookup dataset 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

☑️🟣

prer…

When uploading a CSV, TSV, or JSON file,
ensure that the file meets the following
requirements:
The maximum size for the total data to be
imported into a lookup dataset is 30 MB from
the Dataset Management page. Otherwise, the
limit is 50 MB using Cortex Query Language
(XQL) or APIs.
Field names can contain characters from
different languages, special...

👀🟠

imp…

The name of a dataset created from a TSV file
must always include the extension. For example,
if the original file name is mrkdptusrsnov23.tsv,
you can save the dataset with the name
marketing_dept_users_Nov_2023.tsv.

3.10.1.2.2. Download JSON file of lookup dataset 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

🧭➡️

navi…

Open the Settings → Configurations → Data
Management → Dataset Management page

3.10.1.2.3. Set time to live for lookup datasets  ☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

3.10.1.3. Monitor datasets and dataset views

activity 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

3.10.2. Manage compute units  3.10.2.1. Compute units usage 

📝🟢

note

To enable the add-on, select Settings → Cortex
AgentiX LicenseAddons tile, and select the
Compute Unit tile and Enable.

🧭➡️

navi…

Settings → Configurations → Data
Management → Compute Unit Usage

3.10.3. Manage Event Forwarding 

☑️🟣

prer…

Event Forwarding requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Dataset
Management.

📝🟢

note

After exporting logs, you can download them
from GCP for up to 14 days. The Pub/Sub
subscription messages are available for 7 days.

📝🟢

note

Requires the Endpoints Event Forwarding add-
on.

📝🟢

note

Requires the GB Event Forwarding add-on

👀🟠

imp…

Since data is aggregated and compressed, it
may take up to two hours until the data is
available in the forwarding bucket.

🧭➡️

navi…

Under Settings → Configurations → Data
Management → Event Forwarding, activate one
or more of the following: Enable GB Event
Forwarding

Under Settings → Configurations → Data
Management → Event Forwarding, activate one
or more of the following

3.11. Dashboards and Reports 

3.11.1. About dashboards 

📝🟢

note

Dashboards that are based on custom
infrastructure cannot be exported.
When importing a dashboard that already
exists in the system, the imported dashboard
overwrites the existing dashboard. To avoid
overwriting the existing dashboard, duplicate
and rename the existing dashboard before
importing the new dashboard.

👀🟠

imp…

The ability to create, edit, or share custom
dashboards is governed by access
management. If certain options are unavailable,
contact your administrator. For more
information, see Manage access to custom
dashboards.

3.11.1.1. Command Center dashboards 

📝🟢

note

Access to these dashboards requires RBAC
permissions under Dashboards & Reports. The
Dashboards component must be Enabled in
your role and requires View permissions for
Command Center Dashboards. If certain
options are unavailable, contact your
administrator. For more information, Manage
access to custom dashboards.
The dashboards are available in dark...

3.11.1.1.1. AgentiX Command Center  📝🟢

note

Access to the dashboard requires RBAC View
permissions for Dashboards & Reports and
Command Center Dashboards. Non-admin
users who have these permissions can view
shared system and custom agents as well as
their own agents, listed under Other Agents.
They cannot see other users' agents.
The dashboards are available in dark mode
only. They are not...

3.11.1.2. Predefined dashboards  🔒🟣

cont…

Requires the Cortex XSIAM Premium license or
any other XSIAM license with the Threat Intel
Management (TIM) add-on to use this feature.

3.11.2. Reports 

3.11.2.1. Report templates 

3.11.2.2. Run or schedule reports 

📝🟢

note

Report templates that are based on custom
infrastructure cannot be exported.
If you import a report template that already
exists in the system, the imported template will
overwrite the existing template. If you do not
want to overwrite the existing template,
duplicate and rename the existing template
before importing the new template.
You can also...

📝🟢

note

The report name and description will be
displayed in the report header and are not
editable during customization.

📝🟢

note

For agent-related widgets, you can apply an
endpoint scope to refine the displayed data to
only show results from specific endpoint
groups.

Select the menu on the top right corner of the
widget, select Groups, and select one or more
endpoint groups.
For agent-related widgets, you can apply an
endpoint scope to refine the displayed data to
only show...

📝🟢

note

To send reports to Slack, Slack must be
configured as an external application in Cortex
AgentiX. For more information, see Integrate
Slack for outbound notifications

📝🟢

note

When ownership is transferred, any existing
report schedules associated with the report
template are automatically removed. The new
Owner must manually redefine the schedule to
resume automated report generation.

🧭➡️

navi…

Under Settings → Configurations → General →
Notifications, click + Add Forwarding
Configuration

3.11.3. Build custom dashboards and reports 

3.11.3.1. Build a custom dashboard  📝🟢

note

For case-related widgets, you can refine the
displayed data to only show results from cases
that match a case starring configuration. A
purple star indicates that the widget is
displaying only starred cases. For more
information, see Case starring.
Custom graphs and charts generated by the
Cortex Agentic Assistant and saved to the
Widget Library can...

3.11.3.2. Manage your Widget Library 

📝🟢

note

Unlike dashboards, individual widgets cannot
be shared with other users directly from the
Widget Library. Yet, if you add a Restricted
widget to a dashboard and then share that
dashboard, other users who have access to the
dashboard will also be able to see that specific
widget.

📝🟢

note

Any dashboards or reports that include the
widget are affected by the changes.

3.11.4. Fine-tune dashboards and reports 

3.11.4.1. Create a custom widget using a script 

📝🟢

note

If you have added arguments to the script,
these appear when creating a widget.

📝🟢

note

Available options are Pie, Column, Line, and
Single Value.
To display the result of the script as a time
duration, choose the graph type Single Value
and enable Show as Time. You can then select
the Time Unit (millisecond, second, minute, or
hour) and the Display format.

3.11.4.1.1. Script-based widget examples 

📝🟢

note

Add the widget tag in the script settings to
make the script available for use in script-
based widgets. For more information, see
Create a script.

📝🟢

note

If your script returns a time duration, configure
the widget with the graph type Single Value
and enable Show as Time..

3.11.4.2. Create a text widget  📝🟢

note

From the Dashboard Manager, you can right-
click an existing dashboard and choose Save as
report template. You can also add text-based
widgets directly to new or existing report
templates, as described below.

3.11.4.3. Create custom XQL widgets 

📝🟢

note

Cortex Query Language (XQL) queries
generated from the Widget Library do not
appear in the Query Center. The results are
used only for creating the custom widget.

📝🟢

note

Whenever the time period is changed in the
query window, the config timeframe is
automatically set to the time period defined,
but this won't be visible as part of the query.
Only if you manually type in the config
timeframe will this be seen in the query.
These time picker options are available in XQL
queries when using the Query Builder, XQL
Widgets,...

📝🟢

note

To display the result of as a time duration,
choose the graph type Single Value and enable
Show as Time. You can then select the Time
Unit (millisecond, second, minute, or hour) and
the Display format.

💡🔵

tip

You can create a generic dashboard for
multiple views of the same dataset by using an
asterisk (*) when defining the dataset in the
XQL widget as dataset = <dataset_name>*. The
placement of the asterisk in the dataset name
ensures that any view containing this prefix text
is displayed in the results.
The dataset in a query is defined as:
If there are...

3.11.4.3.1. Configure filters and inputs for
custom XQL widgets 

☑️🟣

prer…

Fixed filters are based on parameters that are
defined in custom XQL widgets. Before you can
configure fixed filters, take the following steps:
Create custom XQL widgets with parameters.
For more information, see Create custom XQL
widgets.
Add the widgets to a custom dashboard. For
more information, see Build a custom
dashboard.

📝🟢

note

This option only appears if the dashboard
contains custom XQL widgets with defined
parameters.

📝🟢

note

You can define up to four parameter filters on a
single dashboard or report.

📝🟢

note

If you specify more than one field, only the first
field value is used.

💡🔵

tip

After the initial setup, when you access your
dashboard the filters and inputs might need
further refinement. You can make changes to
the configured parameters in the XQL widgets,
and update the Filters & Inputs on your
dashboard until you are satisfied with the
results.

3.11.4.4. Configure dashboard drilldowns 

☑️🟣

prer…

To configure drilldowns your dashboard must
contain custom XQL widgets. In addition, if you
want to configure in-dashboard drilldowns your
custom XQL widget must contain one or more
parameters. For more information about
configuring parameters in custom XQL widgets,
see Create custom XQL widgets.

📝🟢

note

If the selected parameter is configured in other
XQL widgets on the dashboard, these widgets
are also affected by the drilldown.

3.11.4.4.1. Variables in drilldowns 

4. Investigation and Response 

4.1. Overview of cases 

☑️🟣

prer…

To work with cases, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Playbooks: This component (under Investigation
& Response → Automations) must be set to
Enabled first. Role-level permissions determine
your ability to create new playbooks or edit
those marked as...

4.1.1. What are cases? 

4.1.2. Resolving cases with AI 

4.1.3. Case lifecycle 

4.1.4. Case thresholds 

4.1.5. Case scope and impact 

4.1.6. Case and issue domains 

4.2. Case concepts 

4.2.1. Issues 

4.2.2. Case grouping 

4.2.3. Case scoring 

🧭➡️ 4.2.4. Case starring 

navi…

You can manage all starring configurations
under Case & Issues → Case Configuration →
Starred Issues

4.2.5. SLAs and tracking 

4.3. Analyze and resolve cases 

📝🟢 4.3.1. Review all cases 

note

Administrators can enable or disable legacy
mode. Go to Configurations → General →
Server Settings → Case display modes.

4.3.2. Start case analysis 

📝🟢

note

This section describes the product in default
mode and using the Split view. If you are using
legacy mode, see Detailed View.

4.3.2.1. Agentic Assistant- Case Investigation

agent 

4.3.3. Establish case context 

4.3.3.1. AI-generated case summaries 

📝🟢

note

The AI-generated title and description is a
calculated value that is regenerated each time
you open a case.

This value is not a saved static description,
therefore it is not reflected in the saved case
names in the list of cases in the Split view , or in
the Case Name and Case Description columns
in the Table view.

🧭➡️

navi…

Configurations → General → Server Settings →
AI Configuration and enable the following
settings

4.3.3.2. Assess case severity and score 

📝🟢 4.3.3.3. Update case attributes 

note

The defined values are shown in the Case Name
and Case Description columns in the Table view,
and saved to the cases dataset. The case title is
also shown the list of cases in the Split view.
These values do not replace the AI-generated
case title and summary. If the Summarize with
AI toggle is enabled, AI-generated case
summaries are automatically...

4.3.4. Analyze case details 

📝🟢

note

If you prefer a tabular or legacy layout, switch
the case card to the Detailed view.
This view preserves the legacy tab based
format and custom layouts, ensuring full
backward compatibility. You can switch
between the new case experience and the
legacy view based on personal workflow
preferences. For more information, see Detailed
View.

📝🟢 4.3.4.1. Grouping graph 

note

Similar issues are displayed as individual entities
rather than in a parent/child hierarchy.

4.3.4.2. Evidence 

4.3.4.3. Issue feed 

📝🟢 4.3.4.4. Associated assets and artifacts 

note

If you do not have permissions to access an
asset of a case (which is shown as grayed out
and locked), check your scoping permissions in
Manage Users or Manage User Groups.

📝🟢 4.3.4.5. MITRE ATT&CK tactics and techniques 

note

This component is available for cases
associated with the Security domain or custom
domains.

📝🟢 4.3.4.6. Case timeline 

note

If you unmark a record as evidence, the
evidence flag is removed but the underlying
timeline record and its attachments are not
deleted.

4.3.4.7. Detailed View 

📝🟢

note

When an issue is resolved, it remains linked to a
case. Once all of the issues in a case are
resolved, the case is automatically closed.

📝🟢

note

Cloud-related issues are displayed in the User
Name grouping.
Prisma Cloud Compute issues are displayed in
the Host Name grouping.

4.3.4.8. Issue card 

🔒🟣

cont…

Requires the TIM add-on.

📝🟢

note

This section is context-specific and shows data
according to the issue context.

4.3.5. Resolve the case 

📝🟢 4.3.5.1. Resolution Center 

note

Tasks that are already In Progress but require
your input will appear in both the Pending and
In Progress tabs.

4.3.5.2. Collaborative notes and comments 

4.3.5.3. Resolve a case 

📝🟢

note

If a case is resolved with the status Resolved -
Auto Resolved, Cortex AgentiX can reopen the
case for up-to six hours if a new issue is
triggered that matches the case. The six-hour
period is defined by the timestamp of the last
issue that was grouped into the case. After the
six-hour period, any new issues are linked to a
new case for a new...

4.3.5.3.1. Resolution reasons for cases and issues 

📝🟢

note

The displayed resolution reasons are domain-
specific. You can see the resolution reasons
that are defined for a domain under
Configurations → Object Setup → Cases →
Domains.

📝🟢

note

Cases and issues resolved as True Positive and
False Positive help Cortex AgentiX to identify
real threats in your environment by comparing
future cases and associated issues to the
resolved cases. Therefore, the handling and
scoring of future cases is affected by these
resolutions.

📝🟢

note

Cases and issues resolved as True Positive and
False Positive help Cortex AgentiX to identify
real threats in your environment by comparing
future cases and associated issues to the
resolved cases. Therefore, the handling and
scoring of future cases is affected by these
resolutions.

4.3.6. Create a case 

📝🟢

note

To create a case manually, you must have
View/Edit permission for Cases and Issues
selected under Settings → Configurations →
Access Management → Roles → Components
→ Cases & Issues.

📝🟢

note

You can assign a case to a single domain only,
and you cannot change the assigned domain.
For more information, see Case and issue
domains.

📝🟢

note

This option is only relevant for certain domains.

📝🟢

note

This option is only relevant for certain domains.

📝🟢

note

You can't attach files to manually created
cases.

💡🔵

tip

The issues that you link to a case can be linked
to multiple cases, and the issue domains do not
need to match the case domain.

4.4. Investigate issues 

☑️🟣

prer…

To work with issues, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Playbooks: This component (under Investigation
& Response → Automations) must be set to
Enabled first. Role-level permissions determine
your ability to create new playbooks or edit
those marked as...

📝🟢 4.4.1. Overview of the Issues page 

note

Every 12 hours, the system enforces a cleanup
policy to remove the oldest issues once the
maximum limit is exceeded. The default issue
retention period in Cortex AgentiX is 186 days.

4.4.2. Issue card 

🔒🟣

cont…

Requires the TIM add-on.

📝🟢

note

This section is context-specific and shows data
according to the issue context.

4.4.3. Resolution actions 

4.4.4. Link or unlink issues from a case 

4.4.5. Run an automation on an issue 

📝🟢

note

In addition to automation, some playbooks
contain manual tasks that prompt the analyst
for input. This enables you to enhance an
automation workflow with analyst input.

📝🟢

note

You can also manually select a playbook to run
from the Issue Work Plan tab.

4.4.6. Use the War Room in an investigation 

📝🟢

note

The case War Room is usually used for
communication capabilities, but unlike the issue
War Room, it does not include playbook
specific entries. The case War Room enables
you to investigate an entire case, not just an
issue.

📝🟢

note

Cortex AgentiX does not index notes and chats.

💡🔵

tip

In the Playground, you can clear the context
data, if needed, which deletes everything in the
Playground context data, but does not affect
the actual issue or case. To clear the context,
run !DeleteContext all=yes' from the CLI or
click Clear Context Data while viewing the
context data.

💡🔵

tip

You can use the up/down arrow buttons in the
CLI to do a reverse history search for previous
commands with the same prefix.

🧭➡️

navi…

Investigation & Response → Automation →
Playground

👀🟠 4.4.7. Use the Work Plan in an investigation 

imp…

A playbook will not continue its execution path
if a prior task has failed; you must resolve the
failed task before subsequent tasks can run.

4.4.8. Manage synced tickets in issues 

☑️🟣

prer…

An integration must be configured before you
can sync issues. For more information, see Set
up an integration for mirroring issues.Issue
syncing with an external integration

📝🟢

note

Multiple tickets can be linked to an issue with
outbound syncing. Issues with inbound syncing
can be linked to a single ticket only.

📝🟢

note

Using issue fields as variables is not currently
supported.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

📝🟢

note

If you change the selected inbound sync
profile, the original sync profile is immediately
overwritten.

⛔🔴

war…

If you leave this field blank, all configured
instances will be used.

💡🔵

tip

You can find a sync profile ID under Settings →
Configurations → Object Setup → Issues →
Sync Profiles. By default, the ID field is not
displayed in the table. Click the three-dot menu
and add it to the layout.

4.4.9. Issue investigation actions 

4.4.9.1. Copy issues 

4.4.9.2. Investigate contributing events 

4.4.9.3. Export issue details to a file 

📝🟢 4.4.9.4. Exclude an issue 

note

Cortex AgentiX supports exclusion of up to
100,000 issues.

4.4.9.5. Query case and issue data 

📝🟢 4.4.9.6. Run indicator extraction in the CLI 

note

Triggering enrichment on a substantial amount
of indicators can take time (since it's activating
all enrichment integrations per indicator) and
can result in performance degradation.

4.4.9.7. Update issue fields 

📝🟢

note

You must create a custom status before you
can select it.

💡🔵

tip

You can create custom issue statuses and
resolution reasons, and use the setIssueStatus
command to set these custom statuses for
issues.
For example, when a user starts investigating an
issue, the issue status is automatically changed
from New to Under Investigation. In some cases,
it is useful to create an interim status, such as
Triage. After you...

4.4.9.8. Use issue timer field commands

manually in the CLI 

📝🟢

note

When defining the values for the slaField use
the machine name for the field, which is
lowercase and without spaces. You can check
the machine name by editing the issue field.

📝🟢

note

Timer fields are not started automatically when
an issue is created unless run in a playbook.

📝🟢

note

After a timer field is stopped, before you can
start the timer again you must reset the timer
using the resetTimer command.
Timers are automatically stopped when an
issue is closed.

📝🟢

note

When running commands in the CLI, you can
specify the alertID to change the timer for a
different issue.

📝🟢 4.4.9.9. Close an issue 

note

When an issue is resolved it remains linked to a
case. Once all of the issues in a case are
resolved, the case is automatically closed.

4.5. Investigate artifacts and assets 

📝🟢 4.5.1. Investigate an IP address 

note

Requires a license key. Select Settings →
Configurations → Integrations → Threat
Intelligence.

📝🟢 4.5.2. Investigate a file and process hash 

note

Requires a license key. Go to Settings →
Configurations → Integrations → Threat
Intelligence.

5. Agentic Assistant chat 

📝🟢 5.1. Get started with Agentic Assistant chat 

note

The Cortex Agentic Assistant is currently
available in limited regions. For more
information see Cortex Agentic Assistant If your
tenant is not within one of those regions, you
have access to the Cortex Assistant.Cortex
Agentic AssistantCortex Assistant
To enable the Cortex Agentic Assistant, go to
Settings → Configurations → General →
Server...

📝🟢 5.2. Choose an Agentic Assistant agent 

note

Only public agents are supported via Slack.

5.3. Chat with an Agentic Assistant agent 

📝🟢

note

Case data is only loaded when you send your
first message. Opening the chat interface
without sending a prompt does not provide the
agent with the case context.

📝🟢

note

An agent's proposed plans and results may
contain inaccuracies or errors. Always review
the results carefully to ensure you fully
understand the proposed action before
proceeding.

💡🔵

tip

You can quickly jump to different product
pages within Cortex AgentiX by typing / in the
prompt area. This shortcut is a built-in
navigation feature that is available even if the
Cortex Agentic Assistant is disabled.

5.4. Chat with the Agentic Assistant from Slack 

☑️🟣

prer…

Before you can interact with the Cortex
Agentic Assistant in Slack, the Slack v3
integration must be set up correctly. For more
information, see the Slack v3 integration
documentation.
To perform actions in Slack, your Slack email
must match your Cortex AgentiX user email.
This ensures the system can strictly follow your
assigned permissions (RBAC). If...

📝🟢

note

Only public agents are supported via Slack.

📝🟢

note

If a team member in the channel sees the
summary and attempts to ask the agent, "Give
me more details on the first SLA issue," the
team member receives an access denied
message because the active session is only
available to you, the initiator.

5.5. Create and run XQL queries with Agentic
Assistant chat 

📝🟢

note

The TextToXQL action is a hidden system action
and does not appear in the Agents Hub.

📝🟢

note

Running XQL queries manually through an
agent does not consume compute units. This
includes scenarios where you prompt the agent
to create and execute a multi-step plan.

5.6. Use natural language to query and visualize
your data 📝🟢 
note

Once the graph or chart is saved the Widget
Library, the link to the chat artifact is severed,
and the agent does not track subsequent
changes made to the widget or the dashboard.

5.7. Manage chat history 

6. Threat Intel Management 

6.1. Get started with Threat Intel Management 

6.1.1. What is Threat Intel Management? 

6.1.2. Threat Intel Management use cases 

6.1.3. Roles and responsibilities in Threat Intel
Management 

📝🟢 6.1.4. Indicator concepts 

note

You can store up to 100,000,000 indicators.

6.1.5. Indicator lifecycle 

📝🟢

note

To run enhancement scripts and reputation
commands, you must configure a relevant
enrichment integration, such as VirusTotal,
IPinfo v2, etc.

💡🔵

tip

We recommend defining your policy for
handling expired indicators.

6.2. Indicator configuration 

6.2.1. Customize indicator fields and types 

6.2.1.1. Create an indicator type 

📝🟢

note

If you miss mapping any field, you can create
additional new indicator fields and either relate
them to all indicator types, or relate them only
to the new indicator type (recommended).

🧭➡️

navi…

Settings → Configurations → Object Setup →
Indicators → Types

6.2.1.1.1. Indicator type profile 

📝🟢

note

Indicator cache expiration rules only apply to
standard enrichment (for example, running the
enrichIndicators command). If you run a
reputation command, such as !ip, the
commands executes even if the cache has not
expired.

💡🔵

tip

To control which integrations run when
explicitly executing a command (and not during
enrichment), add the using-brand argument to
specify only the integrations you want to use.
For example:
!url url="https://www.google.com" using-
brand="AutoFocus V2,VirusTotal"

📝🟢 6.2.1.1.2. File indicators 

note

A new File indicator only affects new indicators
ingested to the Cortex AgentiX platform.
Indicators that were already in Cortex AgentiX
continue to appear as their respective hash-
related indicators.

6.2.1.1.3. Formatting scripts 

📝🟢

note

Formatting scripts must have the indicator-
format tag to appear in the list.

📝🟢

note

Formatting scripts for out-of-the-box indicator
types are system-level, which means that the
formatting scripts for these indicator types are
not configurable. To create a formatting script
for an out-of-the-box indicator type, you need
to disable the existing indicator type and create
a new (custom) indicator type. If you
configured a formatting...

6.2.1.1.4. Enhancement scripts 

📝🟢

note

Enhancement scripts are different from
reputation commands. A reputation command
runs every integration that has that command
within it, to enrich the indicator. The reputation
command ip , for example, runs every IP
integration command in your enabled
integrations, to collect data from multiple
sources. An enhancement script is manually run
after the...

📝🟢

note

Enhancement scripts must have the
enhancement tag applied to appear in the list.

📝🟢 6.2.1.1.5. Reputation scripts 

note

Reputation scripts must have the reputation tag
applied to appear in the list.

📝🟢 6.2.1.1.6. Reputation commands 

note

Running a reputation command directly (such
as !ip) might not apply the result to an
indicator, nor does it use the enrichment cache.
To ensure an indicator is enriched, and to take
advantage of caching, use the enrichIndicators
command or the Enrich button in the UI. This
runs the appropriate reputation
command/script based on the indicator type...

📝🟢 6.2.1.1.7. Map custom indicator fields 

note

Some integrations have indicator mappers and
classifiers, such as AWS. If you want to use an
integration mapper or classifier, see Indicator
classification and mapping.

6.2.1.2. Create an indicator field 

📝🟢

note

Cortex AgentiX IOC fields are based on the
STIX 2.1 specifications. For more information,
see Indicator field structure.

📝🟢

note

The following HTML tags are not permitted:
blockquote, del, dd, div, dl, dt, fieldset, form, h1,
h2, h3, h4, h5, h6, hr, iframe, ins, li, math,
noscript, ol, pre, p, script, style, table, ul,
address, article, aside, canvas, details, dialog,
figcaption, figure, footer, header, hgroup, main,
nav, output, progress, section, video.
The following CSS tags...

🧭➡️

navi…

Settings → Configurations → Object Setup →
Indicators → Fields → New Field

6.2.1.2.1. Indicator field structure 

📝🟢 6.2.1.2.2. Indicator field trigger scripts 

note

Indicator field trigger scripts must have the
field-change-triggered-indicator tag to appear
in the list.

6.2.2. Indicator classification and mapping 

📝🟢

note

When creating a new indicator type, you
classify and map the indicator fields in the
indicator type settings. For more details, see
Map custom indicator fields.

📝🟢

note

You can optionally skip importing data. Click
the pencil on the right of each indicator type
on the right pane to enter the value manually.

📝🟢

note

In the Classification & Mapping page, the
mapping does not indicate for which indicator
types they are configured. When creating a
mapper, it is best practice to add to the
mapper name and the indicator type the
mapper is for. For example, Mail Listener -
Phishing.
When mapping a list, we recommend you map
to a multi-select field. Short text fields do...

6.2.3. Indicator extraction 

📝🟢

note

By default, system-wide automatic indicator
extraction and enrichment is disabled. However,
if you migrated from Cortex AgentiX 2.x to
Cortex AgentiX 3.x, system-wide automatic
indicator extraction and enrichment is enabled.
If you have a Threat Intel Management (TIM)
Add-on, you can enable or disable automatic
indicator extraction system-wide. Go to...

📝🟢

note

Reputation commands, such as !ip and !domain,
can only be used after you configure and
enable a reputation integration instance, such
as VirusTotal and Whois.

📝🟢

note

The inline configuration may delay playbook
execution.

📝🟢

note

While indicator creation is asynchronous,
indicator extraction and enrichment are run
synchronously. Data is placed into the issue
context and is available via the context for
subsequent tasks.
All indicators are automatically extracted and
enriched before a playbook is run. For an on-
field change, extraction occurs before the next
playbook tasks run.

📝🟢

note

When using out of band, the extracted
indicators do not appear in the context. If you
want the extracted indicators to appear select
inline.

6.2.3.1. Set the indicator extraction mode for a

playbook task 

6.2.3.2. Disable indicator extraction for scripts
or integrations 

6.2.4. Configure indicator expiration 

📝🟢

note

You need to run CLI commands in the Case or
Issue War Room.

📝🟢

note

If a feed's expiration method is set to When
removed from the feed, indicators that are
removed from the feed immediately expire.
Note that if the feed is disabled, its expiration
method reverts to that of the indicator type
(time-based).
Time-based expiration is set according to feed
reliability. If the same indicator appears on
multiple feeds, the...

6.2.5. Configure Threat Intelligence feed
integrations 

📝🟢

note

Some third-party services (such as Whois and
MITRE ATT&CK) enforce strict rate limits based
on the source IP address. If you encounter
quota management issues, we recommend
running them on an engine. This routes traffic
through your private network, ensuring the
external service sees a unique, dedicated IP
address assigned to your organization.

📝🟢

note

Some feeds only provide information about
new indicators and do not specify when
indicators are removed. Indicators from these
feeds cannot be automatically expired on
removal.
If a feed's expiration method is set to When
removed from the feed, indicators that are
removed from the feed immediately expire.
Note that if the feed is disabled, its...

6.2.6. Exclude indicators from enrichment 

6.2.7. Generate issues from indicator rules 

📝🟢

note

Indicators should be present in the Threat
Intelligence database (Threat Management →
Threat Intelligence → Indicators) before
creating detection rules.

📝🟢

note

You can't change the Detectable = True and
Status = Active filters, which comply with the
requirements of the supported indicator type
for detection.

📝🟢

note

The Issue source is Threat Intelligence.

📝🟢

note

If an indicator matches multiple indicator rules,
the highest severity rule is used. If all have the
same severity, the rules are used by the first
created.

⚠️🟡 6.2.8. Export indicators 

caut…

You should not run a query on a field that you
might change in the playbook flow. For
example, you shouldn't have a playbook with
query Verdict:Malicious and then change the
indicator verdict as a part of the playbook.

6.2.9. Manage external dynamic lists 

📝🟢

note

Configuring custom certificates or private API
Keys in the EDL integration instance is
supported only on engines, not on the Cortex
AgentiX tenant.
For EDL integrations on the tenant, you must
set a username and password. For long-running
integrations running on an engine, we strongly
recommend setting a username and password,
but it is not required....

📝🟢

note

If an indicator query returns more indicators
than the EDL list size, the list is populated with
the most recent n indicators sorted by their last
seen timestamp, where n is the maximum size of
the EDL.

👀🟠

imp…

The EDL URL must always be prefixed by ext-.

💡🔵

tip

You can also use the IP list and URL lists as part
of a URL Filtering policy, or the domain list as
part of a custom Anti-Spyware profile.

6.3. Indicator management 

6.4. Indicator investigation 

6.4.1. Indicator verdict 

6.4.2. Extract and enrich an indicator 

📝🟢

note

By default, system-wide automatic indicator
extraction and enrichment is disabled. However,
if you migrated from Cortex AgentiX 2.x to
Cortex AgentiX 3.x, system-wide automatic
indicator extraction and enrichment is enabled.
If you have a Threat Intel Management (TIM)
Add-on, you can enable or disable automatic
indicator extraction system-wide. Go to...

📝🟢

note

Triggering enrichment on a substantial number
of indicators can take time (because it's
activating all enrichment integrations per
indicator) and can result in performance
degradation.

📝🟢

note

Reputation commands, such as !ip, !domain can
only be used when you configure and enable a
reputation integration instance, such as
VirusTotal and WHOIS.

📝🟢 6.4.3. Expire an indicator 

note

You need to run these commands in the Case or
Issue War Room.

📝🟢 6.4.4. Manage indicator relationships 

note

You can also add an indicator relationship from
the Quick View when selecting an indicator
from an issue.

⚠️🟡 6.4.5. Delete and exclude indicators 

caut…

Ensure you are using the correct syntax when
defining the values for your exclusion lists.

7. Cortex AgentiX XQL 

7.1. Get started with XQL 

7.1.1. XQL language features 

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion commands and
definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc key.

🧭➡️

navi…

You submit XQL queries to Cortex AgentiX
using the Investigation & Response → Search →
Query Builder user interface

7.1.2.1. Adding comments in queries  7.1.2. XQL Language Structure 

📝🟢 7.1.3. Supported operators 

note

In some cases, using an IN or NOT IN operator
combined with a dataset and filterfilter stage
can be a better alternative to using a join
stage.join

7.1.4. Datasets and presets 

📝🟢

note

INFO issues are not included in this dataset.
This dataset includes issues from the Security
and Health domains. For more information, see
Overview of the Issues page.

📝🟢

note

The fields contained in this dataset are a subset
of the fields in the xdr_data dataset.

📝🟢

note

If the vendor and product are not specified in
the Winlogbeat profile's configuration file,
Cortex AgentiX creates a default dataset
called microsoft_windows_raw.

📝🟢

note

To ensure GlobalProtect access authentication
logs are sent to Cortex AgentiX, verify that
your PANW firewall's Log Settings for
GlobalProtect has the Cortex Data Lake
checkbox selected.

📝🟢

note

The fields contained in this dataset are a subset
of the fields in the xdr_data dataset.

📝🟢

note

Prisma Access firewalls do not send
configuration logs to the Structured Log
Storage (SLS).

📝🟢

note

You can query in XQL for this data and build
widgets based on the xdr_data dataset or using
the preset device_control.
To view in an XQL query these events, the
Device Configuration of the endpoint profile
must be set to Block. Otherwise, the USB events
are not captured. The events are also captured
when a group of device types are blocked on
the...

📝🟢

note

The fields contained in this dataset are a subset
of the fields in the xdr_data dataset.

👀🟠

imp…

This dataset will not be updated after June
2024. Use the health_alerts dataset instead.

🧭➡️ 7.1.5. About examples 

navi…

Investigation & Response → Search → Query
Builder → XQL → Query Library

👀🟠 7.1.6. JSON functions 

imp…

JSON field names are case sensitive, so the key
to field pairing must be identical in an XQL
query for results to be found. For example, if a
field value is "TIMESTAMP" and your query is
defined to look for "timestamp", no results will
be found.

7.1.7. How to filter for empty values in the results
table 

7.1.8. Understanding string manipulation in XQL 

7.2. Build XQL queries 

7.2.1. About the Query Builder 

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query, and can appear with a slight delay.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion commands and
definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc key.

💡🔵

tip

If you prefer to use the Query Builder in Legacy
mode, switch the toggle in the header. In
Legacy mode, the Query Builder searches
predefined datasets only. To search the full
XDM Data Model, switch to New mode or
select XQL Search.

7.2.2. How to build XQL queries 

📝🟢

note

Users with different dataset permissions can
receive different results for the same XQL query.
An administrator or a user with a predefined
user role can create and view queries built with
an unknown dataset that currently does not
exist in Cortex AgentiX. All other users can only
create and view queries built with an existing
dataset.
When you have...

📝🟢

note

You can build a query that investigates data in
both a cold dataset and a hot dataset in the
same query. In addition, as the hot storage
dataset format is the default option and
represents the fully searchable storage, this
format is used throughout this guide for
investigation and threat hunting. For more
information on hot and cold storage, see...

👀🟠

imp…

Forensic datasets are not inlcuded by default in
XQL query results, unless the dataset query is
explicitly defined to use a forensic dataset.

💡🔵 7.2.2.1. Get started with XQL queries 

tip

If you are new to creating queries, you can also
try our simple search templates, which can help
you get started in understanding how queries
work. See Query Builder templates.Query
Builder templates

7.2.2.2. Useful XQL user interface features 

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query and can appear with a slight delay.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

7.2.2.3. XQL Query best practices 

7.2.2.4. Expected results when querying fields 

7.2.2.5. Create XQL query 

📝🟢

note

Whenever the time period is changed in the
query window, the config timeframe is
automatically set to the time period defined for
the entire query, including queries that are part
of the join stage. Yet, this won't be visible as
part of the query. Only if you manually type in
the config timeframe will this be seen in the
query.

These time picker...

📝🟢

note

While datamodel dataset=* is supported in the
query, we recommend that you specify specific
datasets for quicker and more efficient results.

📝🟢

note

Whenever the time period is changed in the
query window, the config timeframe is
automatically set to the time period defined for
the entire query, including queries that are part
of the join stage. Yet, this won't be visible as
part of the query. Only if you manually type in
the config timeframe will this be seen in the
query.

These time picker...

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

💡🔵

tip

While the query is running, you can navigate
away from the page. A notification is sent when
the query has finished. You can also Cancel the
query or run a new query, where you have the
option to Run only new query (cancel previous)
or Run both queries.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

💡🔵

tip

While the query is running, you can navigate
away from the page. A notification is sent when
the query has finished. You can also Cancel the
query or run a new query, where you have the
option to Run only new query (cancel previous)
or Run both queries.

7.2.2.6. Review XQL query results 

📝🟢

note

It's also possible to graph the results displayed.
For more information, see Graph query results.

📝🟢

note

Real time query results are available only in the
Query Builder and in free text query.
Real time results are displayed only for queries
run on hot datasets.
The Sort option is available only after all the
data is retrieved.
When you formulate complex queries, the
results will be displayed when the query has
finished running completely, and not in real...

📝🟢

note

Results are received incrementally for the first
100K records, or up to 100MB worth of records,
whichever comes first. After that, the next
update is when the query has finished running
completely.

📝🟢

note

In order for Cortex AgentiX to provide a
histogram for a field, the field must not contain
an array or a JSON object.

👀🟠 7.2.2.7. Translate to XQL 

imp…

This feature is still in a Beta state and you will
find that not all Splunk queries can be
converted to XQL. This feature will be improved
upon in the upcoming releases to support
greater Splunk query translations to XQL.

7.2.2.8. Graph query results 

📝🟢

note

To display the result of as a time duration,
choose the graph type Single Value and enable
Show as Time. You can then select the Time
Unit (millisecond, second, minute, or hour) and
the Display format.

💡🔵

tip

Alternatively, you can use the Cortex Agentic
Assistant to generate custom graphs and
charts using natural language prompts. By
simply prompting the agent, it will build and
execute the query, returning the visual
representation. For more information, see Use
natural language to query and visualize your
data.

7.2.3. Query Builder templates 

💡🔵

tip

To get started with queries, you can run an
empty template query with no values specified.
The query results will include all of the fields in
the template specific fieldset. Based on the
query results, you can run subsequent queries
to narrow down your search.

7.2.3.1. Get started with Query Builder templates 

7.2.3.2. Considerations for using Query Builder
templates 

7.2.3.3. Create a query from a template 

📝🟢

note

Field names and aliases are listed without their
prefix, for example
xdm.SOURCE.USER.USERNAME is listed as
SOURCE.USER.USERNAME and XDM_ALIAS.ipv4
is listed as ipv4.
Fields that are already included in the query
template are shown as grayed out.
In the Identity and Network templates,
xdm.event.outcome shows as grayed out. In
these templates, the ACTION...

📝🟢

note

This order might change if you include a
filtering field that is listed in the fieldset. In that
case, the field is taken out of the fieldset and
ordered at the top of the list with the other
filtering fields.

7.2.3.4. Run a free text query 

📝🟢

note

Free-text search is also available in XQL
queries. You can use the search stage to query
free-text strings in specific datasets, or all of
the datasets in your tenant.

📝🟢

note

Free text search is limited to the last 90 days of
data. Specifying a time frame outside of this
limitation will cause the query to fail.

🧭➡️ 7.2.3.5. Query Builder template examples 

navi…

Investigation & Response → Search → Query
Builder

7.2.4. Overview of the Query Center 

📝🟢

note

Very short queries might not be listed.
You cannot cancel correlation queries.
The default retention period for historic queries
is aligned with issue retention.

7.2.4.1. Edit and run queries in Query Center 

📝🟢

note

If query limits are applied to your tenant, the
number of concurrent running queries is limited
per user. If query usage is reaching the defined
limit, a system message warns you that a high
query load is impacting performance. If you
exceed the limit, new queries are blocked until
query usage drops. You can view all active
queries under Query Center...

📝🟢

note

You can cancel your own queries. To cancel
queries run by other users, you must have
View/Edit permissions for Configurations →
Query Management. By default, Instance
administrators have View/Edit permission.

📝🟢

note

Cancelled queries show a Canceled status. You
can see details of all canceled queries in the
Query History tab.
You cannot cancel correlation rule queries.
If you cancel a scheduled query, only the
current query is cancelled. Future recurrences
of the scheduled query are not affected.

🧭➡️

navi…

Investigation & Response → Search → Query
Center → Query History

📝🟢 7.2.4.1.1. Query Center reference information 

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

7.2.5. Manage scheduled queries 

🧭➡️

navi…

Investigation & Response → Search →
Scheduled Queries

📝🟢 7.2.5.1. Scheduled Queries reference information 

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

🧭➡️ 7.2.6. Manage your personal query library 

navi…

From the Query Builder Select Investigation &
Response → Search → Query Builder → XQL

Investigation & Response → Search → Query
Builder → XQL

Investigation & Response → Search → Query
Center

7.2.7. Legacy Query Builder 

📝🟢

note

We recommend using the Query Builder in New
mode to take advantage of the Query Builder
templates and the ability to search the full
Cortex Data Model (XDM).
In Legacy mode, the Query Builder searches
predefined datasets only. To search the full
XDM, switch to New mode or select XQL Search.

7.2.7.1. Create authentication query 

7.2.7.2. Create event log query 

7.2.7.3. Create file query 

7.2.7.4. Create image load query 

7.2.7.5. Create network connections query 

📝🟢 7.2.7.6. Create network query 

note

When you run the query, depending on the
outcome of the results, the value specified in
this field might be displayed in the dst_ip field
in the query results. This occurs if an RDP event
is recorded whereby a user connected from the
source IP to the destination IP.

7.2.7.7. Create process query 

👀🟠 7.2.7.8. Create registry query 

imp…

Ensure the KEY NAME is entered as a real
registry key name, and not as a symbolic link.
Otherwise, the query will not retrieve results.
Instead of
HKEY_LOCAL_MACHINE\System\CurrentControl
Set, which is a symbolic link, use
KEY_LOCAL_MACHINE\System\ControlSet001.
Instead of HKEY_CURRENT_USER, use
HKEY_USERS\<SID>, where SID is either a SID of
the current...

7.2.7.9. Query across all entities 

7.3. Cortex XQL syntax, parameters, and
examples 

8. Troubleshoot 

8.1. About health issues 

☑️🟣

prer…

For Cortex AgentiX to monitor data ingestion
health and create health issues, you must
enable the following settings under
Configurations:
Cortex - Analytics: Go to Configurations →
Cortex - Analytics. For more information, see
Enable the Analytics Engine and Identity
Analytics.Enable the Analytics Engine and
Identity Analytics

📝🟢

note

The Health Issues page displays issues that
were triggered after July 2024. To see health
issues that were triggered before this date,
click Legacy Health Issues.

📝🟢

note

Cortex AgentiX enforces the dedup logic to
health issues. This logic reduces the likelihood
of identical health issues from flooding the
issues dataset.

8.1.1. Monitor data ingestion health 

📝🟢 8.1.2. Monitor correlation rules 

note

In the dataset, the Query start time and Query
end time indicate the timeframe of the data
that was queried. The actual start and end
times of the correlation rule execution are
recorded in the _time field for the Initiated and
Completed entries.

8.1.3. Investigate and resolve health issues 

8.2. In-product support case creation 

📝🟢

note

When opening a support case through the
Customer Support Portal, you need to manually
select Cortex AgentiX as the product. While
there may be discrepancies between the
categories in this wizard and the Customer
Support Portal process, that's because this
wizard is designed specifically to focus on
options relevant to Cortex AgentiX.

📝🟢

note

Selecting an endpoint from the endpoint table
and retrieving TSF requires full Retrieve
Endpoint Data permissions under Endpoint
Administration.

📝🟢

note

Console recordings cannot exceed 10 minutes.
The current recording time is displayed at the
top of the window.

9. Reference 

9.1. Cortex AgentiX API 

⚠️🟡

caut…

You can't view the API Key again after you
complete this step so ensure that you copy it
before closing the notification.

🧭➡️

navi…

Settings → Configurations → Integrations →
API Keys

9.2. Fair Usage policy for Cortex AgentiX 