Cortex Cloud

Posture Management
Documentation -

Detailed

1. Get started with Cortex Cloud 

1.1. What is Cortex Cloud?  1.1.1. Key features 

1.2. What is Cortex Cloud Posture Management? 

1.3. Agentic AI in Cortex Cloud 

💡🔵

tip

The system Help Center agent delivers fast,
context-aware assistance to answer your
questions. You can ask natural language
questions, such as "How do I create a
dashboard?" or "Where can I review my data
retention policies?" and the agent retrieves
concise, relevant information from the
documentation. If a question remains
unresolved, the agent...

🧭➡️

navi…

Settings → Configurations → General → Server
Settings → Agentic Assistant

1.3.1. Agentic Assistant use cases  📝🟢

note

If a team member in the channel sees the
summary and attempts to ask the agent, "Give
me more details on the first SLA issue," the
team member receives an access denied
message because the active session is only
available to you, the initiator.

1.3.2. Agentic Assistant security 

1.4. Supported web browsers  📝🟢

note

Cortex Cloud is optimized for best
performance using the HTTPv2 protocol. If you
are using a non-supported browser version or a
network middleware that uses a protocol other
than HTTPv2, you may experience slowness in
loading and response times between the
browser and the Cortex Cloud service.
Consult with your browser or middleware
vendor to see how...

1.5. Use the interface 

1.6. In-product support case creation 

📝🟢

note

When opening a support case through the
Customer Support Portal, you need to manually
select Cortex Cloud as the product. While
there may be discrepancies between the
categories in this wizard and the Customer
Support Portal process, that's because this
wizard is designed specifically to focus on
options relevant to Cortex Cloud.

📝🟢

note

Console recordings cannot exceed 10 minutes.
The current recording time is displayed at the
top of the window.

1.7. Understand your user persona 

1.8. Fair Usage policy for Cortex Cloud 

1.9. Understand license plans 

📝🟢

note

A Cloud Runtime license includes both Posture
Scanning and Runtime Protection on the same
asset. Usage, including any overflow, is tracked
automatically to ensure accurate reporting
across both licenses without duplicate counting.

1.9.1. Data retention 

📝🟢

note

Case data is retained according to the Last
Updated date.
Issue data is retained according to the
Observation Time. Data collected within these
dates is kept and displayed for 186 days. To
ensure the accuracy of issues, Cortex Cloud
provides a grace period of up to 31 days for
issues displayed in the Issues View, Issues table,
and Cases View.

📝🟢

note

Requires the Forensics add-on.

2. Onboard and configure Cortex
Cloud 

2.1. Plan and prepare 

2.2. Deployment steps and checklist 

2.2.1. Activate Cortex Cloud 

☑️🟣

prer…

The Cortex Cloud activation email.
A Customer Support Portal (CSP) account.
You need to set up your CSP account. For more
information, see How to Create Your CSP User
Account.
When you create a CSP account, you can set
up two-factor authentication (2FA) to log into
the CSP by using an Email, Okta Verify, or
Google Authenticator (non-FedRAMP
accounts)....

📝🟢

note

Activation can take about an hour and does not
require you to remain on the activation page.
Cortex Cloud sends a notification to your email
when the process is complete.

📝🟢

note

You can only change the subdomain once, and
it cannot be undone.
After deleting the subdomain, you can reuse it
after 7 days.

2.2.1.1. Cortex Cloud supported regions 

2.2.1.2. Enable access to required PANW
resources



📝🟢

note

<tenant-name> refers to the selected
subdomain of your Cortex Cloud tenant, and
<region> is the region in which your tenant is
deployed. For more information, see Cortex
Cloud supported regions.

📝🟢

note

Refer to Regional Docker registry mapping for
your specific tenant location and
corresponding Docker registry URL.

🧭➡️

navi…

SSL decryption: If you use SSL decryption and
experience difficulty connecting the Cortex
XDR agent to the server, we recommend that
you add the FQDNs required for access to your
SSL Decryption Exclusion list in Device →
Certificate Management → SSL Decryption
Exclusion

2.2.1.2.1. Regional egress resources 

2.2.1.2.2. Engines IP addresses (outbound) 

2.2.1.2.3. Inbound source resources 

2.2.1.2.4. FedRamp and the US Federal
Government required resources



2.2.2. Upgrade from Prisma Cloud to Cortex
Cloud 

2.2.2.1. About the Upgrade Helper  📝🟢

note

Your Prisma Cloud tenant remains available and
fully operational during the agreed upgrade
period.

2.2.2.2. Link Cortex Cloud to Prisma Cloud  🧭➡️

navi…

Settings → Configurations → Upgrade Helper

2.2.2.3. Copy content 

2.2.2.3.1. Copy Global configurations 

📝🟢

note

Keep the following caveats in mind:
Scope-Based Access Control (SBAC)
configurations, such as resources or account
lists are not copied. You can manually assign
scope-based access to the relevant users or
groups.

When migrating permission groups and roles,
the total count of items successfully copied
may be lower than the initial number selected.
This...

🧭➡️

navi…

Settings → Configurations → Access
Management

2.2.2.3.2. Copy CSPM configurations  🧭➡️

navi…

Investigation & Response → Automation →
Automation Rules

Posture Management → Compliance →
Standards

2.2.2.3.3. Copy CWP configurations  🧭➡️

navi…

Posture Management → Rules & Policies →
Cloud Workload

2.2.2.3.4. Copy Cortex Cloud Application
Security configurations  👀🟠
imp…

If labels are imported into Cortex Cloud less
than three hours after they were last added or
modified in Prisma Cloud, they may not be
included in the import.

2.2.2.4. Migrate Cortex CLI  ☑️🟣

prer…

Before you begin, ensure you have the following:
Cortex Cloud API key: An active API key for
your Cortex Cloud tenant with associated CLI
role permissions. Refer to Manage API keys for
more information
Install the Cortex CLI. You can find the
installation instructions here

2.2.3. Set up users and roles 

📝🟢

note

The user who activated the Cortex product is
assigned the Account Admin role. You cannot
create additional Account Admin roles in the
Cortex Cloud tenant. If you do not want the
user to have Account Admin permission, you
must remove the Account Admin role in Cortex
Gateway.

📝🟢

note

Relevant for Behavioral Analytics & Identity
Threat Detection and Response (ITDR).

📝🟢

note

Relevant for Behavioral Analytics & Identity
Threat Detection and Response (ITDR).

📝🟢

note

Relevant for Identity Posture & Cloud
Infrastructure Entitlements (CIEM).

📝🟢

note

Relevant for Identity Posture & Cloud
Infrastructure Entitlements (CIEM).

💡🔵

tip

To quickly see exactly which pages and actions
a role allows, click on the role name, which
opens a read-only view of all checked
permissions.

🧭➡️

navi…

Settings → Configurations → Access
Management → Roles

2.2.3.1. User group management 

📝🟢

note

If users have been created in the CSP, but you
want them to access the tenant through SSO
only, skip this field and add only SAML group
mapping after SSO is set up, otherwise, users
can access the tenant through both the CSP
and SSO.
If you have not yet created any users, skip this
field and add them later. See Set up
authentication .

📝🟢

note

Capitalization is vital. String evaluation enforces
absolute case mapping rules. The name must
match your IdP's string configuration exactly.
When using Microsoft Entra ID for SSO, the
SAML group mapping needs to be provided
using the group object ID (GUID) and not the
group name.

Relevant strategy: For functional context and
the session mechanics of...

📝🟢

note

Visibility of Security domain Issues that refer to
assets with agents is controlled by the
Endpoints scoping configuration.

📝🟢

note

To automatically synchronize group
membership with your organization's Active
Directory, you can import an AD group. When
someone joins or leaves a team in AD, their
Cortex permissions update automatically.
The Import AD Group feature is only enabled
when the Cloud Identity Engine (CIE) is
connected and configured.

📝🟢

note

Only CSP and SSO users already existing in
Cortex will be imported.

📝🟢

note

When using Microsoft Entra ID for SSO, the
SAML group mapping needs to be provided
using the group object ID (GUID) and not the
group name.

📝🟢

note

If an imported group is later deleted from your
Active Directory, Cortex Cloud automatically
deletes the corresponding user group at the
next sync cycle.

👀🟠

imp…

The Cloud Identity Engine (CIE) is used
exclusively for directory group management; it
is not utilized for individual user account
management, provisioning, or authentication
workflows. Consequently, disabling, suspending,
or removing user objects directly within CIE
does not automatically disable, restrict, or
delete those corresponding users inside...

👀🟠

imp…

By default, Enable Scope Based Access
Control is disabled in Settings →
Configurations → General → Server Settings,
and granular scoping is not enforced. Before
enabling SBAC, we recommend that an
administrator or a user with Access
Management permissions first ensures that the
users, user groups, and API Keys defined in
Cortex Cloud are granted...

2.2.3.2. Assign user roles and groups 

📝🟢

note

If an existing user in the Cortex Gateway no
longer has a role or a user group assigned, the
user is revoked. Any roles, user groups, or
egress configurations created by that user are
shown as created by Revoked user instead of
the user's email address.

📝🟢

note

Visibility of Security domain Issues that refer to
assets with agents is controlled by the
Endpoints scoping configuration.

👀🟠

imp…

Setting Cortex Query Language (XQL) dataset
access permissions for a user role can only be
performed from Cortex Cloud Access
Management. For more information, see
Manage user roles.

👀🟠

imp…

Before configuring, ensure that you review
Understand scoping in the Manage user scope
section.

👀🟠

imp…

By default, Enable Scope Based Access
Control is disabled in Settings →
Configurations → General → Server Settings,
and granular scoping is not enforced. Before
enabling SBAC, we recommend that an
administrator or a user with Access
Management permissions first ensures that the
users, user groups, and API Keys defined in
Cortex Cloud are granted...

💡🔵

tip

To apply the same settings to multiple users,
select them, and then right-click and select Edit
Users Permissions.

2.2.4. Manage API keys 

📝🟢

note

Visibility of Security domain Issues that refer to
assets with agents is controlled by the
Endpoints scoping configuration.

👀🟠

imp…

Before configuring, ensure that you review
Understand scoping in the Manage user scope
section.

👀🟠

imp…

By default, Enable Scope Based Access
Control is disabled in Settings →
Configurations → General → Server Settings,
and granular scoping is not enforced. Before
enabling SBAC, we recommend that an
administrator or a user with Access
Management permissions first ensures that the
users, user groups, and API Keys defined in
Cortex Cloud are granted...

👀🟠

imp…

You will not be able to view the API key again
after you complete this step. Ensure that you
copy the API key before closing the notification.

2.2.5. Set up authentication 

💡🔵

tip

You should have at least one user in the
Customer Support Portal for backup, in case of
any authentication issues with your IdP provider.

2.2.5.1. Authenticate users through the
Customer Support Portal 

📝🟢

note

You must be assigned the Super User role in the
Customer Support Portal to add users in the
Customer Support Portal.
The first Super User who logs into Cortex
Gateway is automatically assigned the Account
Admin role and has access to the tenant. The
user who activates the Cortex Cloud tenant will
also be assigned the Account Admin role (if
there is no...

📝🟢

note

After users are added to the Customer Support
Portal and they accept the invitation, you can
manage them in Cortex Gateway or the Cortex
Cloud tenant.

2.2.5.2. Authenticate users using SSO 

📝🟢

note

To set up SSO authentication in the tenant, you
must be assigned an Instance Administrator or
Account Admin role.
SAML 2.0 users must log in to Cortex Cloud
using the FQDN (full URL) of the tenant. To
allow login directly from the IdP to , you must
set the relay state on the IdP to the FQDN of
the tenant.
If you have multiple tenants, you must set up...

📝🟢

note

The first SSO cannot be deleted, it can only be
deactivated by toggling SSO Enabled to off.
The Domain parameter is predefined for the
first SSO.
If you add additional SSO providers, you must
provide the email Domain in the SSO Integration
settings for all providers except the first.
Cortex Cloud uses this domain to determine to
which identity provider...

📝🟢

note

Cortex Cloud requires the IdP to send the
group membership as part of the SAML token.
Some IdPs send values in a format that include
a comma, which is not compatible with Cortex
Cloud. In that case, you must configure your IdP
to send a single value without a comma for
each group membership. For example, if your
IdP sends the Group DN (a...

👀🟠

imp…

PROGRAMMATIC CONTRAINT:
There is no public API endpoint available to
provision or de-provision users
programmatically within Cortex Cloud. All
target accounts must be initialized or explicitly
managed using the native interactive Single
Sign-On (SSO) or Customer Support Portal
(CSP) interface workflows defined in this guide.
To review the list of...

2.2.5.3. Set up Okta as the Identity Provider
Using SAML 2.0  📝🟢
note

When using SAML 2.0, users are required to
authenticate by logging in directly at the tenant
URL. They cannot log in via Cortex Gateway.

2.2.5.4. Set up Microsoft Entra ID as the Identity
Provider Using SAML 2.0 

📝🟢

note

The default attributes shown on the main single
sign-on page in Microsoft Entra ID are not the
values you need. You must click Edit next to
Attributes and Claims to view and copy the
actual values.

📝🟢

note

The attribute claim name must exactly match
the value sent by your IdP. In some cases, this
may be the full attribute name/namespace,
depending on the configuration of our IdP

📝🟢

note

When using SAML 2.0, users are required to
authenticate by logging in directly at the tenant
URL. They cannot log in via Cortex Gateway.

👀🟠

imp…

When copying the Single Sign-On URL value,
remove idp/saml and leave the trailing /.
For example, if the Single Sign-On URL is
https://clientname.panproduct.region.paloaltone
tworks.com/idp/saml, just copy
https://clientname.panproduct.region.paloaltone
tworks.com/.

2.2.6. Cloud service provider (CSP) onboarding 

2.2.6.1. Cloud service provider onboarding 

2.2.6.2. Amazon Web Services cloud onboarding 

2.2.6.2.1. Onboard Amazon Web Services 

2.2.6.2.2. Prerequisites  📝🟢

note

The permissions included in Sid:
CortexBaseModuleDeployment are always
required for the discovery module and stack
lifecycle management.
The permissions included in Sid:
CortexAuditLogsModule are required only
when deploying the audit log collection
module. It enables creating S3, SQS, SNS,
CloudTrail, and KMS resources.

2.2.6.2.3. How to onboard Amazon Web Services 

📝🟢

note

Scanning with an outpost may require
additional AWS permissions and may incur
additional CSP costs.

👀🟠

imp…

It is critical to ensure that your KMS key region
and SNS topic region are the exact same as the
AWS region where you are deploying the
CloudFormation stack.

2.2.6.2.4. Deploy the CloudFormation template
in AWS 

2.2.6.2.5. AWS post-deployment verification 

2.2.6.3. Microsoft Azure cloud onboarding 

2.2.6.3.1. Onboard Microsoft Azure 

2.2.6.3.2. Prerequisites 

2.2.6.3.3. How to onboard Microsoft Azure 

📝🟢

note

Scanning with an outpost may require
additional GCP permissions and may incur
additional CSP costs.

📝🟢

note

Cortex Cloud performs a live verification of
each tenant's approval status against Azure
each time the list is displayed. If a previously
approved tenant no longer shows a green
checkmark, the Cortex Cloud application may
have been removed from the Azure tenant.
Click Approve in Azure to re-authorize.

👀🟠

imp…

It is critical to ensure that your namespace and
Event Hub belong to the specific Azure
subscription being onboarded. Cross-
subscription or centralized logging is not
currently supported.

2.2.6.3.4. Finalize Microsoft Azure onboarding
by executing the authentication template 

☑️🟣

prer…

Before you begin, ensure you have:
An Azure subscription.
A user with the required permissions for the
relevant scope (subscription, management
group, tenant). We recommend you create a
dedicated role.
Tenant ID and subscription ID. You can view
these in Microsoft Azure Portal in Management
groups.

Installed Terraform on your local machine. You
can...

☑️🟣

prer…

Before you begin, ensure you have:
An Azure subscription.
A user with the required permissions for the
relevant scope (subscription, management
group, tenant). We recommend you create a
dedicated role.
Tenant ID and subscription ID. You can view
these in Microsoft Azure Portal in Management
groups.
Installed the Azure CLI tool.
Authorization to create...

☑️🟣

prer…

Before you begin, ensure you have:
An Azure subscription.
A user with the required permissions for the
relevant scope (subscription, management
group, tenant). We recommend you create a
dedicated role.
Tenant ID and subscription ID. You can view
these in Microsoft Azure Portal in Management
groups.
Installed the Azure CLI tool.
Authorization to create...

👀🟠

imp…

Do not delete or move the Terraform files from
this folder. It will prevent you from being able
to edit your cloud instance in the future.

2.2.6.4. GCP cloud onboarding 

2.2.6.4.1. Onboard Google Cloud Platform 

2.2.6.4.2. Prerequisites 

2.2.6.4.3. How to onboard Google Cloud
Platform



📝🟢

note

Scanning with an outpost may require
additional GCP permissions and may incur
additional CSP costs.

📝🟢

note

If you want to connect to your GCP Workspace,
you must first complete onboarding with the
option disabled. Once the GCP cloud instance
is created, perform the steps detailed in
Connect Google Workspace with your GCP
cloud instance.

2.2.6.4.4. Deploy the Terraform authentication
template in GCP 

☑️🟣

prer…

Before you begin, ensure you have:
A GCP account.
Permission to create the required resources in
Google Cloud Deployment Manager.
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the GCP gcloud CLI tool.
Reviewed the...

📝🟢

note

The directory you create must be a
subdirectory of the home directory.

👀🟠

imp…

You must not delete or move the Terraform files
from this folder. It will prevent you from being
able to edit your cloud instance in the future.

2.2.6.4.5. Connect Google Workspace with your
GCP cloud instance 

☑️🟣

prer…

Ensure you have the Super Admin role in
Google Workspace.

☑️🟣

prer…

Ensure you have the organization ID of the
Google Workspace you want to connect:
Log in to your Google Admin Console. and
navigate to Account → Account settings →
Profile. Next to Customer ID is your
organization ID.
Ensure you have the organization ID of the
Google Workspace you want to connect:
Log in to your Google Admin Console. and
navigate...

2.2.6.4.6. Monitor GCP resources inside service
perimeters

 📝🟢

note

Each GCP cloud instance is assigned a scope
within GCP. If the scope, whether it be
organization, folder, or project, includes any
projects with a service perimeter, this
procedure must be performed for that cloud
instance to authorize Cortex Cloud to scan the
resources in the perimeter.

2.2.6.5. Oracle Cloud Infrastructure cloud

onboarding 

2.2.6.5.1. Onboard Oracle Cloud Infrastructure 

2.2.6.5.2. Prerequisites 

2.2.6.5.3. How to onboard Oracle Cloud
Infrastructure



2.2.6.5.4. Deploy the Terraform authentication
template in OCI  ☑️🟣
prer…

Before you begin, ensure you have:
An Oracle Cloud Infrastructure account and
the tenancy OCID.
Permission to deploy a custom template and
create its resources in OCI.
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed...

2.2.6.6. Alibaba Cloud cloud onboarding 

2.2.6.6.1. Security capabilities and deployment
planning 

2.2.6.6.2. Resource inventory 

2.2.6.6.3. Security model and authentication 

2.2.6.6.4. Onboard Alibaba Cloud 

2.2.6.6.5. Prerequisites 

2.2.6.6.6. How to onboard Alibaba Cloud 

☑️🟣

prer…

Before you begin, ensure you have:
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the Alibaba Cloud CLI tool.

📝🟢

note

If the following error appears, contact support:
"Validation failed: No valid managed outpost
found for cloud provider ALIBABA_CLOUD".

👀🟠

imp…

Do not delete or move the Terraform files from
this folder. It will prevent you from being able
to edit your cloud instance in the future.

2.2.6.6.7. Alibaba Cloud post-deployment
verification 

2.2.6.7. Manually connect a cloud instance 

2.2.6.8. Manage cloud instances  📝🟢

note

Currently, automation permission errors or
missing automation permissions do not affect
the Automation health status. You can view any
permission errors or missing permissions in the
the Permissions health status.

2.2.6.9. Pending cloud instances 

2.2.6.10. Edit your onboarded CSP configuration  👀🟠

imp…

When using Terraform authentication
templates, you must execute the updated
Terraform template from the same folder where
the original Terraform template was executed.

2.2.6.11. Update cloud permissions after Cortex
release updates 

☑️🟣

prer…

Ensure that the user account used to modify
permissions has the necessary privileges within
both the Cortex platform and your cloud
environment, for example, AWS or Azure.
You received a notification regarding a new
version available that requires permission
updates, or viewed a Needs Update status in
the Data Sources & Integrations page.

📝🟢

note

Instances requiring updates will not change
their connection status, for example,
Connected, Warning, Error, Disabled, due to the
pending permission update.

2.2.6.12. Troubleshoot errors on cloud instances  📝🟢

note

Errors related to collection of audit logs in the
cloud instance are recorded in the
collection_auditing dataset. For more
information, see Audit logs fields and query
examples.

2.2.6.13. Outposts 

👀🟠

imp…

Outpost scan is an alternative to the
recommended standard cloud scan. Cloud scan
is recommended because it is fully managed by
Palo Alto Networks and incurs no compute
costs for your organization. Outpost scan is an
advanced deployment model reserved for
specific data residency or architectural
requirements.

2.2.6.13.1. Outpost fundamentals and planning 

📝🟢

note

Notes:

Before you create your outpost, verify that
your internet connection is active. An active
internet connection is necessary for the
notification to be sent to Cortex Cloud to
create the new outpost.
(Azure) Due to limitations in Terraform, the
Azure subscription name cannot contain blanks.
Take this into account while onboarding.

👀🟠

imp…

While outposts provide maximum control over
the scanning environment, cloud scan mode is
the recommended default for most
organizations.

2.2.6.13.2. Create an outpost 

☑️🟣

prer…

Before you begin, ensure you have:
An AWS account
Permission to create a stack and its resources in
AWS
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the AWS CLI tool and configured your
profile with the aws...

☑️🟣

prer…

Before you begin, ensure you have:
A GCP account
Permission to create the required resources in
Google Cloud Deployment Manager
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the GCP gcloud CLI tool

☑️🟣

prer…

Before you begin, ensure you have:
An active Azure subscription.
Installed the Azure CLI tool.
Permission to deploy a custom template and
create its resources in Microsoft Azure
("Owner" or "Contributor" on the designated
outpost subscription scope, and Active
Directory "Cloud Application Administrator" or
"Application Administrator" privileged...

📝🟢

note

Verify that your internet connection is active.
An active internet connection is necessary for
notifications to be sent to Cortex Cloud for
creating the new outpost. If you are unable to
establish an internet connection, contact
customer support for a manual workaround.

📝🟢

note

When creating an outpost for a specific Azure
subscription, the outpost account must be in
the same Azure organization as the monitored
subscriptions.

📝🟢

note

Due to limitations in Terraform, the Azure
subscription name cannot contain blanks.

👀🟠

imp…

While outposts provide maximum control over
the scanning environment, cloud scan mode is
the recommended default for most
organizations. For details, see When to choose
outpost scan.

2.2.6.14. Introduction to Terraform for Cloud

service provider (CSP) onboarding 

👀🟠

imp…

Always run these commands in the same folder
where the original .tf files and .terraform folder
live—this is where the state is stored.

👀🟠

imp…

This .tfstate state file is critical because it is
needed by the terraform destroy command to
clean up created resources. It is critical that
you never delete this file.

2.2.6.15. Cloud service provider permissions 

⚠️🟡

caut…

Reducing permissions below the recommended
level may cause specific commands to fail or
limit functionality in future pack updates.

2.2.6.15.1. Amazon Web Services provider
permissions 

📝🟢

note

Unified Cortex platform cloud content packs
require a specific set of automation permissions
to enable full integration with your cloud
environment. Before configuring access for
these packs, review the automation permission
scope guidelines.

👀🟠

imp…

All conditional capabilities documented below
require the mandatory Base and Discovery
Engine permissions to be deployed alongside
them. Base provides the foundational
CortexPlatformRole and AWS-managed read-
only baseline. Discovery Engine extends that
baseline with the asset-inventory coverage that
every other capability assumes.

2.2.6.15.2. Microsoft Azure provider permissions  📝🟢

note

Unified Cortex platform cloud content packs
require a specific set of automation permissions
to enable full integration with your cloud
environment. Before configuring access for
these packs, review the automation permission
scope guidelines.

2.2.6.15.3. Google Cloud Platform provider
permissions  📝🟢
note

Unified Cortex platform cloud content packs
require a specific set of automation permissions
to enable full integration with your cloud
environment. Before configuring access for
these packs, review the automation permission
scope guidelines.

2.2.6.15.4. Oracle Cloud Infrastructure provider
permissions 

2.2.7. Onboard the Kubernetes Connector 

📝🟢

note

This option is not supported for Fargate.

📝🟢

note

Enabling Realtime Protection installs the agent
on your Kubernetes clusters as a DaemonSet.

📝🟢

note

For Fargate, you must provide the cluster
resource identifier.
The format of the identifier is arn:aws:eks:
<region>:<account-id>:cluster/<cluster-name>.

📝🟢

note

Basic authentication is only supported in
Posture Management. If deploying Realtime
Protection, select None .

📝🟢

note

If you install the latest version but select the N-
1 strategy, this policy will take effect starting
from the next upgrade cycle (it will not
immediately downgrade your current
installation).
If you choose an older version and keep the
latest strategy, the latest version will be
installed.

2.2.7.1. What's new in Kubernetes Connector? 

2.2.7.2. Supported Kubernetes distributions 

📝🟢

note

Does not include Autopilot.

📝🟢

note

In Google Container-Optimized OS release 100
and earlier, where the FANOTIFY EXEC flag is
not supported, the Kernel configuration may be
partial for the user mode agent to properly
function. In such cases, the agent will fallback
to asynchronous mode.
In RHCOS version 4.12 and earlier, the Kernel
configuration may be partial for the user mode
agent to...

2.2.8. FedRAMP overview 

2.2.8.1. Cortex Cloud federal compliance 

2.2.8.2. Onboarding & configuration  📝🟢

note

If you choose an outpost scan, you must select
an outpost that has the environment type
defined as Government to maintain compliance
and connectivity.

2.2.8.3. Limitations & supported regions 

2.3. Post-deployment steps 

2.3.1. Set up your environment 

2.3.1.1. Configure server settings 

📝🟢

note

Keyboard shortcuts, timezone, and timestamp
format are not set universally and only apply to
the user who sets them.

📝🟢

note

The Cortex Agentic Assistant and AI case
summarization are currently available for users
in limited regions. For more information, see
Agentic AI in Cortex Cloud.
For multi-tenant/MSSP environments, the
Cortex Agentic Assistant and AI case
summarization are not available in the main
tenant.

📝🟢

note

If the Password Protection (for downloaded
files) setting under Settings → Configuration →
General → Server Settings is enabled, enter the
password 'suspicious' to download the file.

📝🟢

note

In MSSP environments, this option is not
available on the parent tenant.

🧭➡️

navi…

Settings → Configurations → General → Server
Settings

2.3.1.2. Configure security settings 

2.3.1.3. Data and log forwarding 

2.3.1.3.1. Forward logs and data from Cortex
Cloud to external services 

2.3.1.3.1.1. Configure external applications for
forwarding  📝🟢
note

There are two options for configuring external
applications. To configure relevant external
applications before you begin creating
forwarding notifications, follow the steps in the
following topics using the menu path Settings
→ Configurations → Integrations → External
Applications. You can also configure an
external application as part of the...

2.3.1.3.1.2. Configure notification forwarding 

☑️🟣

prer…

Before you can select an external service for
notification forwarding, you must integrate the
external service with Cortex Cloud. For more
information, see Configure external
applications for forwarding. No prior
configuration is required to send data to an
email distribution list.

📝🟢

note

Forwarding destinations: Only issues and cases
can be forwarded to Slack, Splunk, Amazon
SQS, Amazon S3, or Webhook.
Notification forwarding by domain: To
configure notification forwarding for issues by
domain, select Issues and filter the Issues table
by Issue Domain.
Alert vs. issue format:By default, new
configurations use the issue format, but you...

📝🟢

note

Not all data and log types can be sent to all
external services. For more information, see
Forward logs and data from Cortex Cloud to
external services.Forward logs and data from
Cortex Cloud to external services

📝🟢

note

The Grouping Timeframe defines the time
frame, in minutes, of how often Cortex Cloud
sends notifications. Every 20 issues or 20 events
aggregated within this time frame are sent
together in one notification, sorted according
to severity. To send a notification when one
issue or event is generated, set the time frame
to 0.

2.3.1.3.1.3. Monitor administrative activity 

2.3.1.3.2. Data and log notification formats 

📝🟢

note

Issues can be forwarded to email, syslog
servers, and Slack in the alert format, if you
prefer. The alert format can be selected when
you configure your forwarding notification.

2.3.1.3.2.1. Management audit log messages 

2.3.1.3.2.2. Issue notification format  📝🟢

note

For issues with relevant assets, issue
notifications sent to Amazon S3, Amazon SQS,
Webhook, Splunk, and email provide asset and
remediation information including the asset
name, cloud resource name, asset tags,
account name, region, and evidence.

2.3.1.3.2.3. Management Audit log notification
format 

2.3.1.3.2.4. Log format for IOC and BIOC issues 

2.3.1.3.2.5. Analytics log format 

2.3.2. Cortex MCP server  2.3.2.1. Cortex MCP server overview 

📝🟢

note

The Cortex MCP Server empowers you to
integrate AI into your security workflows using
natural language. When using LLM-based
suggestions, always review and approve actions
suggested by the AI before they're executed.
We recommend deploying the Cortex MCP
server in a secure environment where access is
limited to authorized users.

2.3.2.1.1. Install the Cortex MCP server 

☑️🟣

prer…

If you are running the Cortex MCP server in a
Poetry virtual environment, you must have
Python 3.13 or higher.
If you plan to run the Cortex MCP server in a
Docker container, you must have Docker
installed.

📝🟢

note

The MCP Server uses public APIs to
communicate and is limited by the license
quotas available in your tenant. This is
particularly relevant when running XQL queries.
For more information on running XQL query
APIs, see Run XQL query APIs.

📝🟢

note

It is critical to avoid assigning excessive
permissions when creating an API key for the
Cortex MCP Server. Since the key has both read
and write capabilities, overly broad permissions
can lead to unintended actions and potentially
compromise your environment. Ensure the key
follows the principle of least privilege and is
granted only the minimum...

📝🟢

note

When using Docker, we recommend using an
.env file to set the Cortex API credentials as
environment variables. While the credentials
can be provided in the MCP client configuration
settings, the .env file provides safer handling of
API credentials and makes your configuration
easily reproducible.

📝🟢

note

By default, stdio (standard input/output) is
used. You can also configure Streamable HTTP,
to send requests directly to the tenant instead
of through the MCP client. Streamable HTTP
can be useful for testing in the browser without
a MCP client and to bypass limits that may be in
place for your MCP client. For Docker, you can
include the Streamable HTTP...

👀🟠

imp…

To configure the Cortex MCP Server, you need
the Cortex API URL, Cortex API key, and Cortex
API key ID. You will not be able to view the API
key again after you complete this step. Ensure
that you copy the API key before closing the
notification.

2.3.2.1.2. Configure the MCP client  🧭➡️

navi…

Settings → Developer → Edit Config

2.3.2.1.3. Use the Cortex MCP server  📝🟢

note

The built-in tools retrieve information, but do
not write to the tenant. You can create your
own tools that include write actions. The
examples below include both.

2.3.2.1.4. Create custom Cortex MCP server

tools  📝🟢
note

Any new or updated components provided by
Cortex are automatically downloaded into the
remote_components folder. During each
update, the folder is fully replaced and all
existing contents are recreated. Do not add
custom tools to this directory, as it is managed
entirely by Cortex and is overwritten at every
update.

2.3.3. Manage user roles and access
management 

☑️🟣

prer…

Managing users, roles, scopes, user groups,
authentication settings in Cortex Cloud Access
Management requires View/Edit RBAC
permissions for Access Management (under
Configurations). Account Admin and Instance
Administrator roles are granted this permission
by default. For more information, see
Predefined user roles in Set up users and roles.

📝🟢

note

To remove users added to your CSP account,
you must do so in the CSP, not in Cortex
Gateway.

📝🟢

note

For more information on assigning user roles
when generating an API key, see Manage API
keys.

📝🟢

note

Some features are license-dependent.
Accordingly, users may not see a specific
feature if the feature is not supported by the
license type or if they do not have access
based on their assigned role or scope.

📝🟢

note

For features where scoping is not applicable,
Role-Based Access Control (RBAC) is used and
can be configured when managing user roles.
For more information, see Manage user roles.

2.3.3.1. Manage user roles  ☑️🟣

prer…

Managing user roles in Cortex Cloud Access
Management requires View/Edit RBAC
permissions for Access Management (under
Configurations). Account Admin and Instance
Administrator roles are granted this permission
by default. For more information, see
Predefined user roles in Set up users and roles.

2.3.3.2. Manage user access 

☑️🟣

prer…

Managing users, roles, scopes, user groups,
authentication settings in Cortex Cloud Access
Management requires View/Edit RBAC
permissions for Access Management (under
Configurations). Account Admin and Instance
Administrator roles are granted this permission
by default. For more information, see
Predefined user roles in Set up users and roles.

📝🟢

note

You can only reduce the permissions of an
Account Admin user via Cortex Gateway.

📝🟢

note

User permissions for components and datasets
are based on the access permissions set in the
user role. For more information on editing these
user role permissions, see Manage user roles.

📝🟢

note

Visibility of Security domain Issues that refer to
assets with agents is controlled by the
Endpoints scoping configuration.

👀🟠

imp…

Before configuring, ensure that you review
Understand scoping in the Manage user scope
section.

👀🟠

imp…

By default, Enable Scope Based Access
Control is disabled in Settings →
Configurations → General → Server Settings,
and granular scoping is not enforced. Before
enabling SBAC, we recommend that an
administrator or a user with Access
Management permissions first ensures that the
users, user groups, and API Keys defined in
Cortex Cloud are granted...

💡🔵

tip

To apply the same settings to multiple users,
select them, and then right-click and select Edit
User Permissions.

💡🔵

tip

To apply the same settings to multiple users,
select them, and then right-click and select Edit
User Permissions.

💡🔵

tip

To apply the same settings to multiple users,
select them, and then right-click and select Edit
User Permissions.

🧭➡️

navi…

Settings → Configurations → Access
Management → Users

2.3.3.2.1. User access reference information 

2.3.3.3. Manage user scope 

☑️🟣

prer…

Configuring user scopes in Cortex Cloud
Access Management requires View/Edit RBAC
permissions for Access Management (under
Configurations). Account Admin and Instance
Administrator roles are granted this permission
by default. For more information, see
Predefined user roles in Set up users and roles.
By default, Enable Scope Based Access
Control is...

📝🟢

note

Use the existing Realm attribute whenever you
need to scope based on the Account ID.

📝🟢

note

This configuration can impact the visibility of
the related Security domain in the Cases and
Issues scope area, but with not affect asset
visibility.

📝🟢

note

XQL-based dashboard widgets may require a
few hours to initially reflect changes to the list
or definitions of asset groups used for scoping.
To view the most current data immediately,
refresh the dashboard or its XQL widgets.

📝🟢

note

For Cases and Issues domains, a workaround is
to create a Dataset View for each required
combination of domains, and allow the relevant
entity access only to this Dataset View, not to
the underlying cases and issues datasets.

📝🟢

note

Make sure to assign the required default
granular scoping for users. This depends on the
structure and divisions within your organization
and the particular purpose of each
organizational unit to which scoped users
belong.

📝🟢

note

Visibility of Security domain Issues that refer to
assets with agents is controlled by the
Endpoints scoping configuration.

👀🟠

imp…

Some areas and features in Cortex Cloud do
not respect SBAC. In these cases, use RBAC
permissions to restrict access.

👀🟠

imp…

Before configuring, ensure that you review the
Understand scoping section.

2.3.3.4. Manage access to objects 

📝🟢

note

Only the Owner or an Administrator can delete
a custom object.

👀🟠

imp…

While Per-object access controls the visibility
of the object (such as a dashboard or saved
query), the underlying data remains governed
by Scope-Based Access Control (SBAC). A user
must have the appropriate SBAC permissions to
view the data available through an object.

2.3.3.4.1. Manage access to custom dashboards 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for for all custom
objects, including dashboards:
Enable "Owners can Share objects they
created": Grants owners...

📝🟢

note

Only Account Admins and Instance
Administrators have the authority to change the
owner of an object.

📝🟢

note

If you're designated as an Editor, you can
always duplicate the widget and make your
changes on the copy.

2.3.3.4.2. Manage access to report templates 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for all custom
objects, including report templates:
Enable "Owners can Share objects they
created": Grants...

📝🟢

note

Only Account Admins and Instance
Administrators have the authority to change the
owner of an object.

2.3.3.4.3. Manage access to playbooks and
scripts 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for all custom
objects, including report templates:
Enable "Owners can Share objects they
created": Grants...

📝🟢

note

Playbooks can only be Enabled when Scripts
are Enabled first.

📝🟢

note

To detach a playbook or script, a user must
have both the Create and the Edit Public sub-
permissions assigned to their role.
Reattaching an existing automation to a
Marketplace content pack can be performed
by the playbook or script Owner or an
Administrator.

2.3.3.4.4. Manage access to saved queries 

☑️🟣

prer…

Configure tenant-level settings: An
administrator must first establish the sharing
framework under Settings → Configurations →
Access Management → Objects.
The configuration of these settings defines the
authorized sharing workflows for saved queries
in the Query Library, including the options that
appear to users when clicking the three dot,...

📝🟢

note

If the role of a user is set to Edit Public Queries
but not Create Queries, they can update
existing public queries, but the Save as drop-
down menu will be hidden, preventing them
from creating new Query Library entries.

📝🟢

note

When the tenant-level setting Owners and
editors can change the general access is
unselected, the drop-down is disabled and only
an administrator can configure this option.

2.3.4. Configure the Cortex Agentic Assistant 

2.3.4.1. Agentic Assistant components and
concepts  📝🟢
note

Agents are bound by the same rules and robust
permissions as a human user. In addition, you
can mark actions that make real-world changes
in production systems as sensitive, requiring a
quick manual review and confirmation, ensuring
peace of mind before critical system changes
are made.

2.3.4.2. Agents Hub 

📝🟢

note

To manage agents in the Agents Hub, you must
have the proper permissions. For more
information, see Agentic Assistant role-based
access control.

2.3.4.2.1. Manage actions 

📝🟢

note

To manage actions in the Agents Hub, you must
have the correct permissions. For more
information, see Agentic Assistant role-based
access control.

💡🔵

tip

System actions may rely on content packs that
need to be installed and configured.

2.3.4.2.2. Register actions 

📝🟢

note

To register actions, ensure you have the correct
permissions. For more information, see Agentic
Assistant role-based access control.

📝🟢

note

Mandatory arguments cannot be deselected.

📝🟢

note

The execution of system or custom actions that
are based on integration commands can be
restricted using integration permissions.

👀🟠

imp…

A single content item can be registered as
different actions, with each action using
different inputs and outputs from the same
script. Only register the same script, command,
or AI prompt as a new action if it is required for
your use case, as providing an agent with many
actions with overlapping abilities can reduce
the ability of the agent to choose...

2.3.4.2.3. Manage agents 

📝🟢

note

In some cases, an agent may include actions
with content items that are not relevant for all
licenses. If that occurs, the grey circle appears,
but you are not able to install the related
content.

💡🔵

tip

System agents can include actions that require
additional content packs to be installed and
configured. To view all actions assigned to a
system agent, including actions not available
due to missing content, click on the system
agent in the Agents Hub. There may be actions
assigned to a system agent that are not
relevant to your organization. For...

2.3.4.2.4. Build agents 

📝🟢

note

To build agents in the Agents Hub, you must
have view/edit permissions. For more
information, see Agentic Assistant role-based
access control.

📝🟢

note

If you clear the checkbox for a role, all actions
associated with that role are also cleared. The
exception is if another role is also selected,
which is associated with the same actions.
If you clear the checkbox for an action, all roles
associated with that action are cleared. For
example, if you select the Investigator role, and
Send Mail and Tavily...

2.3.4.2.5. Expand agent capabilities with MCP
integrations 

📝🟢

note

For agentic actions to be created from tools on
an MCP server, the tools must have input and
output descriptions on the MCP server. These
descriptions are required for Cortex Cloud to
understand the tools' capabilities and the
correct use cases.

📝🟢

note

If you are using OAuth-based authentication,
the Test button returns an error containing the
command to run in the playground in order to
test the connection.

📝🟢

note

If an MCP tool is removed from the MCP server,
the action will be unavailable due to missing
content. If the tool is restored on the MCP
server, the action is automatically reenabled.

2.3.4.3. Agentic Assistant role-based access
control  📝🟢
note

Agents are limited by the individual permissions
of the user. For example, if users do not have
sufficient permissions to isolate an endpoint,
they cannot use an agent to isolate an endpoint.
The execution of system or custom actions that
are based on integration commands can be
restricted to specific roles using integration
permissions.

2.3.5. XQL query management 

☑️🟣

prer…

Setting query limits requires View/Edit
permissions for Configurations → Query
Management.

📝🟢

note

A scheduled query or report is run on behalf of
the user that originally created it, even if it is
edited and run by another user.

📝🟢

note

Queries run by correlation rules are not
restricted by the query limit.
Very short queries do not count towards
concurrent queries.

📝🟢

note

To ensure optimal system performance, all
queries (user-generated and otherwise) adhere
to a default timeout limit of 60 minutes that is
defined by Palo Alto that takes priority over the
administrator defined value. Therefore,
regardless of the value specified in this field,
queries will be stopped after 60 minutes.
You can override the default timeout...

📝🟢

note

Query visibility is subject to Role-Based Access
Control (RBAC); users can't see queries for
datasets they do not have permission to access.

🧭➡️

navi…

You can find Query Management options under
Settings → Configurations → General → Query
Management

2.3.6. Customize cases and issues  2.3.6.1. Customize cases and issues 

2.3.6.1.1. Set up case scoring 

📝🟢

note

For scoped users, a small lock icon indicates
that you don't have permissions to edit a rule.

💡🔵

tip

Right-click an issue field to add it as match
criteria.

2.3.6.1.2. Create a starring configuration  💡🔵

tip

Right-click an issue field to add it as match
criteria.

2.3.6.1.3. Service Level Agreements (SLAs) for
issue resolution  📝🟢
note

These filters support filtering of whole days
only, for example Resolution SLA > 1 filters for
issues that have a Resolution SLA of greater
than one day.

2.3.6.1.4. Create issue exceptions 

☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete issue exception
approvers and to turn off and on approver
functionality.
You must have Exception Management Admin
permission to create issue exception rules,
delete or disable exception rules, and view the
All Exception Rules page.

📝🟢

note

Issue exceptions apply specifically to issues,
not individual findings.
Issue exceptions do not automatically update
standard dashboards, reporting widgets, or
compliance reporting profiles. To view the
aggregate impact of exceptions, you must
manually filter the issues table.
Exceptions do not alter Attack Path policies.
Exceptions do not bypass any...

2.3.6.1.4.1. Configure the issue exception
approval workflow



☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete approvers.

☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete approvers.

☑️🟣

prer…

You must have the Exception Approver Admin
permission to disable or re-enable the issue
exception approval workflow.

2.3.6.1.4.2. Create an issue exception rule  📝🟢

note

You can create a maximum of 10 issue
exceptions per day.

2.3.6.1.4.3. Create an exception rule from an
issue  📝🟢
note

You can create a maximum of 10 issue
exception rules per day.

2.3.6.1.4.4. View issue Exception Rules 

2.3.6.1.4.5. Disable issue exception rules  📝🟢

note

You cannot re-enable an issue exception rule
after it has been disabled.

2.3.6.1.4.6. View excepted issues 

2.3.6.1.5. Optimize case grouping in correlations  📝🟢

note

Field mapping does not guarantee grouping.
Cortex Cloud grouping uses machine learning
and platform logic that evaluates confidence,
context, relationships, and additional case
signals. Field mapping helps influence grouping
by contributing relevant artifacts, but final
grouping decisions are determined by the
platform's overall grouping logic.
For...

2.3.7. Dashboards and reports 

3. Review inventory and explore

your cloud environment  3.1. Inventory management  3.1.1. Asset management 

3.1.1.1. Asset inventory overview 

3.1.1.2. All assets  📝🟢

note

To maintain an accurate and clutter-free asset
inventory, an automated background cleanup
process periodically removes outdated assets.

3.1.1.3. All cloud assets  3.1.1.3.1. Discovery Engine 

3.1.1.4. Asset classes 

🔒🟣

cont…

This feature is included with the AppSec add-
on.

🔒🟣

cont…

This feature is included with the Attack Surface
Management (ASM) add-on.

3.1.1.4.1. AI assets  🧭➡️

navi…

Inventory → All Assets → AI

3.1.1.4.2. API assets 

3.1.1.4.3. Application assets 

📝🟢

note

The topology graph is available only when all
application components (code, pipeline, build
and deploy), are configured.

🧭➡️

navi…

Inventory → All Assets → Business Applications

3.1.1.4.4. Code and CI/CD assets 

3.1.1.4.4.1. IaC resources assets  📝🟢

note

Scope: The IaC asset represents individual
infrastructure resources defined in Terraform,
CloudFormation, or Kubernetes manifests. The
IaC asset does not represent the physical cloud
resource in the runtime environment; those are
managed under the Cloud asset class.

3.1.1.4.4.2. Repository assets 

📝🟢

note

Scope: The repository asset represents a VCS
repository onboarded into Cortex Cloud. The
repository asset does not represent container
image repositories, artifact registries, or cloud
resource inventories; those asset categories are
managed under the Compute and Cloud asset
classes respectively.

📝🟢

note

For detailed information on investigating and
remediating issues, refer to Code Security
scanners.

🧭➡️

navi…

Inventory, select All Assets → Code →
Repositories

3.1.1.4.4.3. VCS organization assets 

📝🟢

note

Scope: The VCS organization asset represents a
VCS organization discovered through an active
VCS integration. It captures the organizational
identity, provider type, and aggregated security
posture across all child entities. It does not
represent individual repositories, CI/CD
pipelines, or CI/CD instances, nor does it
represent business applications.

🧭➡️

navi…

Inventory, select All Assets → Code → VCS
Organizations

3.1.1.4.4.4. CI/CD pipelines assets 

📝🟢

note

Scope: The CI/CD pipeline asset represents a
CI/CD pipeline definition associated with an
onboarded repository or CI/CD integration. The
CI/CD pipeline asset captures the pipeline
configuration and build activitypipeline
configuration and build activity as discovered
through CI/CD scanning. The CI/CD pipeline
asset does not represent individual...

📝🟢

note

This requires active CI/CD integrations and
successful build log analysis. Pipelines without
successful build log analysis display only the
repository and pipeline nodes

📝🟢

note

Navigate to the dedicated Application Security
→ Issues → CI/CD Risks page to manage the
remediation lifecycle at scale through bulk
status updates, team assignments, and SLA
tracking for compliance monitoring.

🧭➡️

navi…

Inventory, select All Assets → Code → CI/CD
Pipelines

3.1.1.4.4.5. CI/CD instances assets 

📝🟢

note

Scope: The CI/CD instance asset represents a
CI/CD platform instance discovered through an
active CI/CD integration. The CI/CD instance
asset captures the platform identity, version,
and aggregated security postureplatform
identity, version, and aggregated security
posture across all pipelines hosted on the
instance. The CI/CD instance asset does not...

📝🟢

note

Navigate to the dedicated Application Security
→ Issues → CI/CD Risks page to manage the
CI/CD risks remediation lifecycle at scale
through bulk status updates, team assignments,
and SLA tracking for compliance monitoring.

🧭➡️

navi…

Inventory, select All Assets → Code → CI/CD
Instances

3.1.1.4.4.6. Software packages assets 

📝🟢

note

Scope: The software package asset represents
an open-source or third-party dependency
discovered through SCA scanning of
onboarded repositories. The software package
asset does not represent first-party application
code, container image layers, or cloud runtime
packages; those asset categories are managed
under their respective asset...

🧭➡️

navi…

Inventory, select All Assets → Code →
Software Packages

3.1.1.4.5. Compute assets 

📝🟢

note

The Vulnerabilities tab is only available if the
cluster you wish to analyze has a K8s connector.

🧭➡️

navi…

Inventory → All Assets → Compute

3.1.1.4.5.1. Container images assets  🧭➡️

navi…

Inventory, select All Assets → Compute →
Container Images

3.1.1.4.5.2. Serverless functions assets  📝🟢

note

For details of all serverless function issues
generated by Cortex Cloud from vulnerability
findings, refer to Serverless function usage.

3.1.1.4.5.3. VM images assets 

📝🟢

note

The data in the Scanners column is accurate
only for Core Image assets. Ignore the Scanners
value for assets categorized as Registry, Build,
or Runtime images, as it may not reflect an
accurate scan status.

📝🟢

note

The Last Scan value is only accurate for Core
Image assets; ignore the Last Scan values for
Registry, Build, and Runtime images, as they
may be incorrect.

📝🟢

note

If the VM image is not used to create any VM
instance, the graph section will show no results.
This feature enables you to precisely identify
the registry and repository source of any
running image, directly linking runtime security
findings to their origin. As a result, you can
answer audit and security questions, such as
determining which registry...

🧭➡️

navi…

Inventory, select All Assets → Compute → VM
Images

3.1.1.4.6. Data assets 

📝🟢

note

The data assets inventory requires a Cortex
XSIAM Premium license, or a base Cortex
XSIAM license that includes the Cloud Posture
Security or Cloud Runtime Security add-on.

🧭➡️

navi…

Inventory → All Assets → Data → All Data
Assets

Inventory → All Assets → Data → Backups

3.1.1.4.7. Device assets 

📝🟢

note

The device assets inventory requires deployed
Cortex XDR agents, which are included with
Cortex XSIAM Enterprise and Premium licenses,
or available as an add-on for Cortex XSIAM NG-
SIEM.

🧭➡️

navi…

Inventory → All Assets → Device → General
Devices

3.1.1.4.8. External Surface assets  🔒🟣

cont…

Requires the Attack Surface Management
(ASM) add-on

3.1.1.4.9. Identity assets 

3.1.1.4.10. Network Assets 

3.1.1.4.11. Security services assets 

3.1.1.5. Asset groups 

📝🟢

note

If an asset group uses fields outside of this
supported list, it cannot be used for scoping in
Access Management.

📝🟢

note

Note: You cannot create SBAC based on static
groups. When using dynamic asset groups for
SBAC, you can limit access based only on the
following attributes: Asset Class, Category,
Provider, Region, Organization, Realm, Business
Application Names, Kubernetes Cluster,
Kubernetes Namespace, Code Repository, and
Asset Tags.

📝🟢

note

When you create or edit an Asset Group, the
changes are applied immediately to new assets
and to existing assets that have been updated.
However, it may take a few hours for the
changes to appear on existing assets that have
not been updated.

3.1.1.6. Network configuration 

🔒🟣

cont…

This feature is included with the Attack Surface
Management (ASM) add-on.

📝🟢

note

You can add a new range that is fully contained
within an existing range, but you cannot add a
new range that partially intersects with another.

🧭➡️

navi…

Inventory → Assets → Network Configuration
to define the boundaries of your organization's
network

3.1.1.6.1. Configure your network parameters 

📝🟢

note

You can add a range that is fully contained in
an existing range, however, you cannot add a
new range that partially intersects with another
range.

🧭➡️

navi…

In Cortex Cloud , select Assets → Network
Configuration → Internal Domain Suffixes

In Cortex Cloud , select Assets → Network
Configuration → Trusted Networks

3.1.1.7. Asset Roles 

📝🟢

note

Asset Roles are available only if the Identity
Threat Module add-on is enabled.

3.1.1.7.1. Manage Asset Roles for Endpoints  📝🟢

note

The Identity Security Module add-on is required
in order have the ability to explicitly edit the
host lists assigned to asset roles.

3.1.1.7.2. Manage Asset Roles for Users 

📝🟢

note

User Role Management is available only if the
Identity Threat Module add-on is enabled.

📝🟢

note

Deleting a manually added user removes the
user from the Included Users list. If this user is
detected automatically as having this asset role
in the future, it will appear in the Included Users
list again.
Excluding a manually added user ensures that
even if in the future the user is detected as
having this asset role, this detection is
overridden and...

3.1.1.7.2.1. Honey user  ☑️🟣

prer…

The honey user role is available only if the
Identity Threat Module add-on is enabled.

3.1.1.8. Manage Asset Scores 

📝🟢

note

The Hosts tab is available if the Identity Threat
Module add-on is enabled.

📝🟢

note

As new alerts are associated with incidents, the
User and Host Scores are recalculated. You can
view the latest User and Host Scores on the
Asset Scores page, or track the Score trend on
the User Risk View and Host Risk View.

📝🟢

note

Some User Associated Insights may not appear
as part of the User Associated Incidents due to
the insight generation mechanism. For example,
when an insight related to one of the assets in
an incident is generated a few days after the
associated incident, the insight may not be
associated with the incident.

4. Review and prioritize posture
issues 

4.1. Cases and issues 

4.1.1. Overview of cases 

☑️🟣

prer…

To work with cases, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Playbooks: This component (under Investigation
& Response → Automations) must be set to
Enabled first. Role-level permissions determine
your ability to create new playbooks or edit
those marked as...

4.1.1.1. What are cases? 

4.1.1.2. Resolving cases with AI 

4.1.1.3. Case lifecycle 

4.1.1.4. Case thresholds 

4.1.1.5. Case scope and impact 

4.1.1.6. Case and issue domains 

4.1.2. Case concepts 

4.1.2.1. Issues, findings, and events 

4.1.2.1.1. Issues 

4.1.2.1.2. Findings and events 

4.1.2.2. Case grouping 

4.1.2.3. Case scoring 

4.1.2.4. Case starring  🧭➡️

navi…

You can manage all starring configurations
under Case & Issues → Case Configuration →
Starred Issues

4.1.2.5. What is Causality?  📝🟢

note

There are no CGOs in the Cloud Causality
View, when investigating cloud Cortex Cloud
alerts and Cloud Audit Logs, or SaaS Causality
View, when investigating SaaS-related alerts for
501 audit events, such as Office 365 audit logs
and normalized logs.

4.1.3. Analyze and resolve cases 

4.1.3.1. Review all cases  📝🟢

note

Administrators can enable or disable legacy
mode. Go to Configurations → General →
Server Settings → Case display modes.

4.1.3.2. Start case analysis 

📝🟢

note

This section describes the product in default
mode and using the Split view. If you are using
legacy mode, see Detailed View.

4.1.3.2.1. Agentic Assistant- Case Investigation

agent 

4.1.3.3. Establish case context 

4.1.3.3.1. AI-generated case summaries 

📝🟢

note

The AI-generated title and description is a
calculated value that is regenerated each time
you open a case.

This value is not a saved static description,
therefore it is not reflected in the saved case
names in the list of cases in the Split view , or in
the Case Name and Case Description columns
in the Table view.

🧭➡️

navi…

Configurations → General → Server Settings →
AI Configuration and enable the following
settings

4.1.3.3.2. Assess case severity and score 

4.1.3.3.3. Update case attributes  📝🟢

note

The defined values are shown in the Case Name
and Case Description columns in the Table view,
and saved to the cases dataset. The case title is
also shown the list of cases in the Split view.
These values do not replace the AI-generated
case title and summary. If the Summarize with
AI toggle is enabled, AI-generated case
summaries are automatically...

4.1.3.4. Analyze case details 

📝🟢

note

If you prefer a tabular or legacy layout, switch
the case card to the Detailed view.
This view preserves the legacy tab based
format and custom layouts, ensuring full
backward compatibility. You can switch
between the new case experience and the
legacy view based on personal workflow
preferences. For more information, see Detailed
View.

4.1.3.4.1. Grouping graph  📝🟢

note

Similar issues are displayed as individual entities
rather than in a parent/child hierarchy.

4.1.3.4.2. Evidence 

4.1.3.4.3. Issue feed 

4.1.3.4.4. Associated assets and artifacts  📝🟢

note

If you do not have permissions to access an
asset of a case (which is shown as grayed out
and locked), check your scoping permissions in
Manage Users or Manage User Groups.

4.1.3.4.5. MITRE ATT&CK tactics and techniques  📝🟢

note

This component is available for cases
associated with the Security domain or custom
domains.

4.1.3.4.6. Case timeline  📝🟢

note

If you unmark a record as evidence, the
evidence flag is removed but the underlying
timeline record and its attachments are not
deleted.

4.1.3.4.7. Detailed View  📝🟢

note

When an issue is resolved, it remains linked to a
case. Once all of the issues in a case are
resolved, the case is automatically closed.

4.1.3.5. Resolve the case 

4.1.3.5.1. Resolution Center  📝🟢

note

Tasks that are already In Progress but require
your input will appear in both the Pending and
In Progress tabs.

4.1.3.5.2. Collaborative notes and comments 

4.1.3.5.3. Resolve a case  📝🟢

note

If a case is resolved with the status Resolved -
Auto Resolved, Cortex Cloud can reopen the
case for up-to six hours if a new issue is
triggered that matches the case. The six-hour
period is defined by the timestamp of the last
issue that was grouped into the case. After the
six-hour period, any new issues are linked to a
new case for a new...

4.1.3.5.4. Resolution reasons for cases and issues 

📝🟢

note

Cases and issues resolved as True Positive and
False Positive help Cortex Cloud to identify
real threats in your environment by comparing
future cases and associated issues to the
resolved cases. Therefore, the handling and
scoring of future cases is affected by these
resolutions.

📝🟢

note

Cases and issues resolved as True Positive and
False Positive help Cortex Cloud to identify
real threats in your environment by comparing
future cases and associated issues to the
resolved cases. Therefore, the handling and
scoring of future cases is affected by these
resolutions.

4.1.3.6. Additional case actions 

4.1.3.6.1. Create a case 

📝🟢

note

To create a case manually, you must have
View/Edit permission for Cases and Issues
selected under Settings → Configurations →
Access Management → Roles → Components
→ Cases & Issues.

📝🟢

note

This option is only relevant for certain domains.

📝🟢

note

This option is only relevant for certain domains.

📝🟢

note

You can't attach files to manually created
cases.

💡🔵

tip

The issues that you link to a case can be linked
to multiple cases, and the issue domains do not
need to match the case domain.

4.1.3.6.2. Merge a case 

4.2. Investigation and response 

4.2.1. Investigate issues 

☑️🟣

prer…

To work with issues, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Playbooks: This component (under Investigation
& Response → Automations) must be set to
Enabled first. Role-level permissions determine
your ability to create new playbooks or edit
those marked as...

4.2.1.1. Overview of the Issues page  📝🟢

note

Every 12 hours, the system enforces a cleanup
policy to remove the oldest issues once the
maximum limit is exceeded. The default issue
retention period in Cortex Cloud is 186 days.

4.2.1.2. Issue card  📝🟢

note

This section is context-specific and shows data
according to the issue context.

4.2.1.3. Resolution actions 

4.2.1.4. Link or unlink issues from a case 

4.2.1.5. Run an automation on an issue 

📝🟢

note

In addition to automation, some playbooks
contain manual tasks that prompt the analyst
for input. This enables you to enhance an
automation workflow with analyst input.

📝🟢

note

You can also manually select a playbook to run
from the Issue Work Plan tab.

4.2.1.6. Use the War Room in an investigation 

📝🟢

note

The case War Room is usually used for
communication capabilities, but unlike the issue
War Room, it does not include playbook
specific entries. The case War Room enables
you to investigate an entire case, not just an
issue.

📝🟢

note

Cortex Cloud does not index notes and chats.

💡🔵

tip

In the Playground, you can clear the context
data, if needed, which deletes everything in the
Playground context data, but does not affect
the actual issue or case. To clear the context,
run !DeleteContext all=yes' from the CLI or
click Clear Context Data while viewing the
context data.

💡🔵

tip

You can use the up/down arrow buttons in the
CLI to do a reverse history search for previous
commands with the same prefix.

🧭➡️

navi…

Investigation & Response → Automation →
Playground

4.2.1.7. Use the Work Plan in an investigation  👀🟠

imp…

A playbook will not continue its execution path
if a prior task has failed; you must resolve the
failed task before subsequent tasks can run.

4.2.1.8. Issue syncing 

☑️🟣

prer…

You must set up an integration before you can
sync issues. For more information, see Set up an
integration for mirroring issues.

☑️🟣

prer…

You must set up an integration before you can
sync issues. For more information, see Set up an
integration for mirroring issues.

☑️🟣

prer…

You must set up an integration before you can
sync issues. For more information, see Set up an
integration for mirroring issues.

📝🟢

note

Multiple tickets can be linked to an issue with
outbound syncing. Issues with inbound syncing
can be linked to a single ticket only.

📝🟢

note

Using issue fields as variables is not currently
supported.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

📝🟢

note

Using issue fields as variables is not currently
supported.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

📝🟢

note

The inbound syncing flow runs every two
minutes, and the outbound syncing flow runs
every five minutes.
In a bi-directional set-up, if the same field is
updated in both tickets, the most recently
updated value is used.
In the external ticket, the logged history shows
updates to the ticket. The user name that is
logged with the history reflects the user...

📝🟢

note

After an issue is resolved, ticket syncing
remains active for up-to seven days. Therefore,
you still update, change, or reopen the issue or
external ticket and the tickets will continue to
sync.

📝🟢

note

If you change the selected inbound sync
profile, the original sync profile is immediately
overwritten.

📝🟢

note

You can only define a single inbound profile. If
you change the inbound sync profile the
current profile is overwritten.
You can define multiple outbound profiles; one
issue can update multiple tickets.

⛔🔴

war…

If you leave this field blank, all configured
instances will be used.

⛔🔴

war…

If you leave this field blank, all configured
instances will be used.

💡🔵

tip

You can find a sync profile ID under Settings →
Configurations → Object Setup → Issues →
Sync Profiles. By default the ID field is not
displayed in the table. Click the three dot menu
and add it to the table layout.

4.2.1.9. Issue deduplication 

📝🟢

note

Issues are automatically excluded from
deduplication if the agent_id is missing, the
hash_id is missing, or the hash_id is an all-zero
SHA256 string.

🧭➡️

navi…

File Identification (Hash): Use the SHA256
hierarchy (Action File → Action Process →
Actor Process

4.2.1.10. Issue investigation actions 

4.2.1.10.1. Copy issues 

4.2.1.10.2. Update issue fields 

📝🟢

note

You must create a custom status before you
can select it.

💡🔵

tip

You can create custom issue statuses and
resolution reasons, and use the setIssueStatus
command to set these custom statuses for
issues.
For example, when a user starts investigating an
issue, the issue status is automatically changed
from New to Under Investigation. In some cases,
it is useful to create an interim status, such as
Triage. After you...

4.2.1.10.3. Export issue details to a file 

4.2.1.10.4. Exclude an issue  📝🟢

note

Cortex Cloud supports exclusion of up to
100,000 issues.

4.2.1.10.5. Query case and issue data 

4.2.2. Review findings  4.2.2.1. Findings card  📝🟢

note

The information in this card is context specific,
therefore some sections are not available for all
findings.

4.2.3. Investigate artifacts and assets 

4.2.3.1. Investigate an IP address  📝🟢

note

Requires a license key. Select Settings →
Configurations → Integrations → Threat
Intelligence.

4.2.3.2. Investigate an asset  📝🟢

note

The Asset view is available for hosts with a
Cortex XDR agent installed.

4.2.3.3. Investigate a file and process hash  📝🟢

note

Requires a license key. Go to Settings →
Configurations → Integrations → Threat
Intelligence.

4.2.3.4. Investigate a user 

📝🟢

note

Cortex Cloud normalizes and displays case and
issue times in your time zone. If you're in a half-
hour time zone, the activity in the Issues &
Insights Heatmap is displayed in the whole-hour
time slot preceding it. For example, if you're in a
UTC +4.5 time zone, the time displayed for the
activity will be UTC +4.5, however, the
visualization in the...

💡🔵

tip

You can also see a list of all users under
Inventory → Assets → Asset Scores.

4.2.4. Cortex Assistant 

📝🟢

note

If you are in an eligible region and have enabled
the Cortex Agentic Assistant, the Cortex
Agentic Assistant replaces the Cortex Assistant.
The Cortex Assistant is available if you do not
have access to the Cortex Agentic Assistant
based on the tenant region or you have not
enabled it. For more information, see Cortex
Agentic Assistant.Cortex Agentic...

🧭➡️

navi…

To increase usability, you can create a
personalized keyboard shortcut: Settings →
Configurations → Server Settings → Keyboard
Shortcuts and choose the shortcut you want to
use

4.2.4.1. Cortex Assistant layout 

4.2.4.2. Cortex Assistant capabilities 

📝🟢

note

In some cases, if the prompt does not include at
least one recognizable entity such as an IP,
hash, user, asset, domain, case, or XQL query,
no response is returned.

📝🟢

note

When you choose an option from the Respond
column, Cortex Assistant will always prompt
you to approve the action before executing.

4.2.5. Automation 

4.2.5.1. Automation in Cortex Cloud  📝🟢

note

You can build end-to-end automation
workflows from within the playbook editor,
including creating automation rules,
configuring integration instances, and creating
and editing tasks. For more information, see
Playbooks.

4.2.5.2. Quick Actions  📝🟢

note

Quick Actions appear as War Room entries, but
do not appear in the Work Plan.

4.2.5.3. Automation Exclusion Center 

📝🟢

note

By default, all users have read and edit
permissions to lists. When creating a list of
critical assets, we recommend limiting the read
and edit permissions to specific roles.

4.2.5.3.1. Manage automation exclusion policies 

📝🟢

note

For the IAM User Hard Remediation and User
Soft Remediation policies, we recommend
including username, email, and ID for each user
you want to exclude. Example: username1,
user@example.com, userID112.

📝🟢

note

You can also right click on a policy from the
main Automation Exclusion Center page to
disable or enable the policy.
If you click on a list name in the Exclude
column, that list opens in the Lists page.

4.2.5.4. Playbooks 

☑️🟣

prer…

To work with playbooks, an administrator must
configure your user role with specific RBAC
permissions. Permissions must be enabled in the
following order:
Scripts: This component (under Investigation &
Response → Automations) must be set to
Enabled first. It is the foundational permission
for all automation; if Scripts are not enabled,
you cannot...

4.2.5.4.1. Playbooks overview 

4.2.5.4.2. Access to playbooks 

4.2.5.4.3. Playbook development checklist 

4.2.5.4.4. Plan your playbook 

4.2.5.4.5. Manage playbooks  📝🟢

note

The library by default shows only playbooks
that are not adopted. Click the Show Adopted
checkbox to show the adopted playbooks,
indicated by an Adopted mark.
The library shows the most updated playbook
version. Adopting an older version than shown
should be done through Marketplace.
Adopting a playbook does not make it run.
Some content packs include...

4.2.5.4.6. Build your playbook 

👀🟠

imp…

The ability to create, edit, or share custom
playbooks is governed by access management.
If certain options are unavailable, contact your
administrator. For more information, see
Manage access to playbooks and scripts.

4.2.5.4.6.1. Task 1. Choose from existing
playbooks or create your own



📝🟢

note

When a task inside of a sub-playbook is not
configured, the alert is propagated to the main
playbook. If multiple sub-playbooks are nested,
and any of the sub-playbooks have non-
configured components, the alert appears in
the main playbook as well as in the sub-
playbooks. Alerts also appear for the individual
non-configured tasks within the...

📝🟢

note

You must have the required permissions to
create playbooks to view this button. For more
information, see Manage access to playbooks
and scripts.

💡🔵

tip

If there are additional relevant playbooks in
Marketplace that are not in your Org repository,
you can click Explore them now to see them in
the Playbook Catalog and choose to adopt.

💡🔵

tip

To open multiple playbooks at the same time,
edit the first playbook and then click New next
to the playbook name to create a new tab. You
can either create a new playbook, or add an
existing one.
You can view recently modified or deleted
playbooks by clicking version history for all
playbooks .

🧭➡️

navi…

Open the Investigation & Response →
Automation → Playbooks page to find an
existing playbook, customize it, or create a
playbook

4.2.5.4.6.2. Task 2. Configure playbook settings 

📝🟢

note

If the playbook has inputs and outputs, the
Playbook Starts task will show back and forth
arrows. Clicking them opens the Playbook
Settings pane Inputs/Outputs tab.

📝🟢

note

For more information on Automation Rules, see
Create an automation rule.

📝🟢

note

This rule will trigger the playbook to run if no
other Automation Rule triggers the playbook
first. You can view and edit the order the rules
run in the Automation Rules page.

📝🟢

note

If you do not add any fields, the group will be
deleted when you click Save.

4.2.5.4.6.3. Task 3. Add objects from the Task
Library



4.2.5.4.6.4. Task 4. Add custom playbook
features



4.2.5.4.6.5. Task 5. Test and debug the playbook 

4.2.5.4.6.6. Task 6. Manage playbook content 

4.2.5.4.7. Customize your playbook 

4.2.5.4.7.1. Configure a sub-playbook loop 

📝🟢

note

Consider the following when adding a loop:
The maximum number of loops (default is 100).
A high number of loops or a high wait time
combined with a large number of issues may
affect performance.
Periodically check looping conditions to ensure
they are still valid for the data set.
If you want a sub playbook task to loop over an
array passed into its...

💡🔵

tip

Balance between the number of iterations and
the interval so you do not overload the server.

4.2.5.4.7.2. Filter and transform data 

⚠️🟡

caut…

You can change the context data root to filter,
but it is not recommended to select a different
root, as it affects the filter results. The drop-
down list displays the filter root for backward
compatibility.

🧭➡️

navi…

Investigation & Response → Automation →
Scripts → New Script

4.2.5.4.7.3. Extend context 

4.2.5.4.7.4. Extract indicators 

📝🟢

note

This configuration may delay playbook
execution (issue creation).
While indicator creation is asynchronous,
indicator extraction and enrichment are run
synchronously. Data is placed into the issue
context and is available via the context for
subsequent tasks.

📝🟢

note

When using out of band, the extracted
indicators do not appear in the context. If you
want the extracted indicators to appear select
inline.

4.2.5.4.7.5. Update issue fields with playbook
tasks

 📝🟢

note

The setIssue script includes all available fields;
use the scroll bar to see all the fields.
The name field has a limit of 600 characters. If
there are more than 600 characters, you can
shorten the name field to under 600 characters
and then include the full information in a long
text field such as the description field.

4.2.5.4.8. Test your playbook 

📝🟢

note

The debugger does not support using
parentIncidentFields.

📝🟢

note

Using an existing issue in the debugger does not
affect the original issue or change the original
context data.

4.2.5.4.8.1. Troubleshoot playbook performance 

4.2.5.4.9. Manage playbook content 

📝🟢

note

Opening a playbook does not immediately lock
it. A playbook is considered to be in edit mode
only when a user makes a modification that
causes the Save button to become available.

👀🟠

imp…

Manually unlocking a playbook may cause
version conflicts. This turns off concurrent
editing protection, and if the original user is still
editing, their changes might be lost or
overwritten.

🧭➡️

navi…

Investigation & Response → Automations →
Playbooks

4.2.5.4.10. Best practices 

📝🟢

note

If you reattach a detached playbook, any
customizations you have made to the playbook
will be overwritten when the playbook updates
to the current version.

📝🟢

note

Retries are not supported for data collection
tasks that have errors sending emails (indicated
by a server timeout). This is because retries only
work on automation execution failures, not on
email delivery issues.

4.2.5.5. AI Prompts 

📝🟢

note

In MSSP environments, this option is not
available on the main tenant.

🧭➡️

navi…

On the AI Prompts page (Investigation →
Response → Automation → AI Prompts), you
can view, edit, and manually create prompts

4.2.5.5.1. AI prompts role-based access control 

4.2.5.5.2. Use existing prompts  🧭➡️

navi…

Investigation & Response → Automation → AI
Prompts and in the Prompts Library search for
the prompt you want to use

4.2.5.5.3. Create a prompt  🧭➡️

navi…

Investigation & Response → Automation → AI
Prompts and click + New Prompt

4.2.5.5.4. Write effective prompts 

📝🟢

note

An AI prompt cannot create new context keys;
it can only refer to existing ones. The only key
an AI prompt can create is the output key.
Therefore, if a variable like ${PostKey} is not
already mapped to existing context data, it
cannot be used in the prompt.

📝🟢

note

The LLM's response is always returned in a text
format, even if you explicitly ask the LLM to
return a structured format such as JSON.
Therefore, the subsequent task in your
playbook will have to parse the response to
extract the relevant data.

4.2.5.6. Create an automation rule 

📝🟢

note

You can also define the conditions that trigger
a specific playbook in the playbook editor. For
more information, see Task 2. Configure
playbook settings

📝🟢

note

Quick Actions, by default, run using all
available integration instances that contain the
command. When selecting a Quick Action for
an automation rule, you can instead choose one
specific integration instance to use.

4.2.5.7. Scripts 

4.2.5.7.1. Use existing scripts  🧭➡️

navi…

Investigation & Response → Automation →
Scripts and in the Scripts Library search for the
script you want to use

4.2.5.7.2. Create a script 

📝🟢

note

You can enable/disable a script in the Settings
without having to duplicate the script.
You can view recently modified or deleted
scripts by clicking the version history for all
scripts .

👀🟠

imp…

If you choose Python, from the Agentic
Assistant you can use the Automation Engineer
agent.

💡🔵

tip

If there is an error, you can copy the error
message from the test result into the Agentic
Assistant prompt and ask the Automation
Engineer agent to correct the error.

4.2.5.7.3. Use the Automation Engineer agent to
accelerate script development and deployment 

📝🟢

note

The Automation Engineer agent is available
with the Cortex Agentic Assistant, for users
with script editing permissions. For more
information, see Cortex Agentic Assistant and
Agentic Assistant role-based access
control.Cortex Agentic Assistant

📝🟢

note

If you make manual edits in the script and don't
save the changes and you then modify the
script with the Automation Engineer agent, you
are prompted to confirm overwriting the
manual edits.

💡🔵

tip

Use detailed prompts, for example, Get failed
logins from last 24 hours and return as a table.
Clearly define argument names, types, and
default values.
Mention if you expect the output in a specific
format, such as a table, JSON, or plain text.

💡🔵

tip

If there is an error, you can copy the error
message from the test result into the Agentic
Assistant prompt and ask the Automation
Engineer agent to correct the error.

4.2.5.7.4. Change the Docker image in a script 

4.2.5.7.4.1. Connect an engine to an image
registry  📝🟢
note

This procedure uses the --username and --
password command line options to pass the
username and password directly. For
environments where command history or logs
are visible to others, consider more secure
methods like Docker configuration files for
handling authentication in production or CI/CD
environments. For more details, see docker
login or...

4.2.5.8. Context data 

4.2.5.8.1. Issue context data 

4.2.5.8.2. Case context data 

4.2.5.8.3. Search context data 

4.2.5.8.4. Add context data to an issue 

4.2.5.8.5. Add context data to a case  📝🟢

note

If you run the command in the issue War Room,
the data is added to the following places:
The case context data.
The issue context data under the case tab.
If you run the command in the Case War Room,
the data is added to the case context data only.

4.2.5.8.6. Delete context data from a case 

4.2.5.8.7. Use context data in a playbook  ⚠️🟡

caut…

Users with Trigger Playbook permissions on a
given issue may still be able to modify the
parent case via commands and scripts, even
without full access to the case.

4.2.5.9. Lists 

📝🟢

note

The maximum size of a list is 209715 characters.

4.2.5.9.1. Create a list 

📝🟢

note

If you intend to use the list as part of an
Automation Exclusion Policy, we recommend
choosing specific roles for read and edit
access. For more information on using lists in
exclusion policies, see Manage automation
exclusion policies.Manage automation exclusion
policies

📝🟢

note

If you want to edit a list from a content pack,
you need to duplicate or detach a list.
Detached lists do not receive updated content
in subsequent Cortex Cloud content releases.
To retain an updated list, reattach it.

4.2.5.9.2. List commands  📝🟢

note

The list must already exist before you use the
getList, addToList, setList, or removeFomList
commands.

4.2.5.9.3. Use cases: JSON lists  🧭➡️

navi…

Settings → Configurations → Object Setup →
Lists → Add a List

Investigation & Response → Automation →
Playbooks → New Playbook

4.2.5.9.4. Transform a list into an array  🧭➡️

navi…

Investigation & Response → Automation →
Playbooks and create or edit a playbook

4.2.5.10. Integrations 

4.2.5.10.1. Integrations in Cortex Cloud 

4.2.5.10.2. Add an integration instance 

4.2.5.10.3. Use integration commands in the CLI  💡🔵

tip

In the Playground, you can clear the context
data, if needed, which deletes everything in the
Playground context data, but does not affect
the actual issue or case. To clear the context,
run !DeleteContext all=yes' from the CLI or
click Clear Context Data while viewing the
context data.

4.2.5.10.4. Troubleshoot Integrations  📝🟢

note

If an integration instance consistently
encounters API rate limits when running on the
tenant, consider configuring the instance to run
on an engine to change the source of the
outbound traffic.

4.2.5.10.5. Manage credentials  🧭➡️

navi…

Settings → Configurations → Integrations →
Credentials → New Credential

4.2.5.11. Engines 

4.2.5.11.1. What is an engine? 

📝🟢

note

You cannot share a multiple-engine installation
with a single-engine installation.

📝🟢

note

When you add an engine to a load-balancing
group, you cannot use that engine separately.
The engine does not appear in the engines
menu when configuring an integration instance,
but you can choose the load-balancing group.

4.2.5.11.2. Engine requirements 

📝🟢

note

The Cron package is required to install engines
on a Linux machine.

📝🟢

note

CentOS 8.x reached End of Life (EOL) on
December 31, 2021, and is no longer supported
as an operating system.
CentOS 7.x reached End of Life (EOL) on June
30, 2024, and is no longer supported as an
operating system.

📝🟢

note

If you have configured a range of Approved IP
Ranges under Allowed Sessions on the Security
Settings page, the engine must communicate
through one of the approved IPs.

4.2.5.11.3. Install an engine 

📝🟢

note

If you are using DEB, RPM, or Zip installation,
install Docker or Podman.
Natively running Python or PowerShell
integrations/scripts on Windows or Linux is not
supported on Cortex Cloud engines.

📝🟢

note

When upgrading an engine that was installed
using the Shell installation, you can use the
Upgrade Engine feature in the Engines page. For
Amazon Linux 2 type engines, you need to
upgrade these engine types using a zip-type
engine and not use the Upgrade Engine feature.
If you use the shell installer, Docker/Podman is
automatically installed. We...

📝🟢

note

Use DEB and RPM installation when the shell
installation is not available. You need to
manually install Docker or Podman and any
dependencies.

📝🟢

note

If you receive a permissions denied error, it is
likely that you do not have permission to access
the /tmp directory.
If the installer fails to start due to a permissions
issue, even if running as root, add one of the
following two arguments when running the
installer:
--target <path> - Extracts the installer files
into the specified custom...

👀🟠

imp…

For DEB/RPM engines, Python (including 3.x)
and the containerization platform
(Docker/Podman) must be installed and
configured. For Docker or Podman to work
correctly on an engine, IPv4 forwarding must be
enabled.

💡🔵

tip

For Linux systems, we recommend using the
shell installer. If using Amazon Linux 2, use the
zip installer (see step 4).

4.2.5.11.3.1. Docker 

📝🟢

note

This section is relevant when installing an engine.

📝🟢

note

Docker images can be downloaded together
with their relevant content packs for offline
installation.

4.2.5.11.3.2. Podman  📝🟢

note

When upgrading an engine, the engine keeps
the previously used container management
type (regardless of distribution version).

4.2.5.11.4. Manage engines 

4.2.5.11.5. Upgrade an engine 

4.2.5.11.6. Remove an engine 

4.2.5.11.7. Configure engines 

📝🟢

note

A day is defined as 24 hours and may not
exactly correspond to calendar days due to
daylight savings, leap seconds, etc.

📝🟢

note

In addition, to support engine upgrades from
the UI, edit the /usr/local/demisto/upgrade.conf
file on the engine to include the SERVER_URLS
setting with the new tenant's address. Include
only the host, without https:// or any additional
path at the end of the host name. For example:
SERVER_URLS="api-
example.us.paloaltonetworks.com"

📝🟢

note

If the engine loses communication for longer
than this time, it will disconnect and you need
to restart the service.

4.2.5.11.7.1. Configure the engine to use a web
proxy



📝🟢

note

You need to configure Docker to use a proxy.
When using a BlueCoat proxy, ensure you
encode the values correctly.

📝🟢

note

In an environment with a single engine, go to
/usr/local/demisto/upgrade.conf. In an
environment with multiple engines, go to
/usr/local/demisto/<engine-
name>/upgrade.conf, replacing <engine-name>
with the name of the engine.
Note that the key is in the upgrade.conf file and
must be https_proxy, even if your proxy address
starts with http://.

4.2.5.11.7.2. Configure the engine to call the
server without using a proxy



4.2.5.11.7.3. Configure an engine to use custom
certificates 

4.2.5.11.8. Use an engine in an integration  📝🟢

note

Long-running integrations should not run on
load-balancing groups.

4.2.5.11.9. Run a script using an engine  🧭➡️

navi…

Investigation & Response → Automation →
Scripts

4.2.5.11.10. Troubleshoot engines 

📝🟢

note

Ensure that pop-ups are not blocked by your
browser.

📝🟢

note

If the installer fails to start due to a permissions
issue, even if running as root, add one of the
following two arguments when running the
installer:
--target <path> - Extracts the installer files
into the specified custom path.
--keep - Extracts the installer files into the
current working directory (without cleaning at
the end).
If using installer...

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

📝🟢

note

Any changes made to variables in the
upgrade_engine.sh file are reset after each
upgrade. We recommend instead using the
upgrade.conf file to set variables.

📝🟢

note

You can ignore the following error: Cannot
create folder '/var/lib/demisto'

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

📝🟢

note

If the installer fails to start due to a permissions
issue, even if running as root, add one of the
following two arguments when running the
installer:
--target <path> - Extracts the installer files
into the specified custom path.
--keep - Extracts the installer files into the
current working directory (without cleaning at
the end).
If using installer...

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl status d1_<Engine _name>

4.2.5.11.11. Troubleshoot integrations running on
engines 

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

📝🟢

note

Docker CE installations typically run Docker,
while Docker EE installations typically run
dockerroot.

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

📝🟢

note

If the Allow running multiple engines on the
same machine option is selected, run the
command:
sudo systemctl restart d1_<Engine _name>

4.2.6. Build XQL queries 

📝🟢

note

If you have the Cortex Agentic Assistant, you
can use natural language prompts to create
and run XQL queries within the chat interface.
For more information, see Create and run XQL
queries with Agentic Assistant chat.Cortex
Agentic Assistant

4.2.6.1. About the Query Builder 

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query, and can appear with a slight delay.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion commands and
definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc key.

4.2.6.2. How to build XQL queries 

📝🟢

note

Users with different dataset permissions can
receive different results for the same XQL query.
An administrator or a user with a predefined
user role can create and view queries built with
an unknown dataset that currently does not
exist in Cortex Cloud. All other users can only
create and view queries built with an existing
dataset.
When you have more...

📝🟢

note

You can build a query that investigates data in
both a cold dataset and a hot dataset in the
same query. In addition, as the hot storage
dataset format is the default option and
represents the fully searchable storage, this
format is used throughout this guide for
investigation and threat hunting. For more
information on hot and cold storage, see...

👀🟠

imp…

Forensic datasets are not inlcuded by default in
XQL query results, unless the dataset query is
explicitly defined to use a forensic dataset.

4.2.6.2.1. Get started with XQL queries 

4.2.6.2.2. Useful XQL user interface features 

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query and can appear with a slight delay.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

4.2.6.2.3. XQL Query best practices 

4.2.6.2.4. Expected results when querying fields 

4.2.6.2.5. Create XQL query 

📝🟢

note

Whenever the time period is changed in the
query window, the config timeframe is
automatically set to the time period defined for
the entire query, including queries that are part
of the join stage. Yet, this won't be visible as
part of the query. Only if you manually type in
the config timeframe will this be seen in the
query.

These time picker...

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

💡🔵

tip

While the query is running, you can navigate
away from the page. A notification is sent when
the query has finished. You can also Cancel the
query or run a new query, where you have the
option to Run only new query (cancel previous)
or Run both queries.

4.2.6.2.6. Review XQL query results 

📝🟢

note

It's also possible to graph the results displayed.
For more information, see Graph query results.

📝🟢

note

Real time query results are available only in the
Query Builder and in free text query.
Real time results are displayed only for queries
run on hot datasets.
The Sort option is available only after all the
data is retrieved.
When you formulate complex queries, the
results will be displayed when the query has
finished running completely, and not in real...

📝🟢

note

Results are received incrementally for the first
100K records, or up to 100MB worth of records,
whichever comes first. After that, the next
update is when the query has finished running
completely.

📝🟢

note

In order for Cortex Cloud to provide a
histogram for a field, the field must not contain
an array or a JSON object.

4.2.6.2.7. Translate to XQL  👀🟠

imp…

This feature is still in a Beta state and you will
find that not all Splunk queries can be
converted to XQL. This feature will be improved
upon in the upcoming releases to support
greater Splunk query translations to XQL.

4.2.6.2.8. Graph query results 

🔒🟣

cont…

Building Cortex Query Language (XQL) queries
in the Query Builder requires a Data Collection
add-on.

📝🟢

note

To display the result of as a time duration,
choose the graph type Single Value and enable
Show as Time. You can then select the Time
Unit (millisecond, second, minute, or hour) and
the Display format.

💡🔵

tip

Alternatively, you can use the Cortex Agentic
Assistant to generate custom graphs and
charts using natural language prompts. By
simply prompting the agent, it will build and
execute the query, returning the visual
representation. For more information, see Use
natural language to query and visualize your
data.

4.2.6.3. XQL query entities 

4.2.6.3.1. Create authentication query 

4.2.6.3.2. Create event log query 

4.2.6.3.3. Create file query 

4.2.6.3.4. Create image load query 

4.2.6.3.5. Create network connections query 

4.2.6.3.6. Create network query  📝🟢

note

When you run the query, depending on the
outcome of the results, the value specified in
this field might be displayed in the dst_ip field
in the query results. This occurs if an RDP event
is recorded whereby a user connected from the
source IP to the destination IP.

4.2.6.3.7. Create process query 

4.2.6.3.8. Create registry query  👀🟠

imp…

Ensure the KEY NAME is entered as a real
registry key name, and not as a symbolic link.
Otherwise, the query will not retrieve results.
Instead of
HKEY_LOCAL_MACHINE\System\CurrentControl
Set, which is a symbolic link, use
KEY_LOCAL_MACHINE\System\ControlSet001.
Instead of HKEY_CURRENT_USER, use
HKEY_USERS\<SID>, where SID is either a SID of
the current...

4.2.6.3.9. Query across all entities 

4.2.6.4. Overview of the Query Center 

📝🟢

note

Very short queries might not be listed.
You cannot cancel correlation queries.
The default retention period for historic queries
is aligned with issue retention.

4.2.6.4.1. Edit and run queries in Query Center 

📝🟢

note

If query limits are applied to your tenant, the
number of concurrent running queries is limited
per user. If query usage is reaching the defined
limit, a system message warns you that a high
query load is impacting performance. If you
exceed the limit, new queries are blocked until
query usage drops. You can view all active
queries under Query Center...

📝🟢

note

You can cancel your own queries. To cancel
queries run by other users, you must have
View/Edit permissions for Configurations →
Query Management. By default, Instance
administrators have View/Edit permission.

📝🟢

note

Cancelled queries show a Canceled status. You
can see details of all canceled queries in the
Query History tab.
You cannot cancel correlation rule queries.
If you cancel a scheduled query, only the
current query is cancelled. Future recurrences
of the scheduled query are not affected.

4.2.6.4.1.1. Query Center reference information  📝🟢

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

4.2.6.5. Manage scheduled queries 

4.2.6.5.1. Scheduled Queries reference

information  📝🟢
note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

4.2.6.6. Manage your query library 

📝🟢

note

The Query to Library option is only available if
your role has the Create Queries capability. For
more information, see Manage access to saved
queries.

📝🟢

note

When the tenant-level setting Owners and
editors can change the general access is
unselected, the drop-down is disabled and only
an administrator can configure this option.

👀🟠

imp…

The ability to create, edit, or share queries is
governed by access management. If certain
options are unavailable, contact your
administrator. For more information, see
Manage access to saved queries.

4.2.7. Quick Launcher  📝🟢

note

For hosts, Cortex Cloud displays results for
exact matches but supports the use of
wildcard (*) which changes the search to return
matches that contain the specified text. For
example, a search of compy-7* will return any
hosts beginning with compy-7 such as compy-
7000, compy-7abc, and so forth.

4.3. Customize cases and issues  4.3.1. Customize cases and issues 

4.3.1.1. Set up case scoring 

📝🟢

note

For scoped users, a small lock icon indicates
that you don't have permissions to edit a rule.

💡🔵

tip

Right-click an issue field to add it as match
criteria.

4.3.1.2. Create a starring configuration  💡🔵

tip

Right-click an issue field to add it as match
criteria.

4.3.1.3. Service Level Agreements (SLAs) for
issue resolution 

📝🟢

note

These filters support filtering of whole days
only, for example Resolution SLA > 1 filters for
issues that have a Resolution SLA of greater
than one day.

🧭➡️

navi…

Settings → Configurations → Object Setup →
Issues

Settings → Configurations → Object Setup →
Issues and select the SLA Rules tab

4.3.1.4. Create issue exceptions 

☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete issue exception
approvers and to turn off and on approver
functionality.
You must have Exception Management Admin
permission to create issue exception rules,
delete or disable exception rules, and view the
All Exception Rules page.

📝🟢

note

Issue exceptions apply specifically to issues,
not individual findings.
Issue exceptions do not automatically update
standard dashboards, reporting widgets, or
compliance reporting profiles. To view the
aggregate impact of exceptions, you must
manually filter the issues table.
Exceptions do not alter Attack Path policies.
Exceptions do not bypass any...

4.3.1.4.1. Configure the issue exception approval
workflow



☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete approvers.

☑️🟣

prer…

You must have Exception Approver Admin
permission to add or delete approvers.

☑️🟣

prer…

You must have the Exception Approver Admin
permission to disable or re-enable the issue
exception approval workflow.

🧭➡️

navi…

Settings → Configurations → Server Settings
and scroll to the Exception Management section

Settings → Configuration → General → Server
Settings and scroll to the Exception
Management section

4.3.1.4.2. Create an issue exception rule  📝🟢

note

You can create a maximum of 10 issue
exceptions per day.

4.3.1.4.3. Create an exception rule from an issue  📝🟢

note

You can create a maximum of 10 issue
exception rules per day.

4.3.1.4.4. View issue Exception Rules  🧭➡️

navi…

Settings → Exceptions Configuration →
Exception Rules

4.3.1.4.5. Disable issue exception rules 

📝🟢

note

You cannot re-enable an issue exception rule
after it has been disabled.

🧭➡️

navi…

Settings → Exceptions → Configuration →
Exception Rules

4.3.1.4.6. View excepted issues 

4.3.1.5. Optimize case grouping in correlations  📝🟢

note

Field mapping does not guarantee grouping.
Cortex Cloud grouping uses machine learning
and platform logic that evaluates confidence,
context, relationships, and additional case
signals. Field mapping helps influence grouping
by contributing relevant artifacts, but final
grouping decisions are determined by the
platform's overall grouping logic.
For...

5. Agentic Assistant chat 

5.1. Get started with Agentic Assistant chat  📝🟢

note

The Cortex Agentic Assistant is currently
available in limited regions. For more
information see Cortex Agentic Assistant If your
tenant is not within one of those regions, you
have access to the Cortex Assistant.Cortex
Agentic Assistant
To enable the Cortex Agentic Assistant, go to
Settings → Configurations → General → Server
Settings → AI...

5.2. Choose an Agentic Assistant agent 

📝🟢

note

Only public agents are supported via Slack.

📝🟢

note

Requires the Exposure Management add-on.

5.3. Chat with an Agentic Assistant agent 

📝🟢

note

Case data is only loaded when you send your
first message. Opening the chat interface
without sending a prompt does not provide the
agent with the case context.

📝🟢

note

An agent's proposed plans and results may
contain inaccuracies or errors. Always review
the results carefully to ensure you fully
understand the proposed action before
proceeding.

💡🔵

tip

You can quickly jump to different product
pages within Cortex Cloud by typing / in the
prompt area. This shortcut is a built-in
navigation feature that is available even if the
Cortex Agentic Assistant is disabled.

5.4. Chat with the Agentic Assistant from Slack 

☑️🟣

prer…

Before you can interact with the Cortex
Agentic Assistant in Slack, the Slack v3
integration must be set up correctly. For more
information, see the Slack v3 integration
documentation.
To perform actions in Slack, your Slack email
must match your Cortex Cloud user email. This
ensures the system can strictly follow your
assigned permissions (RBAC). If...

📝🟢

note

Only public agents are supported via Slack.

📝🟢

note

If a team member in the channel sees the
summary and attempts to ask the agent, "Give
me more details on the first SLA issue," the
team member receives an access denied
message because the active session is only
available to you, the initiator.

5.5. Create and run XQL queries with Agentic
Assistant chat 

📝🟢

note

The TextToXQL action is a hidden system action
and does not appear in the Agents Hub.

📝🟢

note

Running XQL queries manually through an
agent does not consume compute units. This
includes scenarios where you prompt the agent
to create and execute a multi-step plan.

5.6. Use natural language to query and visualize
your data  📝🟢
note

Once the graph or chart is saved the Widget
Library, the link to the chat artifact is severed,
and the agent does not track subsequent
changes made to the widget or the dashboard.

5.7. Manage chat history 

6. Review and report your security

posture and progress

 6.1. Monitor dashboards and reports 

6.1.1. About dashboards 

📝🟢

note

Dashboards that are based on custom
infrastructure cannot be exported.
When importing a dashboard that already
exists in the system, the imported dashboard
overwrites the existing dashboard. To avoid
overwriting the existing dashboard, duplicate
and rename the existing dashboard before
importing the new dashboard.

👀🟠

imp…

The ability to create, edit, or share custom
dashboards is governed by access
management. If certain options are unavailable,
contact your administrator. For more
information, see Manage access to custom
dashboards.

6.1.1.1. Command Center dashboards 

📝🟢

note

Access to these dashboards requires RBAC
permissions under Dashboards & Reports. The
Dashboards component must be Enabled in
your role and requires View permissions for
Command Center Dashboards. If certain
options are unavailable, contact your
administrator. For more information, Manage
access to custom dashboards.
The dashboards are available in dark...

6.1.1.1.1. Cortex Agentic Assistant dashboard 

☑️🟣

prer…

The Cortex Agentic Assistant is only available
after it is enabled, and for tenants in certain
regions. For more information, see Cortex
Agentic Assistant.Cortex Agentic Assistant

📝🟢

note

Access to this dashboard requires RBAC
permissions under Dashboards & Reports. The
Dashboards component must be Enabled in
your role and requires View permissions for
Command Center Dashboards. If certain
options are unavailable, contact your
administrator. For more information, see
Manage access to custom dashboards.

6.1.1.1.2. Cloud Security Operations  6.1.1.1.2.1. Cloud Security Operations  📝🟢

note

Command Center data may not match the
counts on the Issues page, and you may
observe inconsistencies. This is because
dashboard data is a snapshot of issues
identified, whereas the Issues page provides the
most up-to-date view of risks across your
cloud assets. In addition, the Issues pages do
not support all the currently available filters on
the...

6.1.1.1.3. Cortex Cloud Command Center 

6.1.1.2. Predefined dashboards  📝🟢

note

Users can access all information on the
dashboard when their user access is scoped to
view All assets or assigned to the Instance
Administrator role. Otherwise, users with
granular scoping set to No assets or Select
asset groups will have limited access to the
dashboard. For more information on Scope-
Based Access Control (SBAC), see Manage
user...

6.1.1.3. Reports  6.1.1.3.1. Report templates 

6.1.2. Build custom dashboards and reports 

6.1.2.1. Build a custom dashboard  📝🟢

note

For agent-related widgets, you can apply an
endpoint scope to refine the displayed data to
only show results from specific endpoint
groups. Select the menu on the top right corner
of the widget, select Groups, and select one or
more endpoint groups.
For case-related widgets, you can refine the
displayed data to only show results from cases
that match a...

6.1.2.2. Manage your Widget Library 

📝🟢

note

Unlike dashboards, individual widgets cannot
be shared with other users directly from the
Widget Library. Yet, if you add a Restricted
widget to a dashboard and then share that
dashboard, other users who have access to the
dashboard will also be able to see that specific
widget.

📝🟢

note

Any dashboards or reports that include the
widget are affected by the changes.

6.1.3. Fine-tune dashboards and reports 

6.1.3.1. Create a custom widget using a script 

📝🟢

note

If you have added arguments to the script,
these appear when creating a widget.

📝🟢

note

Available options are Pie, Column, Line, and
Single Value.
To display the result of the script as a time
duration, choose the graph type Single Value
and enable Show as Time. You can then select
the Time Unit (millisecond, second, minute, or
hour) and the Display format.

6.1.3.1.1. Script-based widget examples 

📝🟢

note

Add the widget tag in the script settings to
make the script available for use in script-
based widgets. For more information, see
Create a script.

📝🟢

note

If your script returns a time duration, configure
the widget with the graph type Single Value
and enable Show as Time..

6.1.3.2. Create custom XQL widgets 

📝🟢

note

Cortex Query Language (XQL) queries
generated from the Widget Library do not
appear in the Query Center. The results are
used only for creating the custom widget.

📝🟢

note

Whenever the time period is changed in the
query window, the config timeframe is
automatically set to the time period defined,
but this won't be visible as part of the query.
Only if you manually type in the config
timeframe will this be seen in the query.
These time picker options are available in XQL
queries when using the Query Builder, XQL
Widgets,...

📝🟢

note

To display the result of as a time duration,
choose the graph type Single Value and enable
Show as Time. You can then select the Time
Unit (millisecond, second, minute, or hour) and
the Display format.

💡🔵

tip

You can create a generic dashboard for
multiple views of the same dataset by using an
asterisk (*) when defining the dataset in the
XQL widget as dataset = <dataset_name>*. The
placement of the asterisk in the dataset name
ensures that any view containing this prefix text
is displayed in the results.
The dataset in a query is defined as:
If there are...

6.1.3.2.1. Configure filters and inputs for custom
XQL widgets 

☑️🟣

prer…

Fixed filters are based on parameters that are
defined in custom XQL widgets. Before you can
configure fixed filters, take the following steps:
Create custom XQL widgets with parameters.
For more information, see Create custom XQL
widgets.
Add the widgets to a custom dashboard. For
more information, see Build a custom
dashboard.

📝🟢

note

This option only appears if the dashboard
contains custom XQL widgets with defined
parameters.

📝🟢

note

You can define up to four parameter filters on a
single dashboard or report.

📝🟢

note

If you specify more than one field, only the first
field value is used.

💡🔵

tip

After the initial setup, when you access your
dashboard the filters and inputs might need
further refinement. You can make changes to
the configured parameters in the XQL widgets,
and update the Filters & Inputs on your
dashboard until you are satisfied with the
results.

6.1.3.2.2. Configure dashboard drilldowns 

☑️🟣

prer…

To configure drilldowns your dashboard must
contain custom XQL widgets. In addition, if you
want to configure in-dashboard drilldowns your
custom XQL widget must contain one or more
parameters. For more information about
configuring parameters in custom XQL widgets,
see Create custom XQL widgets.

📝🟢

note

If the selected parameter is configured in other
XQL widgets on the dashboard, these widgets
are also affected by the drilldown.

6.1.3.2.2.1. Variables in drilldowns 

6.1.4. Run or schedule reports 

📝🟢

note

Report templates that are based on custom
infrastructure cannot be exported.
If you import a report template that already
exists in the system, the imported template will
overwrite the existing template. If you do not
want to overwrite the existing template,
duplicate and rename the existing template
before importing the new template.
You can also...

📝🟢

note

The report name and description will be
displayed in the report header and are not
editable during customization.

📝🟢

note

For agent-related widgets, you can apply an
endpoint scope to refine the displayed data to
only show results from specific endpoint
groups.

Select the menu on the top right corner of the
widget, select Groups, and select one or more
endpoint groups.
For agent-related widgets, you can apply an
endpoint scope to refine the displayed data to
only show...

📝🟢

note

To send reports to Slack, Slack must be
configured as an external application in Cortex
Cloud. For more information, see Integrate
Slack for outbound notifications

📝🟢

note

When ownership is transferred, any existing
report schedules associated with the report
template are automatically removed. The new
Owner must manually redefine the schedule to
resume automated report generation.

🧭➡️

navi…

Under Settings → Configurations → General →
Notifications, click + Add Forwarding
Configuration

7. Monitor and track compliance
adherence



7.1. Cortex compliance flow 

7.2. Choose compliance standards from the
compliance catalog 

🧭➡️

navi…

Cortex provides lists of available standards and
controls in the Standards and Controls
catalogs under Posture Management →
Compliance → Catalogs

7.2.1. Standards Catalog  7.2.1.1. Use a built-in or custom standard 

7.2.2. Controls catalog  7.2.2.1. Use a built-in or custom control 

7.2.2.1.1. Associate a custom control to a

detection rule 

📝🟢

note

Custom rules can only be associated with
custom compliance controls.

📝🟢

note

You can only associate custom compliance
controls to the following cloud security rule
types: ConfigIdentityAI

🧭➡️

navi…

Posture Management → Rules & Policies →
Rules → Cloud Workload or Cloud Security

7.2.2.1.2. Create a new Custom Detection Rule 

📝🟢

note

The custom rule must use the match term (not
deny or others) to function properly.

📝🟢

note

Only users with the following roles can enable
or disable Custom Code Execution:
Account Admin
Instance Administrator
Deployment Admin
Privileged Security Admin

📝🟢

note

Only Custom Detection Rules (not built-in
rules) can be assigned to custom controls.

👀🟠

imp…

The custom Python scripts are intended to be
executed exclusively for compliance checks
and validations. To ensure the scripts are used
properly and no security risks or unintended
changes occur, the system implements the
following restrictions and safeguards:
Only a predefined set of Python libraries and
functions required for compliance checks are...

7.3. Use an assessment profile to run
compliance checks on your assets 

🧭➡️

navi…

Under Posture Management → Compliance →
Assessment Profiles, click Create New
Assessment

7.3.1. Configuring assessments for custom
compliance standards based on custom cloud
security rules



7.4. View and manage compliance assessments
and reports 

📝🟢

note

Compliance assessment results may take up to
six hours to be generated.

7.4.1. Compliance score  📝🟢

note

The status of controls is determined by the
evaluation of the associated rules. If an asset
fails a check against any rule associated with a
control, that control is considered failed for
that asset.

7.4.2. Assessments  📝🟢

note

If there are no rules associated with the control,
the control will be assigned a severity of low.

7.4.3. Reports  🧭➡️

navi…

The Posture Management → Compliance →
Results → Reports page shows a table listing
compliance assessment report files

7.5. View the compliance assessment of an
individual asset 

7.6. Compliance Overview Dashboard  📝🟢

note

If you just created an assessment profile, it may
take up to 10 minutes for it to appear in the
Assessment filter. As a workaround, you can
refresh the URL in the browser.

8. Cortex Cloud AI Security 

8.1. What is Cortex Cloud AI Security? 

8.2. Supported services in Cortex Cloud AI
Security 

8.3. Cortex Cloud AI Security concepts 

8.4. Cortex Cloud AI Security use cases 

8.5. How to perform advanced AI Security
investigations using XQL  📝🟢
note

For more information, see Build XQL queries.

9. Cortex Cloud Application

Security 

📝🟢

note

While some features are included by default,
others require a dedicated add-on purchase.

9.1. Onboard data sources 

📝🟢

note

Third-party SCA ingestion (such as Snyk,
Semgrep) covers CVE vulnerabilities only.
License Miscompliance and Package
Operational Risk findings are produced
exclusively by the native Cortex Cloud SCA
scanner. You must enable native SCA scanning
to achieve license and operational-risk coverage

9.1.1. Onboard version control systems 

📝🟢

note

Cortex Cloud Application Security (which
includes IaC and Secrets scanning), is an add-
on to a license (such as Posture Security) that
must be purchased separately.

📝🟢

note

Disclaimer: When onboarding with third-party
data sources, we outline the required steps for
setup, but we do not monitor these external
resources, and they may change over time.
Always refer to the relevant third-party
documentation for the most current integration
steps.

💡🔵

tip

Navigate to Settings → Data Sources &
Integrations → + Add New → and enter your
VCS data source in the search bar.

9.1.1.1. AWS CodeCommit 

☑️🟣

prer…

Before you begin, ensure the following:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
AWS user permissions: To onboard
CodeCommit, the user logged into the AWS
Console must have permissions to deploy the
CloudFormation stack and authorize the...

💡🔵

tip

The instance ID is identical to the stack ID on
the AWS platform.

🧭➡️

navi…

Modules → Application Security → Periodic
scans

9.1.1.2. Azure DevOps 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Azure DevOps permissions: Ensure the user
performing the integration holds one of the
following roles in Azure DevOps:
Project Administrator: This permission is
required to subscribe...

📝🟢

note

This forces Azure to bypass browser cookies
and issue a token for the correct directory.

📝🟢

note

PATs are static. To onboard a different tenant,
you must log in to that specific environment to
generate a new token.

👀🟠

imp…

When redirected to the Microsoft login screen,
do not immediately enter your email.

🧭➡️

navi…

Settings → Data Sources & Integrations → +
Add New

User Settings → Personal access tokens → +
New Token

9.1.1.2.1. Azure DevOps onboarding system
architecture

📝🟢 

note

Microsoft has announced the deprecation of
legacy OAuth methods by 2026, making Entra
ID the required standard for long-term support.

9.1.1.3. Bitbucket Cloud 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Bitbucket permissions: In Bitbucket, grant the
user performing the Cortex application
authorization the following permissions. The
level of access required depends on the
modules you...

📝🟢

note

Ensure that you receive the Instance
Successfully Created message on this step,
indicating successful instance creation.

9.1.1.4. Bitbucket Data Center 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Bitbucket authorization (PAT): You must
generate a Personal Access Token (PAT) within
your Bitbucket Server account settings to
authenticate the connection. For detailed
instruction on...

📝🟢

note

For more information about the Transporter,
including setup instructions, refer to
Transporter over Broker VM.

📝🟢

note

Ensure that you receive the Instance
Successfully Created message on this step,
indicating successful instance creation.

📝🟢

note

By default, the permissions of the access token
are set according to your current access level.
It is essential to define two levels of
permissions, Project and Repository
permissions. The Repository permissions inherit
from Project permissions, requiring Repository
permissions to match or exceed Project
permissions
Providing read and write permissions...

📝🟢

note

For additional security, it is recommended to
set an expiry automatically. The expiry date of
a token cannot be changed after it is created.
You can see the expiry dates for all your tokens
on Profile picture → Manage account →
Personal access tokens.

👀🟠

imp…

Always refer to the Bitbucket documentation
for information relating to creating a PAT.

9.1.1.5. GitHub Cloud 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
In GitHub, grant the user performing the Cortex
application authorization the following
permissions:
Organization Owner: Only an Organization
Owner can directly authorize and install...

🧭➡️

navi…

Verify integration: On Data Sources, select
Code Providers → GitHub SaaS → View more
and confirm that the status of your integrated
GitHub instance is 'Connected

9.1.1.6. GitHub Enterprise (On-Prem) 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
GitHub permissions: You must have
Organization Owner permissions to install the
Cortex application. Users with only repository-
level admin permissions cannot complete the
installation...

📝🟢

note

The domain is the hostname associated with
your GitHub Enterprise (On-Prem) instance.

📝🟢

note

For more information about the Transporter,
including setup instructions, refer to
Transporter over Broker VM.

📝🟢

note

Ensure that you receive the Instance
Successfully Created message on this step,
indicating successful instance creation.

9.1.1.7. GitLab SaaS 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
In Gitlab, the following permissions are required
to integrate the application:
Maintainer (Project-level). Grants sufficient
permissions to configure external integrations,
manage...

📝🟢

note

A repository can only be integrated with a
single instance. The first instance that connects
with the repository will be the one that the
repository is assigned to. This means that if
multiple integrations attempt to connect to the
same repository, only the first integration to
establish the connection will be associated with
that repository.

9.1.1.8. GitLab Self Managed (On-Prem) 

☑️🟣

prer…

Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Gitlab permissions: Authorize the user
integrating Cortex Cloud Application Security
with your GitLab Self Managed (On-Prem)
instances with the following permissions:
Maintainer permissions. Grants...

📝🟢

note

The domain is the hostname associated with
your GitLab Self Managed (On-Prem) instance.

📝🟢

note

For more information about the Transporter,
including setup instructions, refer to
Transporter over Broker VM.

📝🟢

note

Ensure that you receive the Instance
Successfully Created message on this step,
indicating successful instance creation.

9.1.2. Onboard CI/CD systems 

9.1.2.1. CircleCI for CI/CD pipeline scans 

☑️🟣

prer…

Before you begin:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
CircleCI user requirements:
Permissions: To enable Cortex Cloud visibility
for all CircleCI projects, a version control
system (VCS) user with integration permissions
must be...

📝🟢

note

This integration utilizes a Personal Access
Token (PAT) for authentication
CircleCI onboarding offers both code and
CI/CD scanning. A single integrated instance
supports either code or CI scanning, but not
both. If you require both code and CI scanning
for your CircleCi environment, you must create
two separate integrations, selecting the
appropriate...

9.1.2.2. Jenkins for CI/CD pipeline scans 

☑️🟣

prer…

Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
In Jenkins:
To install and configure the Cortex plugin in
Jenkins, you must be a Jenkins Administrator
with Overall/Administer permissions
Ensure the build server allows outbound HTTPS
(Port 443...

📝🟢

note

Jenkins onboarding offers both code and CI/CD
scanning. A single integrated instance supports
either code or CI scanning, but not both. If you
require both code and CI scanning for your
Jenkins servers, you must create two separate
integrations, selecting the appropriate scanning
type for each. To onboard Jenkins for code
scans, refer to Jenkins for...

📝🟢

note

The integration is added on the console but
integration is pending, and will only be
completed after completing step 5 below. You
can view the pending integration on the Jenkins
Instances page: Select Data Sources → Jenkins
→ View Details. The type of integration is
Pipeline Risks

📝🟢

note

Always refer to the official Jenkins
documentation when installing plugins on
Jenkins servers.

9.1.3. Integrate CI tools 

🧭➡️

navi…

Settings → Data Sources & Integrations →
hover over a CI tool → View Details

Delete an instance: Right-click on an instance
of the CI tool → Delete instance → Delete

9.1.3.1. AWS CodeBuild 

☑️🟣

prer…

Before you begin:
User permissions: Ensure the user performing
the integration has permissions to edit pipeline
configurations (such as YAML files) and
manage secrets/credentials within the CI
platform to store the Cortex Cloud API key
securely

📝🟢

note

Do not change the names of the environment
variables provided by Cortex Cloud. They are
required for proper integration and
functionality.

📝🟢

note

The code is only a reference. Replace the
placeholder values with your build-specific
values.

9.1.3.2. CircleCI for code scans 

☑️🟣

prer…

Before you begin:
User permissions: Ensure the user performing
the integration has permissions to edit pipeline
configurations (such as YAML files) and
manage secrets/credentials within the CI
platform to store the Cortex Cloud API key
securely

📝🟢

note

CircleCI onboarding offers both code and
CI/CD scanning. A single integrated instance
supports either code or CI scanning, but not
both. If you require both code and CI scanning
for your CircleCi environment, you must create
two separate integrations, selecting the
appropriate scanning type for each. To onboard
CircleCI for CI/CD scans, refer to...

📝🟢

note

Do not change the names of the environment
variables provided by Cortex Cloud. They are
required for proper integration and
functionality.

👀🟠

imp…

The cortex-secrets naming convention for the
context is mandatory to ensure functionality
and must not be changed.

9.1.3.3. Connect Cortex CLI 

☑️🟣

prer…

System requirements:
macOS (Intel Core i7, such as Sequoia): To
ensure all functionalities work correctly, you
must install the vectorscan dependency via
Homebrew, using this command: brew install
vectorscan

RHEL 8.10 and Red Hat UBI9. The following
prerequisites must be met:
Install patchelf
Install zstd
Ubuntu 20 requires the prefetch...

📝🟢

note

The Cortex CLI requires an API key with the
Standard security level.

📝🟢

note

On macOS arm 64 architecture you must
unpack the downloaded file to retrieve the
executable.

9.1.3.4. GitHub Actions 

☑️🟣

prer…

Before you begin:
User permissions: Ensure the user performing
the integration has permissions to edit pipeline
configurations (such as YAML files) and
manage secrets/credentials within the CI
platform to store the Cortex Cloud API key
securely

📝🟢

note

Do not change the names of the environment
variables provided by Cortex Cloud. They are
required for proper integration and
functionality.

📝🟢

note

The code is only a reference. Replace the
placeholder values with your build-specific
values.

9.1.3.5. Jenkins for code scans 

☑️🟣

prer…

Grant Administrator permissions to the user
integrating Cortex Cloud Application Security
with Jenkins
Egress path: Create an egress path to establish
the designated route for outbound data
transmission from Cortex Cloud to third party
services. For more information about
configuring egress paths, refer to Egress
configurationsEgress configurations

☑️🟣

prer…

For Cortex Cloud Application Security CI tools,
you must store secrets in Jenkins Credentials
for use in your Jenkins pipelines using either of
these methods:
Plain text storage: Store secrets directly as
plain text in Jenkins Credentials. Access them in
your pipeline using the credentials function,
which retrieves the secret directly as plain...

📝🟢

note

Jenkins onboarding offers both code and CI/CD
scanning. A single integrated instance supports
either code or CI scanning, but not both. If you
require both code and CI scanning for your
Jenkins servers, you must create two separate
integrations, selecting the appropriate scanning
type for each. To onboard Jenkins for CI/CD
scans, refer to Jenkins for...

📝🟢

note

This step is only required for new pipelines
For private repositories, ensure the necessary
credentials are configured in Jenkins Credentials

📝🟢

note

This ensures your build runs within a Docker
environment. If a node without Docker is used,
the build will fail.

9.1.3.6. Terraform Cloud (Run Tasks) 

☑️🟣

prer…

Before you begin:
User permissions: Ensure the user performing
the integration has permissions to edit pipeline
configurations (such as YAML files) and
manage secrets/credentials within the CI
platform to store the Cortex Cloud API key
securely
Procure a Terraform cloud license that is either
a trial license or a TF Cloud license at the
TEAM &...

📝🟢

note

Skip this step if you plan on using an existing
token.

📝🟢

note

Cortex Cloud Application Security performs a
scan of Terraform templates on selected
workspaces based on the Run Stage.

9.1.3.7. Terraform Enterprise (Run Tasks) 

☑️🟣

prer…

Before you begin:
Ensure access to a Terraform Enterprise
console to enable you to provide a user or
team token that authorizes Cortex Cloud
Application Security to access workspaces and
helps regulate run configurations
Terraform Enterprise version compatibility:
Ensure Run Tasks for workspaces on is
compatible with version 1.1.9 and above
Terraform...

📝🟢

note

For container image vulnerabilities, Cortex
Cloud Application Security performs 'Image
Referencer' scans within Terraform Enterprise
(Run Tasks), as full SCA scans are not currently
supported.

📝🟢

note

Skip this step if you plan on using an existing
token.

📝🟢

note

Cortex Cloud Application Security performs a
scan of Terraform templates on selected
workspaces based on the Run Stage.

☑️🟣 9.1.4. CLI pipeline code snippets 

prer…

User permissions: Ensure the user performing
the integration has permissions to edit pipeline
configurations (such as YAML files).

9.1.5. Onboard private package registries 

9.1.5.1. JFrog Artifactory 

9.1.5.2. Onboard JFrog Artifactory 

☑️🟣

prer…

Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
JFrog permissions:
The permissions associated with the user
configured during the onboarding process
determine the scope of scan results. Only
repositories and artifacts the user can access
are...

📝🟢

note

To authorize the scanner to connect through
your firewall, select Enable access by IPs, and
copy the displayed source IPs to your
organization's allowlist.

📝🟢

note

For Maven:
Select Mirror Registry if this repository mirrors
an external repository
Use the Mirror Of value to define the
duplication scope:
* mirrors all requests
Type a request [value]: Mirrors only specific
requests (such as central).
For Maven:
Select Mirror Registry if this repository mirrors
an external repository
Use the Mirror Of...

9.1.6. Ingest third-party data sources 

📝🟢

note

Only onboarded and scanned repositories can
be mapped.

9.1.6.1. Semgrep 

☑️🟣

prer…

Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Ensure that you have a connected version
control system (VCS) and repositories
Semgrep: A valid Semgrep API token
Note
To create a Semgrep API token, in Semgrep,
navigate to Settings → Tokens →...

📝🟢

note

Mapping establishes relationships between
Semgrep projects and Cortex Cloud code
repositories, simplifying access management
and enabling risk analysis at the repository
level, including displaying findings on the tenant
Only mapped projects are ingested

📝🟢

note

Requisite RBAC permissions are required for
edit and delete operations.

👀🟠

imp…

Deletion is irreversible. To restore Semgrep
findings after deletion, reconfigure the
integration and wait for the next polling cycle
to re-ingest findings from Semgrep.

9.1.6.1.1. Semgrep Software Composition
Analysis (SCA) data ingestion 🧭➡️ 
navi…

Application Security → Issues → Vulnerabilities

9.1.6.1.2. Semgrep Static Application Security
Testing (SAST) data ingestion 🧭➡️ 
navi…

Application Security → Issues → Code
Weaknesses

9.1.6.2. Snyk 

☑️🟣

prer…

Permissions: The following user permissions are
required:
Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Ensure that you have a connected version
control system (VCS) and repositories
Snyk permissions and requirements:
Permissions: The...

📝🟢

note

Select Test Connection to verify that Cortex
Cloud can connect to your Snyk organization.

📝🟢

note

SCA requires Snyk Open Source or Snyk
Container projects configured in the
organization
SAST requires Snyk Code enabled and projects
configured in the organization

📝🟢

note

Mapping establishes relationships between
Snyk applications and Cortex Cloud code
repositories, simplifying access management
and enabling risk analysis at the repository
level, including displaying findings on the tenant
Only mapped applications are ingested

📝🟢

note

Requisite RBAC permissions are required for
edit and delete operations.

👀🟠

imp…

Snyk Free tier is not supported. A paid Snyk
plan (Team, Business, or Enterprise) is required
to ingest findings from Snyk.

9.1.6.2.1. Snyk Software Composition Analysis
(SCA) ingestion 🧭➡️ 
navi…

Modules → Application Security →
Vulnerabilities (under Issues

9.1.6.2.2. Snyk Static Application Security
Testing (SAST) data ingestion 🧭➡️ 
navi…

Modules → Application Security → Code
Weaknesses

Modules → Application Security → AppSec
Coverage

9.1.6.3. SonarQube 

☑️🟣

prer…

Cortex Cloud user permissions: Ensure you have
View/Edit permissions for Data Sources and
Integrations (RBAC: AppSec Admin or Instance
Administrator)
Ensure that you have onboarded a version
control system (such as GitHub, GitLab) and
have connected repositories
SonarQube and SonarCloud setup:
Token requirements: The API token for both
SonarCloud...

📝🟢

note

Mapping establishes relationships between
SonarQube Applications and Cortex Cloud
Application Security code repositories,
simplifying access management and enabling
risk analysis at the repository level, including
displaying findings on the tenant
Only mapped projects will be ingested

📝🟢

note

Verify that the Connector Created Successfully
message is displayed on the page.

9.1.6.4. Veracode 

☑️🟣

prer…

Permissions: The following user permissions are
required:
Cortex Cloud: Instance Admin, AppSec Admin
or GRBAC permissions. For more information on
AppSec Admin permissions, refer to Code
Security user roles and permissionsCode
Security user roles and permissions
Veracode: At minimum, Reviewer permissions
are required
Permissions: The following...

📝🟢

note

Only mapped applications will be ingested.

📝🟢

note

This is the recommended option to ensure
complete coverage and successful operation
of all features.

📝🟢

note

Mapping establishes relationships between
Veracode projects and Cortex Cloud
Application Security code repositories,
simplifying access management and enabling
risk analysis at the repository level, including
displaying findings on the tenant
Only mapped projects will be ingested

📝🟢

note

Verify that the Connector Created Successfully
message is displayed on the page.

9.1.6.5. Generic 3rd Party AppSec Collector 

9.1.6.5.1. Tenant (console) workflow 

📝🟢

note

Editing a collector does not regenerate the API
credentials. The existing credentials remain
valid after editing the collector name.

👀🟠

imp…

The Token ID and API Token are displayed only
once. Copy and store the credentials securely
before proceeding. If the credentials are lost,
edit the collector to regenerate the credentials.

⚠️🟡

caut…

Deleting a collector permanently removes the
collector instance and invalidates the
associated API credentials. Existing findings
that were previously ingested through the
deleted collector remain in the platform and
are not affected by the deletion.

🧭➡️

navi…

Settings → Data Sources & Integrations → 3rd
Party AppSec Collector

Settings → Data Sources → 3rd Party AppSec
Collector

📝🟢 9.1.6.5.2. API workflow 

note

The apiUrl returned during collector creation
resolves to the public upload endpoint at
/public_api/appsec/v1/collectors/{collectorId}.
Use the apiUrl value directly in upload requests.

9.1.6.5.3. Upload findings from CI/CD pipelines 

📝🟢

note

Store all collector credentials (Token ID, API
Token, API URL) as encrypted secrets or
credentials in the CI/CD platform. Do not
hardcode credentials in pipeline configuration
files.

👀🟠

imp…

Verify that the upload request returns a 2xx
HTTP status code. A non-2xx response
indicates that the upload failed and findings
were not ingested. Configure the CI/CD pipeline
step to fail on non-2xx responses to prevent
silent ingestion failures. In the cURL examples,
add the --fail flag or check the exit code to
detect upload failures.

9.1.6.5.4. Technical requirements and SARIF
specifications 

9.1.6.5.5. Troubleshooting 

9.1.7. Manage data source integrations 

9.1.8. Transporter over Broker VM 

☑️🟣

prer…

Ensure you have the necessary permissions and
have already set up your Broker VM
Confirm that your Broker is v 28 or above
Whitelist IP addresses to enable access to
Cortex Cloud resources. The IP addresses for
the Transporter are in the Broker VM Resources
section of the Enable access to required PANW
resources document
Open port 4052 (inbound), which...

📝🟢

note

The broker enables multiple connections,
allowing a single Transporter applet to connect
to various VCS instances
Multiple Transporter applets cannot be created
on a single Broker VM. Each Broker VM is
limited to a single Transporter applet instance

⛔🔴

war…

The Transporter applet is not supported for
FedRAMP customers.

9.1.8.1. Set up a Transporter applet on Broker VM 

☑️🟣

prer…

Permissions: To configure and manage
Transporter applet settings, you must have
permissions to manage Broker Service
configurations (such as an Instance
Administrator)

📝🟢

note

You cannot add a Transporter applet through
Clusters.

🧭➡️

navi…

your AppSec Transporter under the App
column → Configure → Add Connection (on
the redirected AppSec Transporter setup page

9.1.8.2. Set up a Transporter on your VCS 

☑️🟣

prer…

Ensure you have user permissions to onboard
data sources.
Note
The dedicated AppSec Admin role includes the
required permissions.
Ensure you have user permissions to onboard
data sources.
Before you begin you must setup a Transporter
applet on your Broker VM. Refer to Set up a
Transporter applet on Broker VM for more
information

📝🟢

note

GitLab Enterprise webhook internal IP
restriction: GitLab Enterprise's security policies
prevent webhook subscriptions to internal IP
addresses (such as broker addresses). These
subscriptions can only be enabled through
specific instance configuration. For more
information, refer to
https://docs.gitlab.com/security/webhooks/.

🧭➡️

navi…

Settings → Data Sources & Integrations → +
Add New → , and search for your VCS on-
premises data source

🧭➡️ 9.1.9. Manage integrations via data source APIs 

navi…

API key pair: Generate an API key (Key ID + Key
Secret) from Settings → Configurations → API
Keys in the Cortex Cloud tenant

9.2. Cortex Cloud Application Security
dashboard 🧭➡️ 
navi…

Modules → Application Security → AppSec
Dashboard

9.3. Application Security Posture Management
(ASPM) 

📝🟢

note

ASPM does not execute scans. Detection and
finding generation belong to Code Security
(native scanners) and the Supply Chain Security
pillar (VCS repository and CI/CD pipeline
analysis). ASPM evaluates the findings against
unified policies and orchestrates the resulting
issues through to remediation.

🧭➡️

navi…

You can view AppSec Admin permissions in the
tenant by navigating to Settings →
Configurations → Roles → AppSec Admin

9.3.1. Code to Cloud 

🧭➡️

navi…

Full Code to Cloud (C2C) coverage is achieved
when Cortex Cloud can resolve the following
chain: Repository → Pipeline → Image →
(optional Registry) → Runtime resources
(including VMs, VM images, and IaC-defined
infrastructure

9.3.1.1. Supported integrations 

🧭➡️ 9.3.1.2. Code to Cloud context and visibility 

navi…

Topology graph: Located in the Business
Applications side card, this tab visualizes the
entire path to production (Code → Build →
Deploy → Run), allowing you

📝🟢 9.3.1.3. Code to Cloud troubleshooting 

note

Empty state messages and troubleshooting
guidance for missing lineage are only visible to
users with an active license.

9.3.2.1. Operational workflows  9.3.2. ASPM Command Center 

9.3.3. Applications 

9.3.3.1. Defining Business Applications 

9.3.3.1.1. Defining business applications by
Criteria 

🧭➡️ 9.3.3.1.2. Manage Criteria via the tenant UI 

navi…

Modules → Application Security → Application
Management → Criteria

9.3.3.1.3. Define applications by VCS Criteria 

☑️🟣

prer…

Data source: Your Version Control System (such
as GitHub, GitLab) must already be onboarded
as a Data Source
Permissions: You must have View/Edit
permissions for Access Management
SBAC: You can only create applications from
VCS entities (Organizations, Projects, or
Repositories) that are already included in your
SBAC Asset Groups

📝🟢

note

This workflow allows you to unify assets across
disparate providers (such as grouping a GitHub
repository and a Bitbucket repository into one
application) if they share naming conventions

9.3.3.1.4. Define applications by cloud tag-based
Criteria 

☑️🟣

prer…

Permissions: You must have View/Edit
permissions for Access Management, or a role
that includes these permissions.

📝🟢

note

You can only create applications based on the
entities from onboarded Cloud accounts listed
in the Cloud card.

📝🟢

note

Kubernetes (K8S) labels: K8S labels are
supported as tags for asset grouping only when
they originate from a supported cloud provider;
AWS, GCP, Azure, or OCI.

9.3.3.1.5. Manage criteria via the public API 

9.3.3.1.6. How to manually build an application 

☑️🟣

prer…

Before you begin, ensure you have connected
the necessary data sources. Refer to How to
onboard data sources for more
information.How to onboard data sources

📝🟢

note

These filters are represented by icons displayed
on the Code pane after selecting a VCS.

📝🟢

note

To edit application assets, click the Clear All
icon before clicking Finish. This clears all
application data, allowing you to restart the
application building process from the beginning.

🧭➡️ 9.3.3.2. Application management and visibility 

navi…

Inventory → All Assets → Business Applications
(under Application

📝🟢 9.3.3.3. Manage applications via public APIs 

note

The application public API does not execute
scans, enforce policies, or create issues.
Scanning, policy enforcement, and issue
creation are performed by their respective
subsystems (Unified Application Security
policies, scanner orchestration).

9.3.3.4. Business application assets 

9.3.3.4.1. Business application expanded asset
details 📝🟢 
note

The topology graph is available only when all
application components (code, pipeline, build
and deploy), are configured.

9.3.3.4.2. Export business application data 

9.3.3.5. Scope user access to applications
(Application SBAC) 

🧭➡️

navi…

Configure user scopes in Cortex Cloud by
navigating to Settings → Configurations →
Access Management

9.3.3.5.1. Enable SBAC in the Cortex Cloud tenant 

☑️🟣

prer…

RBAC permissions: To configure user scopes
you must have Administrator or View/Edit RBAC
permissions for Access Management (under
Configurations).

📝🟢

note

Exclusions (roles not governed by SBAC):
Certain roles cannot have SBAC applied. For
these roles, access and permissions are
managed through Role-Based Access Control
(RBAC). You must manually ensure that these
roles have all necessary base permissions (for
example Edit/View permissions to assets),
because SBAC is bypassed and does not
impose its usual...

🧭➡️

navi…

Settings → Configurations → General → Server
Settings → Enable Scope Based Access Control

9.3.3.5.2. Create an application-based Asset
Group

📝🟢 

note

You cannot create SBAC based on static
groups.

9.3.3.5.3. Scope user access to an application 

📝🟢

note

For more information about User Groups, refer
to User group management.

🧭➡️

navi…

Settings → Configurations → User Groups
(under Access Management

Right-click on a group in the table → Edit
Group → select the Scope tab

Assets → Select asset groups → select an
Asset Group associated with applications

Settings → Configurations → Users (under
Access Management

the chevron icon (>) in the Assets field →
Select Asset groups → select the user group
scoped to the application (see above

9.3.3.5.4. Create application-scoped policies 

📝🟢

note

Application-scoped policies apply to both
code and CI/CD configuration policies.

🧭➡️

navi…

Modules → Application Security → AppSec
Polices (under Policy Management) → Add
Policy

Add Filter → Business Application Names →
enter the required application name

9.3.3.5.5. Terraform workflow for Asset Groups 

9.3.4. Repository as an asset 

🧭➡️ 9.3.4.1. Understanding repository assets via the UI 

navi…

To access repository assets, under Inventory,
select All Assets → Code → Repositories

🧭➡️ 9.3.4.2. Investigate repository assets 

navi…

From the Repositories table, right-click a
repository → Open in Agentic Assistant →
select Application Security from the agents
menu, and query repository-specific insights
(for example, scan coverage, risk posture, or
gaps

9.3.4.3. Manage repositories via API 

📝🟢

note

The public API base path is
/public_api/appsec/v1/repositories.

📝🟢

note

The excludedPaths field replaces the entire
excluded paths list; to add a new path, retrieve
the current configuration first, append the new
path, and send the complete array
SAST: SAST scanner configuration is not
available through the public API and must be
managed through the console UI

⛔🔴

war…

Always retrieve the current branch list before
updating. The PUT operation replaces the entire
branch selection, meaning any omitted
branches will be unintentionally removed from
scanning.

🧭➡️ 9.3.5.1. Coverage in the user interface  9.3.5. Coverage 

navi…

Modules, → Application Security → AppSec
Coverage

9.3.6. Urgency 

👀🟠 9.3.6.1. Urgency and code to cloud traceability 

imp…

Without code to cloud traceability,
deployment-related urgency signals cannot be
computed. Scanner-specific signals (EPSS
Score, KEV Status, Reachability, Secret
Validation, CWE Top 25, OWASP Top 10) are
always available regardless of code to cloud
traceability.

9.3.6.2. Prioritize issues by Urgency levels 

9.3.6.3. Urgency metrics 

9.3.6.4. View Urgency in the tenant 

📝🟢

note

EPSS Score, CVSS Score, Exploit Maturity,
Exploit Availability, Fixable, Is Deployed,
Visibility, Validation metrics are displayed even
when the value is false or zero. Other metrics
are displayed only when the value is present
and non-null
For more about Urgency metrics refer to
Urgency metrics

💡🔵

tip

Combine Urgency filters with scanner type
filters to focus on the highest-risk issues for a
specific scanner category (for example, Top
Urgent vulnerability issues or Urgent secrets
issues).

9.3.7. Backlog baseline 

9.3.7.1. Backlog use cases 

9.3.7.2. Issue/Finding classification by scanner 

📝🟢

note

In some cases/vendors, this is not accurate as
findings are deleted every time new findings are
uploaded. In such cases, the feature may not be
accurate or supported
For SAST, the vendor does not support policy in
pull requests

📝🟢

note

Scanner updates and new detections: When a
security scanner is updated to support new
languages, detection rules, or capabilities, any
issues discovered by these new features for
existing code are classified as part of the
backlog.

🧭➡️ 9.3.7.3. Using Backlog 

navi…

Settings → Configuration → Application
Security → AppSec Issues Configurations and
enable Are SBOM issues considered new
Permissions for all other capabilities, such as
viewing issues or applying policies, are defined
by the existing RBAC/SBAC policies and the
user's specific issue management capabilities

Settings → Configuration → Application
Security → AppSec Issues Configurations and
enable Are SBOM issues considered new

9.3.8. Service Lead Agreements (SLA) 

9.3.8.1. Configure and monitor Cortex Cloud
Application Security SLAs

🧭➡️ 

navi…

Settings → Configurations → Application
Configuration (under Application Security

9.3.9. Compliance for Cortex Cloud Application
Security 

9.3.9.1. Infrastructure-as-Code (IaC) compliance 

9.3.9.2. Manage IaC compliance 

☑️🟣

prer…

Add these properties to the IaC Rules table
through the Table Settings Menu, as they are
not exposed by default.

🧭➡️

navi…

Modules → Application Security → AppSec
Polices → Add Policy

9.3.9.3.1. Create CI/CD compliance reports  9.3.9.3. CI/CD Compliance 

📝🟢

note

The CI/CD module supports GitHub and GitLab
provider types.

📝🟢

note

For more information about about Asset
Groups, refer to Asset groups.

📝🟢

note

CIS GitLab Benchmark, CIS GitHub Benchmark,
and the OWASP Top 10 CI/CD Risks standards
are supported.

📝🟢

note

For more information about assessment
profiles, refer to Use an assessment profile to
run compliance checks on your assets.

🧭➡️

navi…

Posture Management → Compliance → Reports

9.3.9.4. Terraform workflow for Compliance
assessments 

9.3.10. Unified Application Security policies 

9.3.10.1. Core concepts 

📝🟢

note

The Policy Type selection determines the
available configuration options for conditions,
filters, and triggers.

📝🟢

note

For a full description of each finding type and
its applicability (Code vs. Image), refer to
Reference A: Finding type details.

📝🟢

note

For a complete list of all available filters by
category, as well as logical groups, refer to
Reference B: Condition filters and logic.

📝🟢

note

To understand the specific calculation logic
and how grace periods align with your
organization's SLAs, refer to Reference E: Grace
period logic and configuration.

📝🟢

note

Application scoping is currently not supported
for Build Images

📝🟢

note

Images do not have a default severity. You must
select a severity to continue when any image
trigger is selected.

9.3.10.2. Tenant (UI) workflow 

📝🟢

note

SBAC restricts table visibility. The policy table
displays only policies scoped to assets within
the user's assigned application boundary.
Policies scoped to assets outside the user's
SBAC scope are not visible

🧭➡️

navi…

Posture Management → Rules & Policies →
Policy → Application Security The policy table
displays all policies (out-of-the-box (OOTB)
and custom

Posture Management → Rules & Policies →
Policy → Application Security

📝🟢 9.3.10.3. Create a policy 

note

To create a new asset group or application,
select the corresponding link in the scope step.
The link opens the asset group or application
configuration in a new tab

9.3.10.4. View and manage policy details 

⚠️🟡

caut…

Deleting a policy closes all issues created by
that policy during the next periodic scan. This
action cannot be undone. Consider disabling
the policy instead of deleting the policy to
preserve issue history.

🧭➡️

navi…

Modules → Application Security → Issues → All

Modules → Application Security → Scans

📝🟢 9.3.10.5. API workflow 

note

The public API remains backward compatible
with existing API consumers. The API returns
additional fields for the unified policy structure,
including the new policy type, trigger
configuration, and Override Severity settings.

9.3.10.6. Cortex CLI workflow 

9.3.10.7. IDE workflow 

9.3.10.8. Terraform workflow for policies 

9.3.10.9. Reference A: Finding type details 

📝🟢 9.3.10.10. Reference B: Condition filters and logic 

note

When multiple finding types are selected within
a single condition group, only the Common
filters (such as Severity or Backlog Status) are
available for selection.

9.3.10.11. Reference C: Scope mapping details 

9.3.10.12. Reference D: Trigger and actions
mapping 

9.3.10.13. Reference E: Grace period logic and
configuration 

9.3.10.14. Reference F: Engine evaluation and
Urgency logic 📝🟢 
note

If code to cloud traceability is absent (the
repository is not linked to a relevant
application), the engine cannot compute
deployment signals, and Urgency will display as
Not Applicable

9.3.10.15. Reference G: Finding type to trigger
mapping 

📝🟢

note

Malware findings are applicable to image scans
only When a CI Scan includes image scanning,
Malware findings are evaluated PR Scan and
Periodic Scan do not evaluate Malware findings

📝🟢

note

CI/CD Configuration Scanners policies support
only the Periodic Scan trigger PR Scan and CI
Scan triggers are not available for the CI/CD
Configuration Scanners policy type

📝🟢

note

Drift Detection Scanner policies support only
the Periodic Scan trigger PR Scan and CI Scan
triggers are not available for the Drift Detection
Scanner policy type

9.3.10.16. Reference H: Action availability by
trigger 

🧭➡️ 9.3.10.17. Reference I: Audit logging 

navi…

Settings → Configurations → Management →
Audit Logs to review policy audit entries

9.3.11. Application Security Rules 

📝🟢

note

The AppSec rules API manages rule definitions
and metadata. It does not execute scans, create
issues, or enforce policies.

📝🟢 9.3.11.1. Roles and permissions 

note

Custom roles can be configured with either
View or View/Edit permissions for AppSec rules
API keys inherit the permissions of the role
assigned during key generation. To perform
write operations (create, modify, delete)
through the AppSec rules API, the API key must
have a role that includes View/Edit permissions

🧭➡️ 9.3.11.2. Rules inventory 

navi…

To filter rules relating to Secrets, select filter
icon → Scanner (from the Select field) →
Secrets (from the Value field

9.3.11.3. Create custom Cortex Cloud

Application Security rules 📝🟢 
note

See Configure YAML file properties below for
more details.

9.3.11.4. Configure YAML file properties 

9.3.11.5. Terraform workflows for rules 

9.3.11.6. Manage custom rules 

9.3.12. Manage code weakness issues 

9.3.12.1. Navigate to SAST code weakness issues 

📝🟢

note

The default sort order is by Severity
(descending). The default filter shows issues
with a status of New and In Progress. Resolved
issues are hidden by default.

📝🟢

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a unified
policy that matches the finding pattern to
generate an actionable issue in the Issues tab.

🧭➡️

navi…

To access SAST code violation issues, under
Modules, select Application Security → Issues
→ Code Weaknesses

9.3.12.2. Understand the Code Weaknesses table 

9.3.12.3. Investigate and remediate code
weakness issues 

👀🟠

imp…

Important: Urgency signals are populated only
when the SAST asset is traced to deployed
cloud resources through the Code-to-Cloud
mapping. If no Code-to-Cloud trace exists,
urgency signals display as Not Applicable

🧭➡️

navi…

From the main issues table: Right-click on an
issue in the table → Change Status → [Select a
status

From the main issues table: Right-click on an
issue in the table → Change Assignee → [Select
a user

From the main issues table: Right-click on an
issue in the table → Change Severity → [Select
a severity

📝🟢 9.3.12.4. Code weakness findings 

note

Findings in the Findings tab are raw scanner
output. They do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a unified
policy that matches the finding pattern to
generate an actionable issue in the Issues tab.

9.4. Software supply chain security 

9.4.1. Supply Chain assets 

9.4.1.1. VCS organization assets 

🧭➡️ 9.4.1.1.1. Understanding VCS organization assets 

navi…

To access VCS organization assets, under
Inventory, select All Assets → Code → VCS
Organizations

9.4.1.1.2. Investigate VCS organization assets 

9.4.1.1.3. Manage VCS organization assets 

9.4.1.2. VCS collaborators as assets 

9.4.1.2.1. In-depth Collaborator asset information 

9.4.1.2.2. Manage Collaborator assets 

9.4.1.3. Repository as an asset 

9.4.1.3.1. Understanding repository assets via the
UI 

9.4.1.3.2. Investigate repository assets 

9.4.1.3.3. Manage Repository assets 

9.4.1.3.4. Export Software Bill of Materials
(SBOM) 

📝🟢 9.4.1.3.5. Manage issues detected in repositories 

note

You can also find the repository issues in the
general issue inventory table, and in the
dedicated inventory of issues for each scanner
type (see step 4 above for details).

9.4.1.4. Technologies as assets 

📝🟢

note

The filter supports wildcard filtering and is
case-insensitive.

📝🟢 9.4.1.4.1. Manage repository technologies 

note

The filter supports wildcard filtering and is
case-insensitive.

9.4.1.4.2. Troubleshooting and FAQs 

9.4.1.5. Software packages as assets 

9.4.1.5.1. Understanding software package assets 

9.4.1.5.2. Investigate software package assets
and security issues 

9.4.1.6. CI/CD instance as an asset 

🧭➡️ 9.4.1.6.1. Understanding CI/CD instance assets 

navi…

To access CI/CD instance assets, under
Inventory, select All Assets → Code → CI/CD
Instances

9.4.1.6.2. Investigate and manage CI/CD instance
assets 📝🟢 
note

Navigate to the dedicated Application Security
→ Issues → CI/CD Risks page to manage the
CI/CD risks remediation lifecycle at scale
through bulk status updates, team assignments,
and SLA tracking for compliance monitoring.

9.4.1.6.3. Manage CI/CD pipeline instances 

9.4.1.7. CI/CD pipeline as an asset 

🧭➡️ 9.4.1.7.1. Understanding CI/CD pipeline assets 

navi…

To access CI/CD pipeline assets, under
Inventory, select All Assets → Code → CI/CD
Pipelines

9.4.1.7.2. Investigate and manage CI/CD pipeline
assets 

📝🟢

note

This requires active CI/CD integrations and
successful build log analysis. Pipelines without
successful build log analysis display only the
repository and pipeline nodes

📝🟢

note

Navigate to the dedicated Application Security
→ Issues → CI/CD Risks page to manage the
remediation lifecycle at scale through bulk
status updates, team assignments, and SLA
tracking for compliance monitoring.

🧭➡️

navi…

From the CI/CD Pipelines table, right-click a
pipeline row → Open in Agentic Assistant →
Application Security from the agents menu

9.4.1.7.3. Manage CI/CD pipeline assets 

9.4.1.8. Tools as an asset 

📝🟢

note

Although attributes are identical across
inventories, their values for the same tool can
differ, most commonly in Risk Factors and Type.
This is because the inventory reflects your live
environment, which may include different
versions or configurations than the catalog—for
example, a package may not have been
upgraded or may be deployed differently.

🧭➡️

navi…

In the inventory table, right-click on a tool →
Change Status → select a status

9.4.1.8.1. Active tools 

🧭➡️ 9.4.1.8.1.1. Supply Chain Tools 

navi…

To access Supply Chain Tools, select Modules
→ Application Security → Supply Chain Tools
(under 3rd Party Tools

9.4.1.8.1.2. Expanded Supply Chain tool
information 📝🟢 
note

AppSec Admin user permissions are required to
perform these actions.

🧭➡️ 9.4.1.8.2. Supply chain catalog 

navi…

To access the Supply Chain Catalog, select
Modules → Application Security → Supply
Chain Catalog (under 3rd party tools

9.4.2. Software Composition Analysis (SCA)
scanners



👀🟠

imp…

Utilizing native SCA scanner capabilities
requires an AppSec add-on license (Posture
Management, Runtime, or Premium Cortex).

9.4.2.1. Supported Software Composition
Analysis (SCA) frameworks and languages 

9.4.2.2. Software Composition Analysis (SCA)
vulnerability issues 

9.4.2.2.1. Understand the Vulnerabilities table 

9.4.2.2.2. Investigate and remediate CVE
vulnerability issues 

📝🟢

note

The Open Fix PR action is available only for
CVE vulnerability issues detected through
periodic scans where an automated fix exists.
The action is not available for issues detected
through PR scans. Verify that the Is Fixable
indicator is True before attempting to open a
fix pull request.

👀🟠

imp…

Important: Urgency signals related to
deployment context are populated only when
the SCA asset is traced to deployed cloud
resources through the Code-to-Cloud
mapping. If no Code-to-Cloud trace exists,
deployment-related urgency signals display as
Not Applicable. SCA-specific signals (EPSS
Score, KEV Status, Reachability, Has Fix) are
always available...

9.4.2.2.3. Investigate CVE vulnerabilities findings 

📝🟢

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a unified
policy that matches the finding pattern to
generate an actionable issue in the Issues tab.

🧭➡️

navi…

To access CVE findings, under Modules, select
Application Security → Issues → Vulnerabilities
→ click the Findings tab

9.4.2.2.4. Manage SCA CVE vulnerability issues 

9.4.2.3. License miscompliance issues 

📝🟢 9.4.2.3.1. Understand the Licenses table 

note

Important: License miscompliance issues do not
have urgency signals. Unlike CVE vulnerability
issues, license issues are legal compliance risks
— not security vulnerabilities. Urgency signals
such as EPSS score, KEV status, reachability
analysis, and Code-to-Cloud deployment
context do not apply to license compliance
evaluation.

9.4.2.3.2. Take action on license miscompliance
issues 👀🟠 
imp…

Automated fix pull requests are not available
for license miscompliance issues. Remediation
requires human judgment to evaluate license
compatibility, negotiate licensing terms, identify
alternative packages, or accept the license with
documented justification.

🧭➡️ 9.4.2.3.3. License miscompliance findings 

navi…

To access license miscompliance findings,
under Modules, select Application Security →
Issues → Licenses → click the Findings tab

9.4.2.3.4. Open-source software license
categories 

9.4.2.3.5. Manage license miscompliance issues 

🧭➡️ 9.4.2.4. Ingest third-party SCA data 

navi…

Modules → Application Security → Issues →
Vulnerabilities

9.4.3. CI/CD Risks 

📝🟢

note

Cortex Cloud Application Security CI/CD
pipeline scans create a comprehensive
inventory of all CI/CD pipelines in your
environment. For more information refer to
CI/CD pipeline as an asset.

🧭➡️ 9.4.3.1. CI/CD pipeline issues 

navi…

To access CI/CD pipeline risks issues, under
Modules, select Application Security → Issues
→ CI/CD Risks

📝🟢 9.4.3.2. Expanded CI/CD risks issue information 

note

Different issue types include different
properties; therefore, not all properties are
available for every issue.

9.4.3.3. VCS and CI/CD pipeline risk findings 

9.4.4. CI/CD Policies 

💡🔵

tip

For Cortex Cloud Code policies, refer to
Unified Application Security policies.

9.4.4.1. CI/CD policies user roles and permissions 

📝🟢 9.4.4.2. CI/CD policies inventory 

note

To view all out-of-the-box (OOTB) policies,
filter by Policy Owner = System.

9.4.4.3. Create CI/CD configuration policies 

📝🟢

note

When you select a category (such as
Application), a list of available asset values for
that type is displayed. The policy will
automatically apply to all assets in the
displayed list unless you apply additional filters
using specific asset attributes (such as the
Business Application Names filter).

📝🟢

note

SBAC scope-based limitations do not apply to
Asset Types

📝🟢

note

Periodic scan is the only trigger that is
supported for CI/CD policies.

9.4.4.3.1. Cortex Cloud Application Security
CI/CD policy Condition attributes 

9.4.4.4. Manage CI/CD policies 

📝🟢

note

You cannot edit out-of-the-box (OOTB)
policies.

📝🟢

note

The duplicated policy will include the word
"clone" in its name and must be renamed.

9.4.5. CI/CD Rules 

📝🟢

note

Out-of-the-box rules cannot be modified
Custom CI/CD rules are not supported

9.4.5.1. CI/CD rules roles and permissions 

9.4.5.2. CI/CD rules inventory 

9.5. Code Security 

📝🟢

note

Code Security covers detection and analysis
engines. For governance and orchestration
layer, see Application Security Posture
Management (ASPM).

🧭➡️

navi…

You can view the granular permissions assigned
to each role in the tenant by navigating to
Settings → Configurations → Roles

9.5.1. Code Security assets 

9.5.2. Software packages as assets 

9.5.2.1. Supported Software Composition
Analysis (SCA) frameworks and languages 

🧭➡️ 9.5.2.2. Understanding software package assets 

navi…

To access software package assets, under
Inventory, select All Assets → Code →
Software Packages

9.5.2.3. Investigate software package assets and
security issues 

9.5.3. Infrastructure-as-Code (IaC) resources as
assets 

9.5.3.1. Supported frameworks and languages 

🧭➡️ 9.5.3.2. Access and filter IaC assets 

navi…

Inventory, select All Assets → Code → IaC
Resources

9.5.3.3. Investigate IaC assets 

9.5.4. Code Security scanners 

9.5.5. Secrets scans 

9.5.5.1. Navigate to secrets issues 

📝🟢

note

The default sort order is by Severity
(descending). The default filter shows issues
with a status of New and In Progress. Resolved
issues are hidden by default.

🧭➡️

navi…

Under Modules, select Application Security →
Issues → Secrets

9.5.5.2. Understand the secrets issues table 

👀🟠 9.5.5.3. Investigate and remediate secrets issues 

imp…

Urgency signals related to deployment context
are populated only when the secrets asset is
traced to deployed cloud resources through
the Code-to-Cloud mapping.

📝🟢 9.5.5.4. Secrets findings 

note

Review excluded findings: Investigate findings
that were excluded by policy filters to confirm
that exclusions are intentional and do not
suppress critical IaC misconfigurations.

9.5.6. Infrastructure as Code (IaC)

misconfiguration scanner 

9.5.6.1. Navigate to IaC misconfiguration issues 

📝🟢

note

The default sort order is by Severity
(descending). The default filter shows issues
with a status of New and In Progress. Resolved
issues are hidden by default.

🧭➡️

navi…

To access IaC misconfiguration issues, under
Modules, select Application Security → Issues
→ IaC misconfiguration

9.5.6.2. Understand the IaC misconfigurations
table 

9.5.6.3. Investigate and remediate IaC
misconfiguration issues 📝🟢 
note

The Open Fix PR action is available only for
issues detected through periodic scans where
an automated fix exists. The action is not
available for issues detected through PR scans.

9.5.6.4. IaC misconfigurationfindings 

📝🟢

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a unified
policy that matches the finding pattern to
generate an actionable issue in the Issues tab.

🧭➡️

navi…

To access IaC misconfiguration findings, under
Modules, select Application Security → Issues
→ IaC Misconfigurations → click the Findings
tab

The Cortex Cloud Application Security IaC
misconfigurations Findings table is a filtered
view of the comprehensive Findings inventory,
located at Cases & Issues → Issues → Findings
tab

9.5.7. IaC Drift Detection scans 

9.5.7.1. Navigate to IaC drift detection issues 

📝🟢

note

The default sort order is by Severity
(descending). The default filter shows drift
findings with a status of New and In Progress.
Resolved drift findings are hidden by default.

🧭➡️

navi…

Modules → Application Security → IaC Drifts

9.5.7.2. Understand the IaC drift detection table 

9.5.7.3. Investigate and remediate IaC drift issues 

📝🟢

note

The Revert to IaC State action requires write
permissions on the target cloud account. Verify
that the Cortex Cloud service account has the
necessary IAM permissions before initiating a
revert operation.

📝🟢

note

Updating the IaC template codifies the current
cloud state as the new desired state. Verify that
the cloud resource configuration complies with
security policies before merging the generated
pull request.

🧭➡️

navi…

From the main drift table:Right-click on a drift
issue in the table → Change Status → Change
Status

From the main drift table:Right-click on a drift
issue in the table → Change Assignee →
Change Status

9.5.7.4. Investigate IaC drift detection findings 

📝🟢

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific drift finding, create or update a
unified policy that matches the finding pattern
to generate an actionable issue in the Issues tab.

🧭➡️

navi…

Modules → Application Security → IaC Drifts
(under Issues) → click the Findings tab

9.5.8. Software Composition Analysis (SCA)
scanners



9.5.8.1. Software Composition Analysis (SCA)
vulnerability issues 

9.5.8.1.1. Navigate to CVE vulnerability issues 

📝🟢

note

The default sort order is by Severity
(descending). The default filter shows issues
with a status of New and In Progress. Resolved
issues are hidden by default.

📝🟢

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a policy
that matches the finding pattern to generate an
actionable issue in the Issues tab.

9.5.8.1.2. Understand the Vulnerabilities table 

9.5.8.1.3. Investigate and remediate CVE
vulnerability issues 

📝🟢

note

The Open Fix PR action is available only for
CVE vulnerability issues detected through
periodic scans where an automated fix exists.
The action is not available for issues detected
through PR scans. Verify that the Is Fixable
indicator is True before attempting to open a
fix pull request.

👀🟠

imp…

Important: Urgency signals related to
deployment context are populated only when
the SCA asset is traced to deployed cloud
resources through the Code-to-Cloud
mapping. If no Code-to-Cloud trace exists,
deployment-related urgency signals display as
Not Applicable. SCA-specific signals (EPSS
Score, KEV Status, Reachability, Has Fix) are
always available...

📝🟢 9.5.8.1.4. Investigate CVE vulnerabilities findings 

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a unified
policy that matches the finding pattern to
generate an actionable issue in the Issues tab.

9.5.8.2. License miscompliance issues 

9.5.8.2.1. Navigate to license miscompliance
issues 

📝🟢

note

The default sort order is by Severity
(descending). The default filter shows issues
with a status of New and In Progress. Resolved
issues are hidden by default.

📝🟢

note

Findings in the Findings tab are raw scanner
output and do not have resolution statuses,
SLA tracking, or assignees. To track remediation
for a specific finding, create or update a unified
policy that matches the finding pattern to
generate an actionable issue in the Issues tab.

🧭➡️

navi…

Application Security → Issues → Licenses

📝🟢 9.5.8.2.2. Understand the Licenses table 

note

Important: License miscompliance issues do not
have urgency signals. Unlike CVE vulnerability
issues, license issues are legal compliance risks
— not security vulnerabilities. Urgency signals
such as EPSS score, KEV status, reachability
analysis, and Code-to-Cloud deployment
context do not apply to license compliance
evaluation.

9.5.8.2.3. Take action on license miscompliance
issues 👀🟠 
imp…

Automated fix pull requests are not available
for license miscompliance issues. Remediation
requires human judgment to evaluate license
compatibility, negotiate licensing terms, identify
alternative packages, or accept the license with
documented justification.

9.5.8.2.4. How license miscompliance issues fit in
the Application Security ecosystem 🧭➡️ 
navi…

Configure unified policies for license
compliance at Modules → Application Security
→ AppSec Policies

9.5.8.2.5. Open-source software license
categories 

9.5.8.3. Package operational risk scanner 

9.5.8.3.1. Understand the package operational
risk table 

9.5.8.3.2. Investigate and remediate package
operational risk issues 

📝🟢

note

Automated fix pull requests are not available
for package operational risk issues. Remediation
requires human judgment to evaluate
alternative packages, assess API compatibility,
and perform the migration.

👀🟠

imp…

Package operational risk issues do not have
urgency signals based on exploit intelligence.
Unlike CVE vulnerability issues, operational risk
issues assess the health and sustainability of a
package, not active security vulnerabilities.
Urgency signals such as EPSS score, KEV status,
and reachability analysis do not apply to
operational risk evaluation....

9.5.8.3.3. How package operational risk issues fit
in the Application Security ecosystem 🧭➡️ 
navi…

Configure unified policies for operational risk at
Modules → Application Security → AppSec
Policies

9.5.9. API workflows for Code Security issues 

📝🟢

note

When blockingPoliciesIds is non-empty, the PR
was blocked, preventing vulnerable code from
merging into the protected branch.

📝🟢

note

The severity query parameter is the only server-
side filter available on the issue endpoint.
Filtering by detection method, resolution status,
or assignee requires client-side processing.

📝🟢

note

License compliance (License Security) and
package operational risk (Operational Risk
Security) are not represented as separate keys
in issuesTypesBreakdown. To retrieve these
issue types, filter client-side by
detection_method from the full issue set.

📝🟢

note

Application issue summaries return counts only,
not individual issue records. To retrieve
individual records, list periodic scans for the
application's repositories, then use GET
/scans/{scanId}/issues.

💡🔵

tip

Use the issuesTypesBreakdown and issuesFound
fields in the scan list response to identify scans
with Critical or High issues before drilling into
individual records, thereby reducing API calls.

🧭➡️

navi…

Generate an API key pair from Settings →
Configurations → API Keys in the Cortex Cloud
console

9.5.10. Application Security scans management 

📝🟢 9.5.10.1. Manage scans through the tenant (UI) 

note

While scans provide a comprehensive inventory
of all issues detected during a scan, dedicated
inventories are also maintained for specific
scan types for more granular management. For
more information, refer to Code Security
scanners.

📝🟢 9.5.10.2. Branch periodic scans 

note

The inventory table displays scan issues for
visibility only; remediation is not available here.
To resolve issues, navigate to the dedicated
issue type inventory, where you can manage
and remediate them.

📝🟢 9.5.10.3. Pull Request scans 

note

The inventory table displays scan issues for
visibility only; remediation is not available here.
To resolve issues, navigate to the dedicated
issue type inventory, where you can manage
and remediate them.

📝🟢 9.5.10.4. CI scans 

note

The inventory table displays scan issues for
visibility only; remediation is not available in
scan management. To resolve issues, navigate
to the dedicated issue type inventory, where
you can manage and remediate them.

📝🟢 9.5.10.5. Scan health and status reference 

note

Periodic scans do not have a scan status.
Periodic scans produce issues and findings but
do not evaluate blocking policies. The scan
health status is the only health indicator for
periodic scans.

9.5.10.6. Manage repository scan configurations 

9.5.10.7. Monitor data source instances health 

9.5.10.8. Manage scans through public APIs 

9.5.11. Application Security CLI 

9.5.12. IDE 

📝🟢

note

Not all remediation options are available for all
findings or all type of scan category.

9.5.12.1. System requirements 

9.5.12.2. Visual Studio (VS) Code and VS Code

compatible IDEs 

☑️🟣

prer…

Before you begin (These prerequisites apply to
both VS Code and VS Code compatibles):
Permissions: CLI Read only permissions. Refer to
Cortex CLI for more information about
permissions
Environment setup
macOS and Windows: Install Python 3.9.x to
3.12.x
Install Node.js version 22 and above for SCA
scans (such as vulnerabilities scans)
Environment...

☑️🟣

prer…

You must insert your API key and API ID values
into the Settings before providing the tenant
URL.

📝🟢

note

Enforcement rules and CA certificates are not
applicable to the open-source project.

👀🟠

imp…

When generating an API key, ensure you select
the Standard security level. IDE and CLI scans
will fail if the security level of the API key is set
to Advanced.

💡🔵

tip

You can access Extensions using the keyboard
shortcuts Ctrl+Shift+X (Windows) or
Cmd+Shift+X (macOS).

🧭➡️

navi…

Settings → Configurations → API Keys (under
Integration) → click Copy API URL

Install through VS Code IDE Access the Activity
bar → Extensions → Cortex Cloud → Install

Access the Activity bar → Extensions → Cortex
Cloud → Install

9.5.12.2.1. How to use the Cortex Cloud

extension in VS Code 

📝🟢

note

Not all remediation options are available for all
issues.

📝🟢

note

The suppression is scoped to the file.

9.5.12.3. JetBrains 

☑️🟣

prer…

Permissions: CLI Read only permissions. Refer to
Cortex CLI for more information about
permissions
Environment setup
macOS and Windows: Install Python 3.9.x to
3.12.x
Install Node.js version 22 and above for SCA
scans (such as vulnerabilities scans)
Environment setup
macOS and Windows: Install Python 3.9.x to
3.12.x
Install Node.js version 22 and...

☑️🟣

prer…

You must insert your API key and API ID values
into the Settings before providing the tenant
URL.

📝🟢

note

The Cortex AppSec code security plugin
supports all JetBrains products.

👀🟠

imp…

When generating an API key, ensure you select
the Standard security level. IDE and CLI scans
will fail if the security level of the API key is set
to Advanced.

🧭➡️

navi…

Settings → Plugins → select the Marketplace
tab → search for Prisma Cloud → Install → OK

Settings → Tools → Cortex Cloud

9.5.12.3.1. How to use the JetBrains Cortex Cloud

extension 

📝🟢

note

After selecting Fix Available, the number of
issues displayed in the issue categories (such as
IaC) reflect the number of fixable issues for
that type.

📝🟢

note

The corresponding file simultaneously opens in
the Code editor, highlighting the issue within its
exact code context. See below for more
information.

📝🟢

note

Not all types of remediation are available for all
issue categories. For example, fixes are not
available for License issues.

📝🟢

note

Secrets and Licenses category issues are
typically mitigated by following the guidance in
the Documentation.

📝🟢 9.5.13. Developer Suppressions 

note

The check IDs and rule names are placeholders
for illustration purposes and do not correspond
to actual checks.

9.6. API endpoints 

9.7. Terraform workflows 

🧭➡️

navi…

To authenticate, you must generate a
dedicated API key in the Cortex Cloud console
under Settings → Configurations → API Keys

📝🟢 9.7.1. Manage resources 

note

For a complete list of supported IaC and CI/CD
compliance standards, as well as rule mapping
logic and available scan types, refer to
Compliance for Cortex Cloud Application
Security.

10. Cortex Cloud Data Classification 

🧭➡️ 10.1. What is Cortex Cloud Data Classification? 

navi…

To access Data Classification management,
click Settings → Configurations → Data
Classification

10.2. How to create and validate a custom data

pattern 

📝🟢

note

Once a custom data pattern is saved, it runs on
all data in the same way as any out-of-the-box
(OOTB) pattern, becoming globally applicable
for all modules using Cortex Cloud Data
Classification.

📝🟢

note

For more information, see How to disable and
enable data patterns in Data Classification.How
to disable and enable data patterns in Data
Classification

10.2.1. Custom data patterns: Guardrails and
syntax guide 

10.3. How to disable and enable data patterns in
Data Classification 

📝🟢

note

Disabling data patterns can cause changes in
your data profile results and stop detection of
these data patterns.

📝🟢

note

For more information about data patterns in
data classification, see What is Cortex Cloud
Data Classification?.

🧭➡️

navi…

In the lower left part of the screen, click
Settings → Configurations → Data
Classification → Data Patterns

10.4. How to create and validate a custom data

profile 

📝🟢

note

If you choose two data locations, only the
filters they have in common will be included in
the possible filter options.

📝🟢

note

For more information, see How to disable and
enable data profiles in Cortex Cloud Data
Classification.

👀🟠

imp…

Deleting a data profile deletes all past data
associated with it in all modules using Cortex
Cloud Data Classification after a warning
notification is displayed.

👀🟠

imp…

Enabling and disabling a data profile removes
or re-adds the data profile results to the data
objects; that is, files and tables.

10.5. How to disable and enable data profiles in
Cortex Cloud Data Classification 

📝🟢

note

Existing results on past scans are updated in
the Asset and Object inventories within two
hours after being disabled or enabled.

📝🟢

note

For more information about data profiles in
Cortex Cloud Data Classification, see What is
Cortex Cloud Data Classification?.What is
Cortex Cloud Data Classification?

🧭➡️

navi…

In the lower left part of the screen, click
Settings → Configurations → Data
Classification → Data Profiles

10.6. How to report a false positive in Cortex
Cloud Data Classification 📝🟢 
note

Alternatively, in Cortex Command Center, click
Help → Submit a Support Case.

11. Cortex Cloud Data Security 

11.1. What is Cortex Cloud Data Security? 

11.2. Supported assets in Cortex Cloud Data
Security 

🔒🟣

cont…

This feature is included with a Cortex XSIAM
Premium license. It is also included with any
other Cortex XSIAM license that has the Cloud
Posture Security or Cloud Runtime Security
add-on.

📝🟢

note

CMK (customer-managed key) in SQL Server is
only supported in outpost mode.

📝🟢

note

The list above refers to only data classification;
however, Cortex Cloud Data Security discovers
and monitors all cloud assets and services for
usage and misconfigurations.

📝🟢

note

For information about configuring data
classification asset types, see How to configure
the scanning settings for supported services.

11.3. Cortex Cloud Data Security concepts 

11.4. Cortex Cloud Data Security use cases 

11.5. Data Inventory 

📝🟢

note

To clear the filter in the filter panel, click Delete.

🧭➡️

navi…

When you select Inventory → All Assets →
Data → All Data Assets, a list of all data assets
appears

11.6. How to review errors in Cortex Cloud Data

Security 

11.7. How to configure the scanning settings for
supported services 📝🟢 
note

If you do not select a scanning cadence, the
default setting is applied. For more information,
contact your Customer Support team.

11.8. How to perform advanced Data Security
investigations using XQL 📝🟢 
note

For more information, see Build XQL queries.

📝🟢 11.9. How to onboard Databricks 

note

If you turn on this feature, both the cloud and
region will be used for scanning, possibly
incurring cost and requiring adherence to
certain compliance policies.

☑️🟣 11.10. How to onboard Microsoft 365 

prer…

You have generated a Globally Unique Identifier
(GUID), also known as a Universally Unique
Identifier (UUID). You will need this ID for the
tenant you want to use for the Microsoft 365
instance.
In order to use Microsoft 365, you must be
registered with Microsoft Azure.

11.11. How to onboard on-premise assets to
Cortex Cloud Data Security 

📝🟢

note

Your data is scanned on the Broker VM itself,
and only the metadata and classification
results are transmitted from the on-premise
environment to Cortex Cloud.

📝🟢

note

For information about activating the DSPM
Fileshare applet, see Activate DSPM Fileshare.

📝🟢

note

For information about activating the DSPM
Database applet, see Activate DSPM Database.

11.12. How to onboard Snowflake 

☑️🟣

prer…

In order to use Snowflake, you must be
registered with one of these cloud providers:
Amazon AWS, Microsoft Azure, or Google
Cloud Platform (GCP).
Ensure you have the necessary account
permissions to onboard. It is recommended to
use Account Admin as the role for the
onboarding.

📝🟢

note

The account identifier can be found using the
user information at the bottom left. Hover over
the account you wish to onboard and select
the copy option at the top right. The account
identifier is usually of the format:
(organization).[account]

11.13. How to use information protection labels
in Cortex Cloud Data Security 📝🟢 
note

Only If you have connected a Microsoft 365
instance to Cortex Cloud, the information
protection labels appear. See How to onboard
Microsoft 365.

12. Cortex Cloud Identity Security 

12.1. What is Cortex Cloud Identity Security? 

12.2. Review and improve your Identity Security
posture



📝🟢

note

Least Privilege Access (LPA) is not supported in
OCI.

💡🔵

tip

Reducing the number of non-human entities
that are either dormant or have administrative
privileges is the most effective way to improve
your security score.

12.3. How does Effective Permission Calculation

work? 📝🟢 
note

Unlike GCP (organizations/folders/projects) or
Azure (management groups/subscriptions), OCI
uses a nested compartment structure.

12.4. Cortex Cloud Identity Security
functionality 

📝🟢

note

Access to the Weak Password widget and
column requires the ITDR add-on.

📝🟢

note

Once a user updates their password to a strong
password, they will no longer appear in the
filtered weak password results.

📝🟢

note

For more information about Effective
Permission Calculation, see How does Effective
Permission Calculation work?.

📝🟢

note

When a policy is analyzed and categorized as
excessive, a relevant finding and highlight is
attached to that policy and the various
identities being granted excessive permissions.

📝🟢

note

Each action must be defined as either being
administrative or not.
The administrative tag is an additional tag for
administrative actions, along with one of the
other access level types. For example, the AWS
action iam:CreateGroup is categorized as
config and has the administrative tag as well.

👀🟠

imp…

If you turn off the audit logs, even briefly, this
temporarily impacts the accuracy of the Last
Access data, potentially showing permissions
as unused when actually they were active. Full
accuracy is restored 90 days after you re-
enable the audit logs.

12.5. Unified Human Identities 

12.6. Achieve the principle of least privilege
access



📝🟢

note

In the case of Amazon AWS, Cortex Cloud
Identity Security also uses AWS Identity and
Access Management (IAM) Access Advisor
insights to expand the coverage of supported
actions.

👀🟠

imp…

If you turn off the audit logs, even briefly, this
temporarily impacts the accuracy of the Last
Access data, potentially showing permissions
as unused when actually they were active. Full
accuracy is restored 90 days after you re-
enable the audit logs.

12.7. Explore permissions using the simple and
advanced access tables 

12.8. Create a custom detection rule in Cortex

Cloud Identity Security 🧭➡️ 
navi…

Posture Management → Rules & Policies →
Rules → Cloud Security

12.9. Perform advanced Identity Security
investigations using XQL 📝🟢 
note

For more information, see Build XQL queries.

☑️🟣 12.10. Ingest logs and data from Okta 

prer…

Administrator privileges: Your Okta user must
have a role capable of creating API tokens,
such as Read-only Administrator, Super
Administrator, or Organization Administrator.
For more information, see the Okta
Administrators Documentation.

12.11. Enable inactive human identity logs on
Azure in Cortex Cloud Identity Security 📝🟢 
note

For more information, see Ingest logs from
Microsoft Azure Event Hub.Ingest logs from
Microsoft Azure Event Hub

12.12. Manage RBAC and SBAC in Cortex Cloud
Identity Security 👀🟠 
imp…

You can use the Cortex Cloud Identity Security
RBAC roles to define access to the various
sections and functionalities of Cortex Cloud
Identity Security, but these roles do not directly
control the specific data a user sees within
those sections. Data visibility is further refined
and limited by scope-based access control
(SBAC) capabilities.

13. Cloud ASM 

13.1. What is Cloud ASM? 

13.2. Cloud ASM concepts 

13.2.1. Scanning 

13.2.1.1. Scanning cadences 

📝🟢 13.2.1.2. Scanning ports and protocols 

note

The following lists are not exhaustive. For
current and complete lists, contact your
customer success team.

13.2.1.3. Scanning activity 

13.2.2. Network mapping 

13.3. Enable Cloud ASM 

📝🟢

note

Discovered services will increase workload
consumption. ASM data will begin to appear in
your tenant and be included in your licensing
workload meter within 72 hours.

🧭➡️

navi…

Settings → Configurations → Attack Surface →
Data Management

13.4. Attack Surface Management detections 

🧭➡️ 13.4.1. Attack surface rules 

navi…

Modules → Attack Surface → Policies →
Attack Surface Rules

📝🟢 13.4.2. Externally inferred CVEs 

note

An externally inferred CVE might impact your
service or asset, but additional investigation is
required to confirm that the CVE is actually
present.

🧭➡️ 13.5. Attack surface assets 

navi…

The internet-facing assets that were discovered
in an attack surface management (ASM) scan
and attributed to your organization are
available in the inventory on the Inventory →
Assets → All Assets → External Surface pages

🧭➡️ 13.6. Review your unmanaged cloud services 

navi…

Inventory → Assets → All Assets → External
Surface → Services

13.7. Review unmanaged cloud issues 

14. Network exposure detection 

14.1. What is Cloud Network Analyzer? 

📝🟢 14.2. Internet exposure detection 

note

Currently, CNA supports only an isolated
architecture with an AWS Gateway Load
Balancer (GWLB) within a single VPC. Other
more centralized topologies including one
security VPC that forwards traffic to other
workload VPCs are currently not supported.

14.3. Outbound exposure detection 

14.4. East-west exposure detection 

14.5. Investigate an internet exposure 

📝🟢

note

The Network tab is currently only available for
virtual machines.

📝🟢

note

The Network tab is only displayed when you
have access to the main asset and associated
ones, such as security groups, VPCs and
subnets. For more information on Scope-Based
Access Control (SBAC) for configuring granular
scoping, see Manage user scope.Manage user
scope

📝🟢

note

The Network tab is currently only available for
virtual machines.

📝🟢

note

The Network tab is only displayed when you
have access to the main asset and associated
ones, such as security groups, VPCs and
subnets. For more information on Scope-Based
Access Control (SBAC) for configuring granular
scoping, see Manage user scope.Manage user
scope

🧭➡️ 14.6. Configure trusted IPs 

navi…

Inventory → Network Configuration → Trusted
Networks

15. Vulnerability management 

15.1. Vulnerability management in Cortex Cloud 

🔒🟣

cont…

Requires the Cortex Cloud Posture Security,
Cortex Cloud Runtime Security, Exposure
Management, Cortex XSIAM Premium or ASM
add-on.

📝🟢

note

Cortex Cloud vulnerability management
provides the ability to identify and assess
runtime vulnerabilities in every asset across
traditional IT and cloud environments. For
vulnerabilities detected in your software
development lifecycle through application
security scans, refer to the Cortex Cloud
Application Security documentation.

📝🟢 15.1.1. Cortex Cloud vulnerability concepts 

note

CVE is an acronym for Common Vulnerabilities
and Exposures, which is a list of publicly
disclosed security threats. We often use the
term "CVE" to refer to a vulnerability that has
been a assigned a CVE ID. Cortex Cloud
identifies CVEs and non-CVE vulnerabilities.

15.1.2. Vulnerability Management dashboard 

🧭➡️ 15.2. Cortex Vulnerability Risk Score 

navi…

Posture Management → Vulnerability
Management → Vulnerability Issues

15.3. Vulnerability policies 

🧭➡️ 15.3.1. Create a vulnerability policy 

navi…

Posture Management → Rules & Policies →
Policies → Vulnerability Management

15.3.2. Update the Ignored CVEs, Asset Groups,
and Assets policy



🧭➡️ 15.3.3. Modify a vulnerability policy 

navi…

Posture Management → Rules & Policies →
Vulnerability Management

🧭➡️ 15.3.4. Configure a block grace period 

navi…

Posture Management → Rules & Policies →
Vulnerability Management → Vulnerability
Policies - Prevention

15.3.5. Enable or disable a vulnerability policy 

15.4. Investigate and remediate vulnerabilities 

🧭➡️ 15.4.1. View all Vulnerabilities 

navi…

Posture Management → Vulnerability
Management → Vulnerabilities

15.4.2. View vulnerability issues 

🧭➡️ 15.4.3. View All Vulnerability Findings 

navi…

Posture Management → Vulnerability
Management → Vulnerability Issues and click
the All Vulnerability Findings button

🧭➡️ 15.4.4. View vulnerable assets 

navi…

Posture Management → Vulnerability
Management → Vulnerable Assets

🧭➡️ 15.5. Vulnerability Intelligence 

navi…

Posture Management → Vulnerability
Management → Vulnerability Intelligence

15.6. Recast CVSS scores and CVSS severities 

16. Cloud Security Rules and Policies 

16.1. Cloud security rules and policies 

📝🟢 16.2. Cloud security rules 

note

Findings are only generated for OOTB rules.

📝🟢 16.3. Cloud security policies 

note

If you have the following Scope Based Access
Control (SBAC) settings in place, User Settings
→ Cases and Issues Scope → Select domains
→ Posture, you may encounter a Case
mismatch in Issues/Cases/Findings counts. This
is because the Case count on the Rules page
captures Cases belonging to the Posture
domain. Whereas Platform pages, capture
Issues...

16.4. Create and manage cloud security rules 

🧭➡️

navi…

You can view and manage cloud posture
security rules from the Posture Management →
Rules & Policies → Rules → Cloud Security
page

16.4.1. Create an attack path rule 

16.4.2. Create a configuration rule 

16.4.2.1. Guidelines for creating cloud security
rules 

16.4.2.2. Cloud security rule status for custom
configuration rules 

16.4.3. Create a data rule 

16.4.4. Create an identity rule 

16.4.5. Create a network exposure rule 

16.4.6. Create an AI rule 

16.4.7. View cloud security rule status 

16.4.8. Edit a cloud security rule 

16.4.9. Enable or disable a rule 

16.4.10. Use an existing rule to create a new one 

16.4.11. Delete a custom cloud security rule 

16.5. Create and manage cloud security policies 

🧭➡️ 16.5.1. Create a cloud security policy 

navi…

Posture Management → Rules & Policies →
Policies → Cloud Security

16.5.2. Edit a cloud security policy 

16.5.3. Enable or disable a policy 

16.5.4. Use an existing policy to create a new one 

16.5.5. Delete a custom cloud security policy 

17. Cloud Workload Policies and

Rules 

☑️🟣

prer…

Users need View/Edit RBAC permissions (under
Policies → Compute Policies) or the Instance
Administrator role to view, edit, and modify
Cloud Workload Policies policies.

👀🟠

imp…

Users with SBAC granular scoping (in addition
to the RBAC permissions required for Cloud
Workload Policies) can only view Cloud
Workload Policies, when their access is scoped
to any of the available options: All assets, No
assets, or Select asset groups. For more
information on granular scoping, see Manage
user scope. When no SBAC restriction is...

17.1. How policies and rules work together 

17.2. Cloud Workload Policies 

📝🟢

note

Issues are automatically resolved when the
finding is no longer applicable to the asset or
when the affected asset is removed from the
inventory.
For more details on Prevent and create issues,
see Cloud Workload Preventive Action.

17.2.1.1. Trusted image cloud workload policies  17.2.1. Types of Cloud Workload Policies 

📝🟢

note

Registry scanning is for finding problems that
are objectively issues regardless of context,
organization, or scope. Trust, however, is
subjective. Depending on the scope and other
factors, an issue may or may not be
problematic, and can change over time
depending on the context.

📝🟢

note

Trusted image policies do not override other
policies and their rules, such as malware
policies, misconfiguration policies, and secrets
policies.

⚠️🟡

caut…

The default behavior for the prevent action
when trust criteria are unavailable is Create an
Issue. Be aware that changing this default will
mean an incomplete deployment file can block
a workload.

17.2.2. Cloud Workload Policies page 

📝🟢

note

Keep the following caveats in my mind when
working with Policies:
Instance Administrators are able to view all
facets of policies without restrictions, even if
Scope Based Access Control (SBAC) roles are
in effect. Learn more about SBAC.Manage user
scope

If you've been assigned a custom role with
View/Edit permissions limited by SBAC, you may
not be...

17.2.2.1.1. Show or hide the widget panel  17.2.2.1. Widgets panel 

17.2.2.2. Change the layout of the policies table 

17.2.2.3. Policy Details Panel 

17.2.3. Enable or disable a Cloud Workload
Policy

🧭➡️ 

navi…

Posture Management → Rules & Policies →
Policies → Cloud Workload

17.2.4. Create a Cloud Workload Policy 

📝🟢

note

The Evaluation stage for Misconfiguration
policies is supported only in the Runtime SDLC
stage and is enforceable through the
Kubernetes Admission Controller for clusters
on-boarded using the Posture Management
(KSPM) Connector.

📝🟢

note

Use Create a new Custom Detection Rule to
define and add new custom rules as required.

📝🟢

note

Each rule may support different actions. While
some include both Create an Issue and Prevent
and Create an Issue, others provide only the
Create an Issue option.

📝🟢

note

All non-relevant (non-compute) assets are
automatically excluded from the included asset
list.

📝🟢

note

Currently, dynamic asset groups for Kubernetes
Prevent Policies support only the following
attributes:
Kubernetes Resource Cluster: The cloud
identifier for the cluster.
Kubernetes Resource Namespace: The
namespace where the resource is located.
Kubernetes Resource Labels: The labels
assigned to the resource.
Kubernetes Resource Category: The category...

📝🟢

note

See Cloud Workload Preventive Action, to learn
more about the Prevent action behavior and
prerequisites.
If the Prevent and create an issue action is
selected, the preventive actions in the CI
pipeline will trigger a Fail Pipeline by returning
an exit code of 2 in the CI tool.

📝🟢

note

Choosing this option results in the automatic
selection of all other asset scopes in the list.

📝🟢

note

See Cloud Workload Preventive Action, to learn
more about the Prevent action behavior and
prerequisites.

📝🟢

note

See Cloud Workload Preventive Action, to learn
more about the Prevent action behavior and
prerequisites.

📝🟢

note

Choosing this option results in the automatic
selection of all other asset scopes in the list.

📝🟢

note

See Cloud Workload Preventive Action, to learn
more about the Prevent action behavior and
prerequisites.

17.2.5. Use an existing policy to create a new
Cloud Workload policy



17.2.6. Edit a Cloud Workload Policy 

17.2.7. Delete a Cloud Workload Policy 

17.2.8. Cloud Workload Preventive Action 

17.3. Cloud Workload Rules 

17.3.1. Default (pre-defined) Rules 

17.3.2. Custom (user-defined) Rules 

17.3.3. Cloud Workload Rules page 

📝🟢

note

Keep the following caveats in my mind when
working with Rules:
Instance Administrators are able to view all
facets of Rules without restrictions, even if
Scope Based Access Control (SBAC) roles are
in effect. Learn more about SBAC.Manage user
scope

If you've been assigned a custom role with
View/Edit permissions limited by SBAC, you may
not be able...

17.3.3.1. Filter page results 

17.3.3.2. Change the layout of the rules table 

17.3.3.3. Rule details panel 

17.3.4. Create a new Custom Detection Rule 

📝🟢

note

The custom rule must use the match term (not
deny or others) to function properly.

📝🟢

note

Only users with the following roles can enable
or disable Custom Code Execution:
Account Admin
Instance Administrator
Deployment Admin
Privileged Security Admin

📝🟢

note

Only Custom Detection Rules (not built-in
rules) can be assigned to custom controls.

👀🟠

imp…

The custom Python scripts are intended to be
executed exclusively for compliance checks
and validations. To ensure the scripts are used
properly and no security risks or unintended
changes occur, the system implements the
following restrictions and safeguards:
Only a predefined set of Python libraries and
functions required for compliance checks are...

17.3.5. Use an existing rule to create a new
Custom Detection Rule 🧭➡️ 
navi…

Posture Management → Rules & Policies →
Rules → Cloud Workload

17.3.6. Edit a Custom Detection Rule 

17.3.7. Delete a Custom Detection Rule 

🧭➡️ 18. Base Images Rule 
navi…

Posture Management → Rules & Policies →
Rules → Base Images

Inventory → Assets → All Assets → Compute
→ Container Images

19. Serverless function posture
security 

19.1. Onboard cloud providers for serverless
functions 

📝🟢

note

When scanning serverless functions with layers,
those layers need to be from the same cloud
account.

📝🟢

note

Only functions containing zip files are
supported.

👀🟠

imp…

Cortex supports Google Cloud Functions: 1st
gen and 2nd gen Cloud Functions API.

19.2. Serverless function posture policies 

19.2.1. Manage serverless function policies 

📝🟢

note

You can select multiple types to view all your
serverless function rules across your cloud
providers.

📝🟢

note

Refer to Create serverless function policies for
more information on how to define the steps of
a policy in the wizard.

19.2.2. Create serverless function policies 

📝🟢

note

You can select multiple options.

📝🟢

note

For more information about rules, refer to
Manage serverless function rules.

19.3. Serverless function posture rules 

19.3.1. Manage serverless function rules 

📝🟢

note

You can select multiple types to view all your
serverless function policies across your cloud
providers.

📝🟢

note

Refer to Create serverless function rules for
more information on how to define the steps of
a rule in the wizard.

19.3.2. Create serverless function rules 

19.3.3. Create an attack path rule for serverless
functions 🧭➡️ 
navi…

Under Posture Management, select Rules &
Policies → Cloud Security (under Rules) →
click Create Rule

in as the operator → enter the CVE ID → Search

> or >= as the operator → Severity level (such
as High, Low) → Search

> or >= as the operator → enter a score →
Search

19.3.4. Create a configuration rule for serverless
functions 

19.3.5. Create a network exposure rule for
serverless functions 

🧭➡️ 19.4. Serverless function usage 

navi…

To access serverless function assets, under
Inventory, select All Assets → Compute →
Serverless Functions

Posture Management → Vulnerability
Management) → Vulnerability Issues

Add Filters → Asset Category → Serverless
Function

Posture Management → Vulnerability
Management) → Vulnerable Assets

Add Filters → Asset Category → Serverless
Functions

20. Data management 

20.1. Broker VM 

📝🟢 20.1.1. What is the Broker VM? 

note

The Broker VM is a closed, hardened appliance.
To maintain its security integrity and
performance standards, third-party agents
cannot be installed on the Broker VM.

20.1.2. Set up and configure Broker VM 

📝🟢

note

The Broker VM comes with a 512 GB disk.
Therefore, deploy the Broker VM with thin
provisioning, meaning the hard disk can grow up
to 512 GB but will do so only if needed.

📝🟢

note

When DHCP is not enabled in your network and
there isn't an IP address for your Broker VM,
configure the Broker VM with a static IP using
the serial console menu.

📝🟢

note

When configuring a Local Agent Settings
applet with installer and content caching, you
need to configure an SSL certificate for the
Broker VM as explained in the task below. For
more information on specific requirements for
the Local Agent Settings applet, see Activate
Local Agent Settings.Activate Local Agent
Settings
Keep in mind that several Broker...

📝🟢

note

When a Broker VM is disconnected for more
than 30 days, it will have to go through a re-
registration process.

📝🟢

note

When DHCP is not enabled in your network and
there isn't an IP address for your Broker VM,
configure the Broker VM with a static IP using
the serial console menu.

📝🟢

note

You can configure another Broker VM as a
proxy server for this Broker VM by selecting the
HTTP type. When selecting HTTP to route
Broker VM communication, you need to add the
IP Address and Port number (set when
activating the Agent Proxy) for another Broker
VM registered in your tenant. This designates
the other Broker VM as a proxy for this Broker
VM.

📝🟢

note

The Palo Alto Networks Broker VM supports
only strong cipher SHA256-based certificates.
MD5/SHA1-based certificates are not
supported.

📝🟢

note

If adding a CA certificate to the Broker VM is
not possible, ensure that you've added the
Broker Service FQDNs to the SSL Decryption
Exclusion list on your firewalls. See Enable
Access to Cortex XDR.

👀🟠

imp…

There can be instances in which the Broker VM
requires up to 50% of the incoming bandwidth
as outgoing. Such instances can be, network
instability between the Broker VM and Cortex
Cloud, or data that is being collected, but not
well compressed.

👀🟠

imp…

The internal network for the Broker VM must be
unique and reserved. Other devices should not
use the same IP as the Broker VM internal
network as it can lead to communication issues
with the Broker VM.

👀🟠

imp…

If you use SSL decryption in your firewalls and
proxies, see the Understanding CA certificate
functionality in Broker VM deployments section
below. In addition, verify that the proxies used
support HTTP/2, gRPC-specific headers, and
HTTP/2 trailers, and the inspection policies
support gRPC traffic. Any devices that you use
with this configuration...

👀🟠

imp…

When configuring more than one network
interface, ensure that only one Default
Gateway is defined. The rest must be set to
0.0.0.0, which configures them as undefined. In
addition, we recommend assigning each
network interface to a different subnet, as
oppose to configuring two interfaces on the
same subnet which can potentially cause
unexpected...

👀🟠

imp…

Internal IP must be:
Formatted as prefix/mask, for example
192.0.2.1/24.
Must be within /8 to /24 range.
Cannot be configured to end with a zero.
For Broker VM version 9.0 and earlier, Cortex
Cloud will only accept 172.17.0.0/16.

👀🟠

imp…

We strongly recommend disabling SSH
connectivity when it's not being used.
Therefore, activate SSH connectivity when it's
needed and disable it right afterwards.
When generating a new SSH key ensure to avoid
embedding the domain-style username, by not
using any backslashes (\) in the comment field,
to ensure the SSH key passes validation.

20.1.2.1. Broker VM image installations 

20.1.2.1.1. Set up Broker VM on Microsoft Azure 

☑️🟣

prer…

Download a Cortex Cloud Broker VM VHD
(Azure) image. For more information, see the
virtual machine compatibility requirements in
Set up and configure Broker VM.

📝🟢

note

Extraction can take up to a few hours.

📝🟢

note

Upload can take up to a few hours.

📝🟢

note

For more information, see the Azure
Documentation.

📝🟢

note

Creating the VM can take up to 15 minutes. The
Broker VM Web UI is not accessible during this
time.

📝🟢

note

For more information on creating a rule in an
Azure VM, see Create a Security Rule in the
Azure Documentation.

20.1.2.1.2. Set up Broker VM on Alibaba Cloud 

☑️🟣

prer…

Download a Cortex Cloud Broker VM QCOW2
image. For more information, see the virtual
machine compatibility requirements in Set up
and configure Broker VM.

📝🟢

note

For Linux and Windows uploads, you can use
Alibaba Cloud's graphical management tool
called ossbrowser.

📝🟢

note

The Object Storage Service must be created in
the same Region as the image of the virtual
machine.

🧭➡️

navi…

Hamburger menu → Elastic Compute Service
→ Instances & Images → Instances

20.1.2.1.3. Set up Broker VM on Amazon Web
Services 

☑️🟣

prer…

Download a Cortex Cloud Broker VM VMDK
image. For more information, see the virtual
machine compatibility requirements in Set up
and configure Broker VM.
You need to set up an AWS VM Import role
(vmimport) before running the import-snapshot
CLI command. If the role vmimport does not
exist or does not have the required permissions,
you can create it...

📝🟢

note

Assigning security groups can take up to 15
minutes.

👀🟠

imp…

A t3.xlarge (16 GB RAM) is the lowest machine
type that can be used as an instance type to
meet the mandatory 4 vCPU requirement.

🧭➡️

navi…

In the Cortex Cloud Web Console , select
Settings → Configurations → Data Broker →
Broker VMs → Add Broker → VMDK

In the AWS Console, select All services →
Storage → S3

All services → EC2 → IMAGES → AMIs and
locate your AMI image using the task ID

20.1.2.1.4. Set up Broker VM on Google Cloud
Platform (GCP) 

☑️🟣

prer…

Download a Cortex Cloud Broker VM VMDK
image. For more information, see the virtual
machine compatibility requirements in Set up
and configure Broker VM.
To complete the set up, you must have G Cloud
installed and have an authenticated user
account.

☑️🟣

prer…

Before importing a GCP image using the gcloud
CLI, ensure that you update the Google Cloud
components to version 371.0.0 and above using
the following command:

📝🟢

note

The import tool uses Cloud Build API, which
must be enabled in your project. For the import
to work, Cloud Build service account must have
compute.admin and iam.serviceAccountUser
roles. When using the Google Cloud console to
import the image, you will be prompted to add
these permissions automatically.

20.1.2.1.5. Set up Broker VM on KVM using
Ubuntu ☑️🟣 
prer…

Download a Cortex Cloud Broker VM QCOW2
image. For more information, see the virtual
machine compatibility requirements in Set up
and configure Broker VM.

20.1.2.1.6. Set up Broker VM on Microsoft Azure 

☑️🟣

prer…

Download a Cortex Cloud Broker VM VHD
(Azure) image. For more information, see the
virtual machine compatibility requirements in
Set up and configure Broker VM.

📝🟢

note

Extraction can take up to a few hours.

📝🟢

note

Upload can take up to a few hours.

📝🟢

note

For more information, see the Azure
Documentation.

📝🟢

note

Creating the VM can take up to 15 minutes. The
Broker VM Web UI is not accessible during this
time.

📝🟢

note

For more information on creating a rule in an
Azure VM, see Create a Security Rule in the
Azure Documentation.

20.1.2.1.7. Set up Broker VM on Nutanix
Hypervisor 

☑️🟣

prer…

Download a Cortex Cloud Broker VM QCOW2
image. For more information, see the virtual
machine compatibility requirements in Set up
and configure Broker VM.

📝🟢

note

Saving the image to Nutanix hypervisor can
take time as it's a large file. We recommend
verifying periodically that the connection is
alive for the upload process to finish
successfully.

📝🟢

note

Creating the VM can take up to 15 minutes. The
Broker VM Web user interface is not accessible
during this time.

20.1.2.1.8. Set up Broker VM on VMware ESXi
using vSphere Client ☑️🟣 
prer…

Ensure you have a virtualization platform
installed that is compatible with an OVA image,
and have an authenticated user account.
Download a Cortex Cloud Broker VM OVA
image. For more information, see the virtual
machine compatibility requirements in Set up
and configure Broker VM.

20.1.2.2. Broker VM data collector applets 

20.1.3. Manage Broker VM 

📝🟢

note

For more information on troubleshooting errors
and warnings for these broker applets, see
Troubleshoot Broker VM applet
errors.Troubleshoot Broker VM applet errors

📝🟢

note

Certain fields are also exposed in the Clusters
tab, when a Broker VM node is added to a High
Availability (HA) cluster, and each cluster node
is expanded to view the Broker VM nodes table.
An asterisk (*) is beside every field that is also
included in the Broker VM nodes table for each
HA cluster.

🧭➡️

navi…

Settings → Configurations → Data Broker →
Broker VMs

20.1.3.1. Edit Broker VM Configuration 

📝🟢

note

For all Broker VM nodes added to a HA cluster,
you can also Configure the Broker VM nodes
from the Clusters tab.

📝🟢

note

Internal IP must be:
Formatted as prefix/mask, for example
192.0.2.1/24.
Must be within /8 to /24 range.
Cannot be configured to end with a zero.
For Broker VM version 9.0 and lower, Cortex
Cloud accepts only 172.17.0.0/16.

📝🟢

note

Make sure you save the password before
closing the window. The only way to re-
generate a password is to disable ssh and re-
enable.

20.1.3.2. Increase Broker VM storage allocated
for data caching 

📝🟢

note

Follow your hypervisor documentation to
understand how to add a persistent disk
storage to your VM.

📝🟢

note

If your disk is not listed and you didn't
shutdown your Broker VM in your hypervisor
before manually adding a disk to the VM, you'll
need to reboot the Broker VM before the disk
details are detected by the Broker VM. This can
be performed either in the hypervisor or
directly in the Broker VMs page.

⛔🔴

war…

This operation is irreversible, and will make the
disk become an integral part of the broker,
where disconnecting the disk will result in
errors and data loss.

20.1.3.3. Monitor Broker VM using Prometheus 

📝🟢

note

For all Broker VM nodes added to a HA cluster,
you can also Configure the Broker VM nodes
from the Clusters tab.

📝🟢

note

For more information on Docker, see the Docker
website.

📝🟢 20.1.3.4. Collect Broker VM Logs 

note

The Download Logs (<TIMESTAMP>) is only
displayed when you've downloaded your logs
previously using Generate New Logs.

👀🟠 20.1.3.5. Upgrade Broker VM 

imp…

After a Broker VM upgrade, your broker may
require a reboot to finish installing important
updates. A notification about this will be sent
to your Cortex Cloud console Notification
Center.

20.1.3.6. Import Broker VM Configuration 

📝🟢

note

If this option is selected, the shutdown process
can take up to 10 minutes to allow the broker to
finish offloading the data cache.

👀🟠

imp…

This option can only be used on Broker VMs
with version 20.0 and later, and is only suitable
for importing a configuration of brokers in the
same version, or from a broker in an older
version to a broker in a newer version.

👀🟠

imp…

If your source Broker VM configuration includes
a WEC applet, you'll need to ensure that you
update the DNS record of this Broker VM's
FQDN to point to the target Broker VM IP
address.

📝🟢 20.1.3.7. Open Live Terminal 

note

You can either restart_routes or reboot the
Broker VM for the changes in the
/etc/network/routes file to take affect.

20.1.3.8. Add Broker VM to cluster 

20.1.3.9. Switchover Primary Node in Cluster 

📝🟢 20.1.3.10. Remove from Cluster 

note

If you've enabled a Load Balancer Health-
Check on the cluster, you need to exclude this
Broker VM from your Load Balancer settings.

20.1.4. Broker VM High Availability Cluster 

📝🟢

note

For Load Balancing, you must install a Load
Balancer in your network, which will distribute
the incoming data between the nodes.

20.1.4.1. Configure High Availability Cluster 

☑️🟣

prer…

Be sure you do the following task before
creating a cluster from an existing Broker VM:
If the Broker VM is explicitly specified in some
Agent Settings profile, which means Cortex
Cloud agents retrieve release upgrades and
content updates from this Broker VM, you must
change the Broker VM's current designated
role. To do this, you need to modify the...

📝🟢

note

The Broker VM Load Balancer Health Check
requires HTTP/1.1 or higher. Legacy HTTP/1.0 is
no longer supported. Ensure your external load
balancer is configured to use HTTP/1.1 or above
when performing health checks.

👀🟠

imp…

You can only create a new cluster from an
existing Broker VM, when the Broker VM version
is 19.0 and later, and the STATUS is Connected.
The Create a Cluster from this Broker option is
only listed if the Broker VM is not already
added to a cluster.

👀🟠

imp…

When the Protocol is set to HTTPS, you may
need to perform a few follow-up steps to
establish a validated secure SSL connection
with the Broker VM.
If you're using your own Certificate Authority
(CA) to sign the certificates, you'll need to
place the CA in the client, such as the Load
Balancer, and upload the certificates to the
Broker VM.
If you're...

20.1.4.2. Manage Broker VM clusters 

🧭➡️

navi…

After you've configured a cluster, you can
manage all your Broker VM clusters from the
Clusters tab on the Broker VMs page (Settings
→ Configurations → Data Broker → Broker VMs
→ Clusters

20.1.4.2.1. View cluster details 

📝🟢

note

For more information on troubleshooting errors
and warnings for these applets, see
Troubleshoot Broker VM applet
errors.Troubleshoot Broker VM applet errors

🧭➡️

navi…

The Clusters tab of the Broker VMs page
(Settings → Configurations → Data Broker →
Broker VMs) enables you

🧭➡️ 20.1.4.2.2. Edit cluster 

navi…

Settings → Configurations → Data Broker →
Broker VMs, and select the Clusters tab

20.1.4.2.3. Add applet to cluster 

📝🟢

note

With Cortex XDR Prevent, it's only relevant to
configure a HA cluster with a Local Agent
Settings applet as this is the only applet
supported for this product license. The other
applets are collector applets, which are only
available in Cortex XDR Pro or Cortex XSIAM.

📝🟢

note

For more information on troubleshooting errors
and warnings for these applets, see
Troubleshoot Broker VM applet
errors.Troubleshoot Broker VM applet errors

20.1.4.2.4. Add Broker VM to cluster 

20.1.4.2.5. Remove cluster 

20.1.5. Broker VM notifications 

📝🟢 20.1.6. Monitor Broker VM activity 

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

20.2. Dataset management 

🔒🟣

cont…

Managing Roles requires an Account Admin or
Instance Administrator role. For more
information, see Predefined user roles.

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

Cold storage, in addition to a cold storage
license, requires compute units (CU) to run cold
storage queries. For more information on CU,
see Manage compute units.

📝🟢

note

Only datasets that are already handled as part
of the GB license are supported for this license.
In addition, the retention configuration is only
available in Cortex Cloud, as opposed to the
public APIs.

📝🟢

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.
Datasets include dataset permission
enforcements in the Cortex Query
Language(XQL), Query Center, and XQL
Widgets. For example, to view or access any of
the endpoints and host_inventory datasets, you
need role-based access control (RBAC)...

📝🟢

note

You'll only be able to save the dataset view if
the query contains no errors; otherwise, the
Save button is disabled.

📝🟢

note

For more information on user roles, see Manage
user roles.

👀🟠

imp…

Cortex Cloud enforces retention on all log-
type datasets excluding Host Inventory,
Vulnerability Assessment, Metrics, and Users.

👀🟠

imp…

This column is updated once a day. Therefore,
if the dataset was created or updated by the
target or lookup flows, it's possible that the
Last Updated value is a day behind when the
queries or reports were run as it was before this
column was updated.

20.2.1. What are datasets? 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

CEF and LEEF logs are often wrapped in a
Syslog envelope (RFC 3164/5424). Cortex
Cloud automatically detects and strips the
Syslog header to extract the reporting device IP
and hostname before parsing the inner CEEF or
LEEF payload.

📝🟢

note

For more information on tasks specific to
lookup datasets, see Lookup datasets.

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query and can appear with a slight delay.

👀🟠

imp…

By default, forensic datasets are not included
in XQL query results, unless the dataset query
is explicitly defined to use a forensic dataset.

20.2.2. Lookup datasets 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

A lookup dataset can only be deleted if there
are no other dependencies. For example, if a
Correlation Rule is based on a lookup dataset,
you wouldn't be able to delete the lookup
dataset until you removed the dataset from the
XQL query of the Correlation Rule.

👀🟠

imp…

When you create or add data to a lookup
dataset using the target stage, the _time field
won't be included by default unless you
explicitly add it with the fields stage.

20.2.2.1. Import a lookup dataset 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

☑️🟣

prer…

When uploading a CSV, TSV, or JSON file,
ensure that the file meets the following
requirements:
The maximum size for the total data to be
imported into a lookup dataset is 30 MB from
the Dataset Management page. Otherwise, the
limit is 50 MB using Cortex Query Language
(XQL) or APIs.
Field names can contain characters from
different languages, special...

👀🟠

imp…

The name of a dataset created from a TSV file
must always include the extension. For example,
if the original file name is mrkdptusrsnov23.tsv,
you can save the dataset with the name
marketing_dept_users_Nov_2023.tsv.

20.2.2.2. Download JSON file of lookup dataset 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

🧭➡️

navi…

Open the Settings → Configurations → Data
Management → Dataset Management page

☑️🟣 20.2.2.3. Set time to live for lookup datasets 

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

20.2.3. Monitor datasets and dataset views

activity 

☑️🟣

prer…

Dataset Management requires View/Edit RBAC
permissions for Data Management (under
Configurations → Data Management), which
are the same permissions required for Parsing
Rules, Data Model Rules, and Event Forwarding.

📝🟢

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

20.3. Manage compute units 

📝🟢

note

For more information, see Import historical data
into cold storage.Import historical data into
cold storage

📝🟢

note

The CU consumption of cold storage queries is
based on the number of days in the query time
frame. For example, when querying 1 hour of a
specific day, the CU of querying this entire day
is consumed. When querying 1 hour that
extends past 2 days, such as from 23�50 to
00�50 of the following day, the CU of querying
these two days is consumed.

20.3.1. Compute units usage 

📝🟢

note

To enable the add-on, select Settings → Cortex
Cloud LicenseAddons tile, and select the
Compute Unit tile and Enable.

🧭➡️

navi…

Settings → Configurations → Data
Management → Compute Unit Usage

20.4. Manage instances 

20.4.1. Add a new data source or instance 

📝🟢

note

If a data source contains multiple integrations,
the integration configured as the default
integration will used by the Data Onboarder.
The default integration of the content pack is
indicated in each content pack's
documentation. The other integrations are
available for configuration in the Data Sources
& Integrations page after installing the content...

📝🟢

note

If you are adding a new instance to an existing
data source, these options are View only.
You can adjust the view-only options on the
relevant page in the system, for example
Correlations, Playbooks, or Scripts.
If you are adding a new instance to an existing
data source, these options are View only.
You can adjust the view-only options on the
relevant...

📝🟢 20.4.2. Manage cloud instances 

note

Currently, automation permission errors or
missing automation permissions do not affect
the Automation health status. You can view any
permission errors or missing permissions in the
the Permissions health status.

20.4.3. Update cloud permissions after Cortex
release updates 

☑️🟣

prer…

Ensure that the user account used to modify
permissions has the necessary privileges within
both the Cortex platform and your cloud
environment, for example, AWS or Azure.
You received a notification regarding a new
version available that requires permission
updates, or viewed a Needs Update status in
the Data Sources & Integrations page.

📝🟢

note

Instances requiring updates will not change
their connection status, for example,
Connected, Warning, Error, Disabled, due to the
pending permission update.

20.4.4. Pending cloud instances 

📝🟢 20.4.5. Manage Kubernetes Connector instances 

note

After uninstalling the Kubernetes connector,
the connector status updates to Not connected
48 hours after the uninstall process is initiated.

📝🟢 20.4.6. Troubleshoot errors on cloud instances 

note

Errors related to collection of audit logs in the
cloud instance are recorded in the
collection_auditing dataset. For more
information, see Audit logs fields and query
examples.

20.4.7. Monitor serverless function scan health

and status 

20.5. About health issues 

☑️🟣

prer…

For Cortex Cloud to monitor data ingestion
health and create health issues, you must
enable the following settings under
Configurations:
Cortex - Analytics: Go to Configurations →
Cortex - Analytics. For more information, see
Enable the Analytics Engine and Identity
Analytics.Enable the Analytics Engine and
Identity Analytics

📝🟢

note

The Health Issues page displays issues that
were triggered after July 2024. To see health
issues that were triggered before this date,
click Legacy Health Issues.

📝🟢

note

Cortex Cloud enforces the dedup logic to
health issues. This logic reduces the likelihood
of identical health issues from flooding the
issues dataset.

20.5.1. Investigate and resolve health issues 

21. Cortex Cloud Data Sources 

21.1. What are Cortex Cloud data sources? 

📝🟢

note

Some content packs provide parsing rules and
data model rules for data sources ingested
using standard data sources, and won't be
listed in the Data Sources & Integrations page.

👀🟠

imp…

We are actively working to evolve the UI to
create a single, unified point of configuration
for all data ingestion and integration points.
This documentation is structured to help you
navigate the current segmented process until
that UI update is available.

21.2. Complete data source catalog 

21.3. Vendor-specific data sources 

👀🟠

imp…

In some cases, the same vendor is available
through one or more of these options: standard
data source, CSP onboarding data source, or
content pack integration. Check the available
descriptions for each data source in both the
user interface and documentation to decide
which option is more suitable for your needs.

21.3.1. Amazon Web Services 

21.3.2. Cloud Next-Generation Firewall 

21.3.2.1. Ingest data from Cloud Next-
Generation Firewall 

🔒🟣

cont…

Requires the Data Collection add-on.

☑️🟣

prer…

Cortex Cloud RBAC permissions: Requires
View/Edit permissions for Data Sources (under
Configurations → Data Collections).
Cloud Service Provider (CSP) account
permissions: Configuration of data ingestion
from multiple accounts and regions requires
Super User permissions on both the Cortex
Cloud tenant and on the device accounts.
Note
Cross CSP...

📝🟢

note

If you change a device region, you must first
disconnect the device from Cortex Cloud and
then reconnect it after the region change is
complete.

👀🟠

imp…

This cross-account support is limited to the
same SFDC hierarchy; it does not extend to
MSSP scenarios where the customer and
provider own separate ends of the solution.

📝🟢 21.3.3.1. How to onboard Databricks  21.3.3. Databricks 

note

If you turn on this feature, both the cloud and
region will be used for scanning, possibly
incurring cost and requiring adherence to
certain compliance policies.

21.3.4. Kubernetes 

21.3.4.1. Onboard the Kubernetes Connector 

📝🟢

note

This option is not supported for Fargate.

📝🟢

note

Enabling Realtime Protection installs the agent
on your Kubernetes clusters as a DaemonSet.

📝🟢

note

For Fargate, you must provide the cluster
resource identifier.
The format of the identifier is arn:aws:eks:
<region>:<account-id>:cluster/<cluster-name>.

📝🟢

note

Basic authentication is only supported in
Posture Management. If deploying Realtime
Protection, select None .

📝🟢

note

If you install the latest version but select the N-
1 strategy, this policy will take effect starting
from the next upgrade cycle (it will not
immediately downgrade your current
installation).
If you choose an older version and keep the
latest strategy, the latest version will be
installed.

21.3.4.2. What's new in Kubernetes Connector? 

21.3.4.3. Supported Kubernetes distributions 

📝🟢

note

Does not include Autopilot.

📝🟢

note

In Google Container-Optimized OS release 100
and earlier, where the FANOTIFY EXEC flag is
not supported, the Kernel configuration may be
partial for the user mode agent to properly
function. In such cases, the agent will fallback
to asynchronous mode.
In RHCOS version 4.12 and earlier, the Kernel
configuration may be partial for the user mode
agent to...

21.3.5. Microsoft Azure 

☑️🟣 21.3.6.1. How to onboard Microsoft 365  21.3.6. Microsoft 365 (Posture) 

prer…

You have generated a Globally Unique Identifier
(GUID), also known as a Universally Unique
Identifier (UUID). You will need this ID for the
tenant you want to use for the Microsoft 365
instance.
In order to use Microsoft 365, you must be
registered with Microsoft Azure.

☑️🟣 21.3.7.1. Ingest logs and data from Okta  21.3.7. Okta 

prer…

Administrator privileges: Your Okta user must
have a role capable of creating API tokens,
such as Read-only Administrator, Super
Administrator, or Organization Administrator.
For more information, see the Okta
Administrators Documentation.

21.3.8. Oracle Cloud Infrastructure 

21.3.9. Salesforce 

21.3.9.1. Ingest and run Salesforce automation
and remediation 

☑️🟣

prer…

RBAC permissions: Requires View/Edit
permissions for Log Collections, Data Sources,
and Integrations (under Configurations & Data
Collections).
Content packs: Ensure the Salesforce and Base
content packs are installed or updated to the
latest version.
Gateway permissions: Requires Account Admin
or Instance Administrator permissions for
configuring...

📝🟢

note

Do not include https:// or trailing slashes.

📝🟢

note

The new Salesforce integration has the
description: Salesforce CRM services for
identity management, automation, remediation
and SaaS Posture Security.

👀🟠

imp…

You cannot proceed with the authentication
configuration until you provide a domain URL
and click Apply.When you click Apply, the
instance is connected to a unique URL in
Salesforce, and it cannot be changed. To
change it, you need to exit the wizard and start
over.

👀🟠

imp…

If a '''connected application'' for Cortex Cloud
data collection already exists, there is no
obligation to migrate to an '''External Client
App'' - but as Connected Applications are
deprecated by Salesforce, it is recommended
to migrate.

21.3.10.1. How to onboard Snowflake  21.3.10. Snowflake 

☑️🟣

prer…

In order to use Snowflake, you must be
registered with one of these cloud providers:
Amazon AWS, Microsoft Azure, or Google
Cloud Platform (GCP).
Ensure you have the necessary account
permissions to onboard. It is recommended to
use Account Admin as the role for the
onboarding.

📝🟢

note

The account identifier can be found using the
user information at the bottom left. Hover over
the account you wish to onboard and select
the copy option at the top right. The account
identifier is usually of the format:
(organization).[account]

21.4. Cloud service provider (CSP) onboarding 

21.4.1. Cloud service provider onboarding 

21.4.2. Understand CSP onboarding tiers and
licensing 

21.4.3. Amazon Web Services cloud onboarding 

📝🟢 21.4.3.1. Prerequisites 

note

The permissions included in Sid:
CortexBaseModuleDeployment are always
required for the discovery module and stack
lifecycle management.
The permissions included in Sid:
CortexAuditLogsModule are required only
when deploying the audit log collection
module. It enables creating S3, SQS, SNS,
CloudTrail, and KMS resources.

21.4.3.2. Onboard Amazon Web Services 

21.4.3.3. How to onboard Amazon Web Services 

📝🟢

note

Scanning with an outpost may require
additional AWS permissions and may incur
additional CSP costs.

👀🟠

imp…

It is critical to ensure that your KMS key region
and SNS topic region are the exact same as the
AWS region where you are deploying the
CloudFormation stack.

21.4.3.4. How to onboard Amazon Web Services

with foundational configuration 

🔒🟣

cont…

Onboarding Amazon Web Services (AWS) using
the foundational configuration is included with
Cortex Cloud NG SIEM, Cortex Cloud
Enterprise, and Cortex Cloud Enterprise+
licenses. For more details on the CSP
onboarding tiers and licensing, see Understand
CSP onboarding tiers and licensing.

👀🟠

imp…

It is critical to ensure that your KMS key region
and SNS topic region are the exact same as the
AWS region where you are deploying the
CloudFormation stack.

21.4.3.5. Deploy the CloudFormation template in
AWS 

21.4.3.6. AWS post-deployment verification 

21.4.4. Microsoft Azure cloud onboarding 

21.4.4.1. Onboard Microsoft Azure 

21.4.4.2. Prerequisites 

21.4.4.3. How to onboard Microsoft Azure 

📝🟢

note

Scanning with an outpost may require
additional GCP permissions and may incur
additional CSP costs.

📝🟢

note

Cortex Cloud performs a live verification of
each tenant's approval status against Azure
each time the list is displayed. If a previously
approved tenant no longer shows a green
checkmark, the Cortex Cloud application may
have been removed from the Azure tenant.
Click Approve in Azure to re-authorize.

👀🟠

imp…

It is critical to ensure that your namespace and
Event Hub belong to the specific Azure
subscription being onboarded. Cross-
subscription or centralized logging is not
currently supported.

21.4.4.4. How to onboard Microsoft Azure with

foundational configuration 

🔒🟣

cont…

Onboarding Microsoft Azure using the
foundational configuration is included with
Cortex Cloud NG SIEM, Cortex Cloud
Enterprise, and Cortex Cloud Enterprise+
licenses. For more details on the CSP
onboarding tiers and licensing, see Understand
CSP onboarding tiers and licensing.

📝🟢

note

Cortex Cloud performs a live verification of
each tenant's approval status against Azure
each time the list is displayed. If a previously
approved tenant no longer shows a green
checkmark, the Cortex Cloud application may
have been removed from the Azure tenant.
Click Approve in Azure to re-authorize.

👀🟠

imp…

It is critical to ensure that your namespace and
Event Hub belong to the specific Azure
subscription being onboarded. Cross-
subscription or centralized logging is not
currently supported.

21.4.4.5. Finalize Microsoft Azure onboarding by
executing the authentication template 

☑️🟣

prer…

Before you begin, ensure you have:
An Azure subscription.
A user with the required permissions for the
relevant scope (subscription, management
group, tenant). We recommend you create a
dedicated role.
Tenant ID and subscription ID. You can view
these in Microsoft Azure Portal in Management
groups.

Installed Terraform on your local machine. You
can...

☑️🟣

prer…

Before you begin, ensure you have:
An Azure subscription.
A user with the required permissions for the
relevant scope (subscription, management
group, tenant). We recommend you create a
dedicated role.
Tenant ID and subscription ID. You can view
these in Microsoft Azure Portal in Management
groups.
Installed the Azure CLI tool.
Authorization to create...

☑️🟣

prer…

Before you begin, ensure you have:
An Azure subscription.
A user with the required permissions for the
relevant scope (subscription, management
group, tenant). We recommend you create a
dedicated role.
Tenant ID and subscription ID. You can view
these in Microsoft Azure Portal in Management
groups.
Installed the Azure CLI tool.
Authorization to create...

👀🟠

imp…

Do not delete or move the Terraform files from
this folder. It will prevent you from being able
to edit your cloud instance in the future.

21.4.5. GCP cloud onboarding 

21.4.5.1. Onboard Google Cloud Platform 

21.4.5.2. Prerequisites 

21.4.5.3. How to onboard Google Cloud Platform 

📝🟢

note

Scanning with an outpost may require
additional GCP permissions and may incur
additional CSP costs.

📝🟢

note

If you want to connect to your GCP Workspace,
you must first complete onboarding with the
option disabled. Once the GCP cloud instance
is created, perform the steps detailed in
Connect Google Workspace with your GCP
cloud instance.

21.4.5.4. How to onboard GCP with foundational

configuration 🔒🟣 
cont…

Onboarding Google Cloud Platform (GCP)
using the foundational configuration is included
with Cortex Cloud NG SIEM, Cortex Cloud
Enterprise, and Cortex Cloud Enterprise+
licenses. For more details on the CSP
onboarding tiers and licensing, see Understand
CSP onboarding tiers and licensing.

21.4.5.5. Deploy the Terraform authentication
template in GCP 

☑️🟣

prer…

Before you begin, ensure you have:
A GCP account.
Permission to create the required resources in
Google Cloud Deployment Manager.
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the GCP gcloud CLI tool.
Reviewed the...

📝🟢

note

The directory you create must be a
subdirectory of the home directory.

👀🟠

imp…

You must not delete or move the Terraform files
from this folder. It will prevent you from being
able to edit your cloud instance in the future.

21.4.5.6. Connect Google Workspace with your
GCP cloud instance 

☑️🟣

prer…

Ensure you have the Super Admin role in
Google Workspace.

☑️🟣

prer…

Ensure you have the organization ID of the
Google Workspace you want to connect:
Log in to your Google Admin Console. and
navigate to Account → Account settings →
Profile. Next to Customer ID is your
organization ID.
Ensure you have the organization ID of the
Google Workspace you want to connect:
Log in to your Google Admin Console. and
navigate...

21.4.6. Oracle Cloud Infrastructure cloud

onboarding 

21.4.6.1. Onboard Oracle Cloud Infrastructure 

21.4.6.2. Prerequisites 

21.4.6.3. How to onboard Oracle Cloud
Infrastructure



21.4.6.4. How to onboard Oracle Cloud

Infrastructure with foundational configuration 🔒🟣 
cont…

Onboarding Oracle Cloud Infrastructure (OCI)
using the foundational configuration is included
with Cortex Cloud NG SIEM, Cortex Cloud
Enterprise, and Cortex Cloud Enterprise+
licenses. For more details on the CSP
onboarding tiers and licensing, see Understand
CSP onboarding tiers and licensing.

21.4.6.5. Deploy the Terraform authentication
template in OCI ☑️🟣 
prer…

Before you begin, ensure you have:
An Oracle Cloud Infrastructure account and
the tenancy OCID.
Permission to deploy a custom template and
create its resources in OCI.
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed...

21.4.7. Alibaba Cloud cloud onboarding 

21.4.7.1. Security capabilities and deployment
planning 

21.4.7.2. Resource inventory 

21.4.7.3. Security model and authentication 

21.4.7.4. Onboard Alibaba Cloud 

21.4.7.5. Prerequisites 

21.4.7.6. How to onboard Alibaba Cloud 

☑️🟣

prer…

Before you begin, ensure you have:
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the Alibaba Cloud CLI tool.

📝🟢

note

If the following error appears, contact support:
"Validation failed: No valid managed outpost
found for cloud provider ALIBABA_CLOUD".

👀🟠

imp…

Do not delete or move the Terraform files from
this folder. It will prevent you from being able
to edit your cloud instance in the future.

21.4.7.7. Alibaba Cloud post-deployment
verification 

21.4.8. Manually connect a cloud instance 

📝🟢 21.4.9. Manage cloud instances 

note

Currently, automation permission errors or
missing automation permissions do not affect
the Automation health status. You can view any
permission errors or missing permissions in the
the Permissions health status.

21.4.10. Pending cloud instances 

👀🟠 21.4.11. Edit your onboarded CSP configuration 

imp…

When using Terraform authentication
templates, you must execute the updated
Terraform template from the same folder where
the original Terraform template was executed.

21.4.12. Update cloud permissions after Cortex
release updates 

☑️🟣

prer…

Ensure that the user account used to modify
permissions has the necessary privileges within
both the Cortex platform and your cloud
environment, for example, AWS or Azure.
You received a notification regarding a new
version available that requires permission
updates, or viewed a Needs Update status in
the Data Sources & Integrations page.

📝🟢

note

Instances requiring updates will not change
their connection status, for example,
Connected, Warning, Error, Disabled, due to the
pending permission update.

📝🟢 21.4.13. Troubleshoot errors on cloud instances 

note

Errors related to collection of audit logs in the
cloud instance are recorded in the
collection_auditing dataset. For more
information, see Audit logs fields and query
examples.

21.4.14. Outposts 

👀🟠

imp…

Outpost scan is an alternative to the
recommended standard cloud scan. Cloud scan
is recommended because it is fully managed by
Palo Alto Networks and incurs no compute
costs for your organization. Outpost scan is an
advanced deployment model reserved for
specific data residency or architectural
requirements.

21.4.14.1. Outpost fundamentals and planning 

📝🟢

note

Notes:

Before you create your outpost, verify that
your internet connection is active. An active
internet connection is necessary for the
notification to be sent to Cortex Cloud to
create the new outpost.
(Azure) Due to limitations in Terraform, the
Azure subscription name cannot contain blanks.
Take this into account while onboarding.

👀🟠

imp…

While outposts provide maximum control over
the scanning environment, cloud scan mode is
the recommended default for most
organizations.

🧭➡️

navi…

Before onboarding your Cortex Cloud with the
cloud service provider (CSP) onboarding
wizard, create an outpost by navigating to
Settings → Data Sources & Integrations →
Outposts

Settings → Data Sources & Integrations → Add
New

View and manage existing outposts by
navigating to Settings → Data Sources &
Integrations → Outposts

21.4.14.2. Create an outpost 

☑️🟣

prer…

Before you begin, ensure you have:
An AWS account
Permission to create a stack and its resources in
AWS
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the AWS CLI tool and configured your
profile with the aws...

☑️🟣

prer…

Before you begin, ensure you have:
A GCP account
Permission to create the required resources in
Google Cloud Deployment Manager
Installed Terraform on your local machine. You
can download Terraform from the official
Terraform website and follow the installation
instructions for your operating system.
Installed the GCP gcloud CLI tool

☑️🟣

prer…

Before you begin, ensure you have:
An active Azure subscription.
Installed the Azure CLI tool.
Permission to deploy a custom template and
create its resources in Microsoft Azure
("Owner" or "Contributor" on the designated
outpost subscription scope, and Active
Directory "Cloud Application Administrator" or
"Application Administrator" privileged...

📝🟢

note

Verify that your internet connection is active.
An active internet connection is necessary for
notifications to be sent to Cortex Cloud for
creating the new outpost. If you are unable to
establish an internet connection, contact
customer support for a manual workaround.

📝🟢

note

When creating an outpost for a specific Azure
subscription, the outpost account must be in
the same Azure organization as the monitored
subscriptions.

📝🟢

note

Due to limitations in Terraform, the Azure
subscription name cannot contain blanks.

👀🟠

imp…

While outposts provide maximum control over
the scanning environment, cloud scan mode is
the recommended default for most
organizations. For details, see When to choose
outpost scan.

🧭➡️

navi…

Start the outpost creation wizard by navigating
to Settings → Data Sources & Integrations →
Outposts and clicking New Outpost

A new outpost is created in pending status and
can be viewed in the Outpost page at Settings
→ Data Sources & Integrations → Outposts

21.4.15. Introduction to Terraform for Cloud

service provider (CSP) onboarding 

👀🟠

imp…

Always run these commands in the same folder
where the original .tf files and .terraform folder
live—this is where the state is stored.

👀🟠

imp…

This .tfstate state file is critical because it is
needed by the terraform destroy command to
clean up created resources. It is critical that
you never delete this file.

21.4.16. Cloud service provider permissions 

⚠️🟡

caut…

Reducing permissions below the recommended
level may cause specific commands to fail or
limit functionality in future pack updates.

21.4.16.1. Amazon Web Services provider
permissions 

📝🟢

note

Unified Cortex platform cloud content packs
require a specific set of automation permissions
to enable full integration with your cloud
environment. Before configuring access for
these packs, review the automation permission
scope guidelines.

👀🟠

imp…

All conditional capabilities documented below
require the mandatory Base and Discovery
Engine permissions to be deployed alongside
them. Base provides the foundational
CortexPlatformRole and AWS-managed read-
only baseline. Discovery Engine extends that
baseline with the asset-inventory coverage that
every other capability assumes.

📝🟢 21.4.16.2. Microsoft Azure provider permissions 

note

Unified Cortex platform cloud content packs
require a specific set of automation permissions
to enable full integration with your cloud
environment. Before configuring access for
these packs, review the automation permission
scope guidelines.

21.4.16.3. Google Cloud Platform provider
permissions 📝🟢 
note

Unified Cortex platform cloud content packs
require a specific set of automation permissions
to enable full integration with your cloud
environment. Before configuring access for
these packs, review the automation permission
scope guidelines.

21.4.16.4. Oracle Cloud Infrastructure provider
permissions 

21.5. Container Registry Scanning 

21.5.1. Overview of container registry scanning 

21.5.1.1. Registry Components 

21.5.1.2. How Container Registry Scanning Works 

21.5.2. Configure registry scanning for cloud
accounts ☑️🟣 
prer…

Ensure that you have performed the all steps
till Additional Security Capabilities as listed in
the onboarding wizard for the required CSP:
Onboard Amazon Web ServicesOnboard
Amazon Web Services
Onboard Google Cloud PlatformOnboard
Google Cloud Platform
Onboard Microsoft AzureOnboard Microsoft
Azure
Onboard Oracle Cloud InfrastructureOnboard
Oracle...

21.5.3. Modify the container registry scanning
scope



21.5.4. Scan re-evaluation process 

21.5.5. Connect Docker Hub registry 

☑️🟣

prer…

Ensure an Outpost is connected to your tenant.

☑️🟣

prer…

Set up and configure Broker VM
Configure High Availability Cluster

📝🟢

note

If you choose Azure as the Cloud Provider, you
must also select the Tenant Id. The Tenant Id is
required to approve Cortex as an enterprise
application in your Azure tenant.

📝🟢

note

If you choose Azure as the cloud provider, only
Outposts associated with the selected tenant
ID are displayed.

📝🟢

note

The list of Broker VMs displays only VMs that
support registry scanning.
The list of high-availability Clusters displays
only clusters that contain at least one VM
supporting registry scanning.
The registry scanning status for each VM
appears in brackets if it was previously
activated for that specific VM.

📝🟢 21.5.5.1. Manage a Docker Hub connector 

note

If you selected Scan with Broker VM mode, you
can't change to a different scan mode (such as
Cloud Discovery or Scan with Outpost) when
you edit the instance.
When editing an instance configured for Scan
with Broker VM, you must re-enter your
authentication credentials, including Username,
Password, and CA certificate.

21.5.6. Connect Docker V2 compliant container
registry 

☑️🟣

prer…

Ensure an Outpost is connected to your tenant.
Outposts

☑️🟣

prer…

Ensure one of the following is configured:
Set up and configure Broker VM.
Configure High Availability Cluster.

📝🟢

note

If you choose Azure as the Cloud Provider, you
must also select the Tenant Id. The Tenant Id is
required to approve Cortex as an enterprise
application in your Azure tenant.

📝🟢

note

If you choose Azure as the cloud provider, only
Outposts associated with the selected tenant
ID are displayed.

📝🟢

note

The list of Broker VMs displays only VMs that
support registry scanning.
The list of high-availability Clusters displays
only clusters that contain at least one VM
supporting registry scanning.
The registry scanning status for each VM
appears in brackets if it was previously
activated for that specific VM.

📝🟢 21.5.6.1. Manage a Docker V2 connector 

note

If you selected Scan with Broker VM mode, you
can't change to a different scan mode (such as
Cloud Discovery or Scan with Outpost) when
you edit the instance.
When editing an instance configured for Scan
with Broker VM, you must re-enter your
authentication credentials, including Username,
Password, and CA certificate.

21.5.7. Connect GitLab container registry 

☑️🟣

prer…

Ensure an Outpost is connected to your tenant.

☑️🟣

prer…

Set up and configure Broker VM
Configure High Availability Cluster

📝🟢

note

When both the group ID and project ID are
provided, the system retrieves container images
from all projects within the specified group as
well as from the specified project.
If neither the group ID nor the project ID is
provided, the system retrieves container images
from all registries (across all groups and
projects) accessible to the authenticated...

📝🟢

note

When both the group ID and project ID are
provided, the system retrieves container images
from all projects within the specified group as
well as from the specified project.
If neither the group ID nor the project ID is
provided, the system retrieves container images
from all registries (across all groups and
projects) accessible to the authenticated...

📝🟢

note

If you choose Azure as the Cloud Provider, you
must also select the Tenant Id. The Tenant Id is
required to approve Cortex as an enterprise
application in your Azure tenant.

📝🟢

note

If you choose Azure as the cloud provider, only
Outposts associated with the selected tenant
ID are displayed.

📝🟢

note

When both the group ID and project ID are
provided, the system retrieves container images
from all projects within the specified group as
well as from the specified project.
If neither the group ID nor the project ID is
provided, the system retrieves container images
from all registries (across all groups and
projects) accessible to the authenticated...

📝🟢

note

When both the group ID and project ID are
provided, the system retrieves container images
from all projects within the specified group as
well as from the specified project.
If neither the group ID nor the project ID is
provided, the system retrieves container images
from all registries (across all groups and
projects) accessible to the authenticated...

📝🟢

note

The list of Broker VMs displays only VMs that
support registry scanning.
The list of high-availability Clusters displays
only clusters that contain at least one VM
supporting registry scanning.
The registry scanning status for each VM
appears in brackets if it was previously
activated for that specific VM.

📝🟢

note

When both the group ID and project ID are
provided, the system retrieves container images
from all projects within the specified group as
well as from the specified project.
If neither the group ID nor the project ID is
provided, the system retrieves container images
from all registries (across all groups and
projects) accessible to the authenticated...

📝🟢

note

When both the group ID and project ID are
provided, the system retrieves container images
from all projects within the specified group as
well as from the specified project.
If neither the group ID nor the project ID is
provided, the system retrieves container images
from all registries (across all groups and
projects) accessible to the authenticated...

21.5.7.1. Manage a Gitlab Container Registry
connector 📝🟢 
note

If you selected Scan with Broker VM mode, you
can't change to a different scan mode (such as
Cloud Discovery or Scan with Outpost) when
you edit the instance.
When editing an instance configured for Scan
with Broker VM, you must re-enter your
authentication credentials, including Username,
Password, and CA certificate.

21.5.8. Connect Harbor registry 

☑️🟣

prer…

Ensure an Outpost is connected to your tenant.

☑️🟣

prer…

Ensure one of the following is configured:
Set up and configure Broker VM.
Configure High Availability Cluster.

📝🟢

note

If you choose Azure as the Cloud Provider, you
must also select the Tenant Id. The Tenant Id is
required to approve Cortex as an enterprise
application in your Azure tenant.

📝🟢

note

If you choose Azure as the cloud provider, only
Outposts associated with the selected tenant
ID are displayed.

📝🟢

note

The list of Broker VMs displays only VMs that
support registry scanning.
The list of high-availability Clusters displays
only clusters that contain at least one VM
supporting registry scanning.
The registry scanning status for each VM
appears in brackets if it was previously
activated for that specific VM.

📝🟢 21.5.8.1. Manage a Harbor connector 

note

If you selected Scan with Broker VM mode, you
can't change to a different scan mode (such as
Cloud Discovery or Scan with Outpost) when
you edit the instance.
When editing an instance configured for Scan
with Broker VM, you must re-enter your
authentication credentials, including Username,
Password, and CA certificate.

21.5.9. Connect JFrog container registry 

☑️🟣

prer…

Ensure an Outpost is connected to your tenant.

☑️🟣

prer…

Set up and configure Broker VM
Configure High Availability Cluster

📝🟢

note

If you choose Azure as the Cloud Provider, you
must also select the Tenant Id. The Tenant Id is
required to approve Cortex as an enterprise
application in your Azure tenant.

📝🟢

note

If you choose Azure as the cloud provider, only
Outposts associated with the selected tenant
ID are displayed.

📝🟢

note

The list of Broker VMs displays only VMs that
support registry scanning.
The list of high-availability Clusters displays
only clusters that contain at least one VM
supporting registry scanning.
The registry scanning status for each VM
appears in brackets if it was previously
activated for that specific VM.

📝🟢 21.5.9.1. Manage a JFrog connector 

note

If you selected Scan with Broker VM mode, you
can't change to a different scan mode (such as
Cloud Discovery or Scan with Outpost) when
you edit the instance.
When editing an instance configured for Scan
with Broker VM, you must re-enter your
authentication credentials, including Username,
Password, and CA certificate.

21.5.10. Connect Sonatype Nexus registry 

☑️🟣

prer…

Ensure an Outpost is connected to your tenant.
Outposts

☑️🟣

prer…

Ensure one of the following is configured:
Set up and configure Broker VM.
Configure High Availability Cluster.

📝🟢

note

If you are using a CA certificate, enter the
server IP address instead of the registry url.

📝🟢

note

If you choose Azure as the Cloud Provider, you
must also select the Tenant Id. The Tenant Id is
required to approve Cortex as an enterprise
application in your Azure tenant.

📝🟢

note

If you choose Azure as the cloud provider, only
Outposts associated with the selected tenant
ID are displayed.

📝🟢

note

If you are using a CA certificate, enter the
server IP address instead of the registry URL.

📝🟢

note

The list of Broker VMs displays only VMs that
support registry scanning.
The list of high-availability Clusters displays
only clusters that contain at least one VM
supporting registry scanning.
The registry scanning status for each VM
appears in brackets if it was previously
activated for that specific VM.

📝🟢

note

If you are using a CA certificate, enter the
server IP address instead of the registry URL.

📝🟢 21.5.10.1. Manage a Sonatype connector 

note

If you had selected Scan with Broker VM mode,
you can't change to a different scan mode
(such as Cloud Discovery or Scan with
Outpost) when you edit the instance.
When editing an instance configured for Scan
with Broker VM, you must re-enter your
authentication credentials, including Username,
Password, and CA certificate.

21.6.1. Broker VM data collector applets  21.6. Generic on-premise data collectors 

21.6.1.1. Activate DSPM Fileshare 

☑️🟣

prer…

Set up and configure Broker VM
Know the complete path to the files and
folders that you want Cortex Cloud to monitor.
Necessary user permissions to access the
network shares. For the SMB connection type,
you need the user name and password.

📝🟢

note

The applet list displays only the applets for
which you have permissions.

📝🟢

note

By default, all configured connections are
saved.

📝🟢

note

For details regarding the connection fields, see
the table above under File Share Connection.

📝🟢

note

You can add multiple connections under a
single instance of the DSPM Fileshare applet by
returning to the File Share Connection screen
and clicking Add Connection. Each new
connection can be of either the NFS or SMB
connection type.

21.6.1.2. Activate Registry Scanner 

☑️🟣

prer…

To initialize registry scanning on your Broker
VM, you must first add the necessary data
connectors. For details, see:
Connect Docker Hub registry
Connect Docker V2 compliant container
registry
Connect GitLab container registry
Connect Harbor registry
Connect JFrog container registry
Connect Sonatype Nexus registry
To initialize registry...

📝🟢

note

You cannot activate the Registry Scanner
directly on a new or existing Broker VM. You
can only activate or deactivate existing
Registry Scanner applets. To activate or
deactivate existing applets, see Step 4 under
Verify Registry Scanner connection section.

21.6.1.3. Activate Transporter 

☑️🟣

prer…

Permissions: To configure and manage
Transporter applet settings, you must have
permissions to manage Broker Service
configurations (such as an Instance
Administrator)
Set up and configure Broker VMSet up and
configure Broker VM
Confirm that your Broker is v 28 or above
Whitelist IP addresses to enable access to
Cortex Cloud resources. The IP addresses...

⛔🔴

war…

The Transporter applet is not supported for
FedRAMP customers.

21.7. Cloud Posture and Runtime Security data
sources



📝🟢 21.7.1. How to onboard Databricks 

note

If you turn on this feature, both the cloud and
region will be used for scanning, possibly
incurring cost and requiring adherence to
certain compliance policies.

21.7.2. How to onboard on-premise assets to
Cortex Cloud Data Security 

📝🟢

note

Your data is scanned on the Broker VM itself,
and only the metadata and classification
results are transmitted from the on-premise
environment to Cortex Cloud.

📝🟢

note

For information about activating the DSPM
Fileshare applet, see Activate DSPM Fileshare.

📝🟢

note

For information about activating the DSPM
Database applet, see Activate DSPM Database.

21.7.2.1. Activate DSPM Fileshare 

☑️🟣

prer…

Set up and configure Broker VM
Know the complete path to the files and
folders that you want Cortex Cloud to monitor.
Necessary user permissions to access the
network shares. For the SMB connection type,
you need the user name and password.

📝🟢

note

The applet list displays only the applets for
which you have permissions.

📝🟢

note

By default, all configured connections are
saved.

📝🟢

note

For details regarding the connection fields, see
the table above under File Share Connection.

📝🟢

note

You can add multiple connections under a
single instance of the DSPM Fileshare applet by
returning to the File Share Connection screen
and clicking Add Connection. Each new
connection can be of either the NFS or SMB
connection type.

21.7.2.2. Activate DSPM Database 

📝🟢

note

The applet list displays only the applets for
which you have permissions.

📝🟢

note

By default, all configured connections are
saved.

📝🟢

note

For details regarding the connection fields, see
the table above under Database Connection.

📝🟢

note

If you get a Test connection failed message,
check your parameters and try again.

📝🟢

note

You can add multiple connections under a
single instance of the DSPM Database applet
by returning to the Database Connection
screen and clicking Add Connection.

⛔🔴

war…

Set up and configure Broker VM
Know the database engine, host, and port for
the database you want to connect.

☑️🟣 21.7.3. How to onboard Microsoft 365 

prer…

You have generated a Globally Unique Identifier
(GUID), also known as a Universally Unique
Identifier (UUID). You will need this ID for the
tenant you want to use for the Microsoft 365
instance.
In order to use Microsoft 365, you must be
registered with Microsoft Azure.

☑️🟣 21.7.4. Ingest logs and data from Okta 

prer…

Administrator privileges: Your Okta user must
have a role capable of creating API tokens,
such as Read-only Administrator, Super
Administrator, or Organization Administrator.
For more information, see the Okta
Administrators Documentation.

21.7.5. Activate Registry Scanner 

☑️🟣

prer…

To initialize registry scanning on your Broker
VM, you must first add the necessary data
connectors. For details, see:
Connect Docker Hub registry
Connect Docker V2 compliant container
registry
Connect GitLab container registry
Connect Harbor registry
Connect JFrog container registry
Connect Sonatype Nexus registry
To initialize registry...

📝🟢

note

You cannot activate the Registry Scanner
directly on a new or existing Broker VM. You
can only activate or deactivate existing
Registry Scanner applets. To activate or
deactivate existing applets, see Step 4 under
Verify Registry Scanner connection section.

21.7.6. How to onboard Snowflake 

☑️🟣

prer…

In order to use Snowflake, you must be
registered with one of these cloud providers:
Amazon AWS, Microsoft Azure, or Google
Cloud Platform (GCP).
Ensure you have the necessary account
permissions to onboard. It is recommended to
use Account Admin as the role for the
onboarding.

📝🟢

note

The account identifier can be found using the
user information at the bottom left. Hover over
the account you wish to onboard and select
the copy option at the top right. The account
identifier is usually of the format:
(organization).[account]

21.7.7. Activate Transporter 

☑️🟣

prer…

Permissions: To configure and manage
Transporter applet settings, you must have
permissions to manage Broker Service
configurations (such as an Instance
Administrator)
Set up and configure Broker VMSet up and
configure Broker VM
Confirm that your Broker is v 28 or above
Whitelist IP addresses to enable access to
Cortex Cloud resources. The IP addresses...

⛔🔴

war…

The Transporter applet is not supported for
FedRAMP customers.

🧭➡️

navi…

Settings → Configurations → Broker VMs
(under Data Broker

the Brokers tab → locate your Broker VM →
hover and click + Add under the Apps column
→ AppSec Transporter

the Appsec Transporter under the App column
→ Deactivate → Confirm when prompted All
existing connections are deleted but their
configurations are saved in the database

the Appsec Transporter under the App column
→ Deactivate → Confirm when prompted

21.8. Administration and troubleshooting 

21.8.1. Manage instances 

21.8.1.1. Add a new data source or instance 

📝🟢

note

If a data source contains multiple integrations,
the integration configured as the default
integration will used by the Data Onboarder.
The default integration of the content pack is
indicated in each content pack's
documentation. The other integrations are
available for configuration in the Data Sources
& Integrations page after installing the content...

📝🟢

note

If you are adding a new instance to an existing
data source, these options are View only.
You can adjust the view-only options on the
relevant page in the system, for example
Correlations, Playbooks, or Scripts.
If you are adding a new instance to an existing
data source, these options are View only.
You can adjust the view-only options on the
relevant...

21.8.1.2. How to configure the scanning settings
for supported services 📝🟢 
note

If you do not select a scanning cadence, the
default setting is applied. For more information,
contact your Customer Support team.

📝🟢 21.8.1.3. Manage cloud instances 

note

Currently, automation permission errors or
missing automation permissions do not affect
the Automation health status. You can view any
permission errors or missing permissions in the
the Permissions health status.

21.8.1.4. Update cloud permissions after Cortex
release updates 

☑️🟣

prer…

Ensure that the user account used to modify
permissions has the necessary privileges within
both the Cortex platform and your cloud
environment, for example, AWS or Azure.
You received a notification regarding a new
version available that requires permission
updates, or viewed a Needs Update status in
the Data Sources & Integrations page.

📝🟢

note

Instances requiring updates will not change
their connection status, for example,
Connected, Warning, Error, Disabled, due to the
pending permission update.

21.8.1.5. Pending cloud instances 

📝🟢 21.8.1.6. Troubleshoot errors on cloud instances 

note

Errors related to collection of audit logs in the
cloud instance are recorded in the
collection_auditing dataset. For more
information, see Audit logs fields and query
examples.

📝🟢 21.8.1.7. Manage Kubernetes Connector instances 

note

After uninstalling the Kubernetes connector,
the connector status updates to Not connected
48 hours after the uninstall process is initiated.

21.8.2. Integrations 

📝🟢

note

In addition to content packs that you install
from Marketplace, related content packs are
automatically downloaded when you adopt
playbooks or edit tasks that require content
items such as scripts or integrations.

📝🟢

note

If the integration was installed from a content
pack, you need to duplicate the integration
before editing.

📝🟢

note

Integration commands are only available when
the integration instance is enabled. Some
commands depend on a successful connection
between Cortex Cloud and third-party
integrations.

💡🔵

tip

In the Playground, you can clear the context
data, if needed, which deletes everything in the
Playground context data, but does not affect
the actual issue or case. To clear the context,
run !DeleteContext all=yes' from the CLI or
click Clear Context Data while viewing the
context data.

🧭➡️

navi…

Settings → Configurations → Marketplace

21.8.2.1. Add an integration instance 

🧭➡️ 21.8.2.2. Configure integration permissions 

navi…

Settings → Configurations → Data Collection
→ Integration Permissions

📝🟢 21.8.2.3. Troubleshoot Integrations 

note

If an integration instance consistently
encounters API rate limits when running on the
tenant, consider configuring the instance to run
on an engine to change the source of the
outbound traffic.

📝🟢 21.8.3. Verify collector connectivity 

note

For more information on troubleshooting data
collector applet errors, see Troubleshoot
Broker VM applet connectivity.

21.8.4. About health issues 

☑️🟣

prer…

For Cortex Cloud to monitor data ingestion
health and create health issues, you must
enable the following settings under
Configurations:
Cortex - Analytics: Go to Configurations →
Cortex - Analytics. For more information, see
Enable the Analytics Engine and Identity
Analytics.Enable the Analytics Engine and
Identity Analytics

📝🟢

note

The Health Issues page displays issues that
were triggered after July 2024. To see health
issues that were triggered before this date,
click Legacy Health Issues.

📝🟢

note

Cortex Cloud enforces the dedup logic to
health issues. This logic reduces the likelihood
of identical health issues from flooding the
issues dataset.

21.8.4.1. Investigate and resolve health issues 

21.8.4.2. Monitor data ingestion health 

22. Marketplace 

🧭➡️ 22.1. Cortex Marketplace 

navi…

Settings → Configurations → Marketplace) or
at Cortex Developer Docs Marketplace

📝🟢 22.2. Content Pack Support Types 

note

Palo Alto Networks is not liable for and does
not warrant or support any content pack
produced by a third-party publisher.
Palo Alto Networks does not support content
packs that do not have official available
documentation.

22.3. Cortex Cloud content 

22.4. Manage content packs 

📝🟢

note

Optional content packs that are already
installed are treated like they are required
content packs to preserve content integrity.

📝🟢

note

Cortex Cloud includes a built-in default mail
sender. You also have the option of installing a
different mail sender content pack, such as
Microsoft Exchange Online.

📝🟢

note

In addition to content packs that you install
from Marketplace, related content packs are
automatically downloaded when you adopt
playbooks or edit tasks that require content
items such as scripts or integrations.

📝🟢

note

Third-party product integrations are developed
and tested against a specific product version.
For products that are on-prem or cloud-based
with specific API versions, the version
developed and tested against will be included
in the integration's documentation. Newer
versions of the product are not always
immediately tested, and it is expected that...

⚠️🟡

caut…

If you want to downgrade, any content that
depends on the content pack including any
customizations may be deleted if it does not
exist in the target content pack version.

⚠️🟡

caut…

If another content pack is dependent on the
content pack you want to delete, it may break
the other content pack. You can reinstall the
content pack, but you cannot restore detached
and customized content.

💡🔵

tip

You can also find content packs that require
updates by going to Settings → Data Sources &
Integrations and filtering by Pack Version =
Update Available. If you click on an integration
in the filtered list, there is a link to the content
pack in Marketplace for updates.

22.5. Marketplace FAQs 

22.6. Content changes when upgrading Cortex
Cloud versions 

23. Cortex CLI 

📝🟢

note

This role is not supported for CWP, as the CWP
system does not support offline mode.

📝🟢

note

There are no preconfigured CLI-specific roles.
Add the CLI Tools permission to an existing role
or create a dedicated custom role.

23.1. Connect Cortex CLI 

☑️🟣

prer…

System requirements:
macOS (Intel Core i7, such as Sequoia): To
ensure all functionalities work correctly, you
must install the vectorscan dependency via
Homebrew, using this command: brew install
vectorscan

RHEL 8.10 and Red Hat UBI9. The following
prerequisites must be met:
Install patchelf
Install zstd
Ubuntu 20 requires the prefetch...

📝🟢

note

The Cortex CLI requires an API key with the
Standard security level.

📝🟢

note

On macOS arm 64 architecture you must
unpack the downloaded file to retrieve the
executable.

🧭➡️

navi…

If needed, retrieve your public API URL by
navigating to Settings → Configurations → API
Keys and clicking Copy API URL, then paste it
into the code

Windows installer → uninstall via Settings →
Apps → Installed apps, or delete the

23.2. Authenticate credentials 

📝🟢 23.3. Cortex CLI usage 

note

For more information about CLI usage for CWP,
refer to Cortex CLI for Cloud Workload
Protection
For more information about CLI usage for API
Security, refer to Cortex CLI for API Security
For more information about CLI usage for
Cortex Cloud Application Security, refer to
Cortex CLI usage for Cortex Cloud Application
Security

23.4. Self-service API keys for CLI scans 

23.5. Cortex CLI common command line

reference guide 

23.6. Cortex CLI for API Security 

☑️🟣

prer…

Ensure you have the required user permissions.
Refer to Cortex CLI for more information
Onboard and install the Cortex CLI. Refer to
Connect Cortex CLI for more information
Ensure your application exposes APIs and
provides a corresponding OpenAPI
Specification file
Ensure that you have installed Java v 11 and
above

23.6.1. Cortex CLI API Security command line
reference guide 

23.7. Cortex CLI for Cloud Workload Protection 

☑️🟣

prer…

Ensure you have the required user permissions.
Refer to Cortex CLI for more information
Onboard and install the Cortex CLI. Refer to
Connect Cortex CLI for more information
Verify that Java version 11 and above is
installed: Run java -version in your terminal. If
not, refer to Java SE Development Kit 11.0.25 for
information about installing Java

☑️🟣

prer…

Before you begin, ensure you have sudo
privileges to execute the image scan.

📝🟢

note

CWP does not support container image secret
scanning for systems running on ARM
architecture.

📝🟢

note

For available CWP commands, refer to Cloud
Workload Protection command line reference.

📝🟢

note

--archive: When used with image scan, sets the
scan source to an archive file. When used with
image sbom, indicates the SBOM should be
exported from an archive file
The --archive flag can also be explicitly set as -
-archive=true
--archive-format <value>: The image archive
format (such as docker-archive or oci-archive).
Default: docker-archive.

23.7.1. Cloud Workload Protection command
line reference



23.8. Cortex CLI for Code Security 

☑️🟣

prer…

For the Cortex CLI binary:
Ensure you have Node.js v22 installed on your
host machine before running any scans with the
Cortex CLI. This is crucial to prevent runtime
errors, as the CLI depends on Node.js for
executing JavaScript analysis
Note

To check your version of Node.js, run node -v
To download Node.js, refer to the official
Node.js...

23.8.1. Cortex CLI usage for Cortex Cloud
Application Security 

23.8.2. Cortex CLI Cortex Cloud Application
Security command line reference



📝🟢

note

The repo-id flag must not end with .config, .log
or .ini. -config is acceptable.
-- repo-id foo.config will be blocked
-- repo-id foo-config will pass

👀🟠

imp…

The Cortex CLI Cortex Cloud Application
Security only supports single occurrences of
each flag. If the same flag is passed multiple
times, only the last provided value will be used.
For example, in the following command, only TF
CloudFormation will be the scanned framework.
./cortexcli --api-base-url <YOUR_API_URL> --
api-key <YOUR_API_KEY> --auth-id...

23.8.3. Cortex CLI pre-commit hooks 

☑️🟣

prer…

These common prerequisites are required for
all types of installation (both local and global)
of the Cortex CLI pre-commit hook.
Ensure you have a license for Cortex Cloud
Application Security
Install the Cortex Cloud CLI binary locally.
Refer to Connect Cortex CLI for information
about onboarding the CLIConnect Cortex CLI
Obtain Cortex Cloud API...

☑️🟣

prer…

For local installation: Install the pre-commit
framework version 3.2.0 or greater. Refer to
https://pre-commit.com/ for installation
instructions.

📝🟢

note

It is recommended you configure credentials
for the Cortex CLI using a configuration file.

📝🟢

note

The pre-commit framework is not required for
global mode.

23.8.4. Pre-commit hook usage 

23.8.5. Cortex CLI pre-receive hooks 

☑️🟣

prer…

Before you begin, ensure you have:
Administrator access to the VCS server and
console
A valid license for Cortex Cloud Application
Security
The Cortex Cloud CLI binary or Docker image
installed on the server (requires GLIBC (GNU C
library) version 2.35 or greater). Refer to
Connect Cortex CLI for information about
onboarding the CLIConnect Cortex...

📝🟢

note

Make sure to create the directory under the
home directory of the Linux user that runs the
Git hooks. This user is typically not the root user.

23.8.6. Pre-receive hook usage 

24. Cortex Cloud XQL 

24.1. Get started with XQL 

24.1.1. XQL language features 

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion commands and
definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc key.

🧭➡️

navi…

You submit XQL queries to Cortex Cloud using
the Investigation & Response → Search →
Query Builder user interface

24.1.2. XQL Language Structure 

📝🟢

note

Specifying a dataset is not required because
Cortex Cloud uses xdr_data as the default
dataset. If you have more than one dataset or
lookup, you can change your default dataset
by navigating to Settings → Configurations →
Data Management → Dataset Management,
right-click on the appropriate dataset, and
select Set as default.

📝🟢

note

You can also build a query that investigates
data in both a cold dataset and hot dataset in
the same query. In addition, as the hot storage
dataset format is the default option and
represents the fully searchable storage, this
format is used throughout this guide. For more
information on hot and cold storage, see
Dataset management.

24.1.2.1. Adding comments in queries 

📝🟢 24.1.3. Supported operators 

note

In some cases, using an IN or NOT IN operator
combined with a dataset and filterfilter stage
can be a better alternative to using a join
stage.join

24.1.4. Datasets and presets 

📝🟢

note

To set up this Cloud Identity Engine (previously
called Directory Sync Service (DSS)) dataset,
you need to set up a Cloud Identity Engine.
Otherwise, you will not have a pan_dss_raw
dataset. For more information, see Set up Cloud
Identity Engine.Set up Cloud Identity Engine

📝🟢

note

INFO issues are not included in this dataset.
The issue fields included in this dataset are
limited to certain fields available in the API. For
the full list, see Get Alerts Multi-Events v2 API.

📝🟢

note

The fields contained in this dataset are a subset
of the fields in the xdr_data dataset.

📝🟢

note

If the vendor and product are not specified in
the Winlogbeat profile's configuration file,
Cortex Cloud creates a default dataset called
microsoft_windows_raw.

📝🟢

note

To ensure GlobalProtect access authentication
logs are sent to Cortex Cloud, verify that your
PANW firewall's Log Settings for GlobalProtect
has the Cortex Data Lake checkbox selected.

📝🟢

note

The fields contained in this dataset are a subset
of the fields in the xdr_data dataset.

📝🟢

note

Prisma Access firewalls do not send
configuration logs to the Structured Log
Storage (SLS).

📝🟢

note

You can query in XQL for this data and build
widgets based on the xdr_data dataset or using
the preset device_control.
To view in an XQL query these events, the
Device Configuration of the endpoint profile
must be set to Block. Otherwise, the USB events
are not captured. The events are also captured
when a group of device types are blocked on
the...

📝🟢

note

The fields contained in this dataset are a subset
of the fields in the xdr_data dataset.

🧭➡️ 24.1.5. About examples 

navi…

Investigation & Response → Search → Query
Builder → XQL → Query Library

👀🟠 24.1.6. JSON functions 

imp…

JSON field names are case sensitive, so the key
to field pairing must be identical in an XQL
query for results to be found. For example, if a
field value is "TIMESTAMP" and your query is
defined to look for "timestamp", no results will
be found.

24.1.7. How to filter for empty values in the
results table 

24.1.8. Understanding string manipulation in XQL 

24.2. Build XQL queries 

📝🟢

note

If you have the Cortex Agentic Assistant, you
can use natural language prompts to create
and run XQL queries within the chat interface.
For more information, see Create and run XQL
queries with Agentic Assistant chat.Cortex
Agentic Assistant

24.2.1. About the Query Builder 

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query, and can appear with a slight delay.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion commands and
definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc key.

24.2.2. How to build XQL queries 

📝🟢

note

Users with different dataset permissions can
receive different results for the same XQL query.
An administrator or a user with a predefined
user role can create and view queries built with
an unknown dataset that currently does not
exist in Cortex Cloud. All other users can only
create and view queries built with an existing
dataset.
When you have more...

📝🟢

note

You can build a query that investigates data in
both a cold dataset and a hot dataset in the
same query. In addition, as the hot storage
dataset format is the default option and
represents the fully searchable storage, this
format is used throughout this guide for
investigation and threat hunting. For more
information on hot and cold storage, see...

👀🟠

imp…

Forensic datasets are not inlcuded by default in
XQL query results, unless the dataset query is
explicitly defined to use a forensic dataset.

24.2.2.1. Get started with XQL queries 

24.2.2.2. Useful XQL user interface features 

📝🟢

note

Schema changes to datasets may not be
reflected in the autocomplete suggestions and
definitions as you type in real time the XQL
query and can appear with a slight delay.

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

24.2.2.3. XQL Query best practices 

24.2.2.4. Expected results when querying fields 

24.2.2.5. Create XQL query 

📝🟢

note

Whenever the time period is changed in the
query window, the config timeframe is
automatically set to the time period defined for
the entire query, including queries that are part
of the join stage. Yet, this won't be visible as
part of the query. Only if you manually type in
the config timeframe will this be seen in the
query.

These time picker...

💡🔵

tip

When creating XQL queries, you can:
Use the up and down arrow keys to navigate
through the auto-suggestion command
suggestions and definitions.
Select an auto-suggestion command by
pressing either the Enter or Tab key.
Press Shift+Enter to add a new line, and easily
ignore the auto-suggestion output.
Close the auto-suggestion output by pressing
the Esc...

💡🔵

tip

While the query is running, you can navigate
away from the page. A notification is sent when
the query has finished. You can also Cancel the
query or run a new query, where you have the
option to Run only new query (cancel previous)
or Run both queries.

24.2.2.6. Review XQL query results 

📝🟢

note

It's also possible to graph the results displayed.
For more information, see Graph query results.

📝🟢

note

Real time query results are available only in the
Query Builder and in free text query.
Real time results are displayed only for queries
run on hot datasets.
The Sort option is available only after all the
data is retrieved.
When you formulate complex queries, the
results will be displayed when the query has
finished running completely, and not in real...

📝🟢

note

Results are received incrementally for the first
100K records, or up to 100MB worth of records,
whichever comes first. After that, the next
update is when the query has finished running
completely.

📝🟢

note

In order for Cortex Cloud to provide a
histogram for a field, the field must not contain
an array or a JSON object.

👀🟠 24.2.2.7. Translate to XQL 

imp…

This feature is still in a Beta state and you will
find that not all Splunk queries can be
converted to XQL. This feature will be improved
upon in the upcoming releases to support
greater Splunk query translations to XQL.

24.2.2.8. Graph query results 

🔒🟣

cont…

Building Cortex Query Language (XQL) queries
in the Query Builder requires a Data Collection
add-on.

📝🟢

note

To display the result of as a time duration,
choose the graph type Single Value and enable
Show as Time. You can then select the Time
Unit (millisecond, second, minute, or hour) and
the Display format.

💡🔵

tip

Alternatively, you can use the Cortex Agentic
Assistant to generate custom graphs and
charts using natural language prompts. By
simply prompting the agent, it will build and
execute the query, returning the visual
representation. For more information, see Use
natural language to query and visualize your
data.

24.2.3. XQL query entities 

🧭➡️ 24.2.3.1. Create authentication query 

navi…

Investigation & Response → Search → Query
Builder

24.2.3.2. Create event log query 

24.2.3.3. Create file query 

24.2.3.4. Create image load query 

24.2.3.5. Create network connections query 

📝🟢 24.2.3.6. Create network query 

note

When you run the query, depending on the
outcome of the results, the value specified in
this field might be displayed in the dst_ip field
in the query results. This occurs if an RDP event
is recorded whereby a user connected from the
source IP to the destination IP.

24.2.3.7. Create process query 

👀🟠 24.2.3.8. Create registry query 

imp…

Ensure the KEY NAME is entered as a real
registry key name, and not as a symbolic link.
Otherwise, the query will not retrieve results.
Instead of
HKEY_LOCAL_MACHINE\System\CurrentControl
Set, which is a symbolic link, use
KEY_LOCAL_MACHINE\System\ControlSet001.
Instead of HKEY_CURRENT_USER, use
HKEY_USERS\<SID>, where SID is either a SID of
the current...

24.2.3.9. Query across all entities 

24.2.4. Overview of the Query Center 

📝🟢

note

Very short queries might not be listed.
You cannot cancel correlation queries.
The default retention period for historic queries
is aligned with issue retention.

24.2.4.1. Edit and run queries in Query Center 

📝🟢

note

If query limits are applied to your tenant, the
number of concurrent running queries is limited
per user. If query usage is reaching the defined
limit, a system message warns you that a high
query load is impacting performance. If you
exceed the limit, new queries are blocked until
query usage drops. You can view all active
queries under Query Center...

📝🟢

note

You can cancel your own queries. To cancel
queries run by other users, you must have
View/Edit permissions for Configurations →
Query Management. By default, Instance
administrators have View/Edit permission.

📝🟢

note

Cancelled queries show a Canceled status. You
can see details of all canceled queries in the
Query History tab.
You cannot cancel correlation rule queries.
If you cancel a scheduled query, only the
current query is cancelled. Future recurrences
of the scheduled query are not affected.

📝🟢 24.2.4.1.1. Query Center reference information 

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

24.2.5. Manage scheduled queries 

🧭➡️

navi…

Investigation & Response → Search →
Scheduled Queries

24.2.5.1. Scheduled Queries reference

information 📝🟢 
note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

24.2.6. Manage your query library 

📝🟢

note

The Query to Library option is only available if
your role has the Create Queries capability. For
more information, see Manage access to saved
queries.

📝🟢

note

When the tenant-level setting Owners and
editors can change the general access is
unselected, the drop-down is disabled and only
an administrator can configure this option.

👀🟠

imp…

The ability to create, edit, or share queries is
governed by access management. If certain
options are unavailable, contact your
administrator. For more information, see
Manage access to saved queries.

24.3. Cortex XQL syntax, parameters, and
examples 

25. Graph Search 

☑️🟣 25.1. What is Graph Search? 

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

25.2. Get started with Graph Search queries 

🔒🟣

cont…

This feature is included with a Cortex XSIAM
Premium license. It is also included with any
other Cortex XSIAM license that has the Cloud
Posture Security or Cloud Runtime Security
add-on.

☑️🟣

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

🧭➡️

navi…

Understand your assets and findings data:
Graph Search queries are based on the current
data that has been collected for assets and
findings from the data sources configured and
then sent to the Unified Asset Inventory (UAI),
which is displayed in the All Assets page
(Inventory → Assets → All Assets

25.3. How to build Graph Search queries? 

☑️🟣

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

🧭➡️

navi…

You submit Graph Search queries using the
Investigation & Response → Search → Query
Builder → Graph Search built-in query interface

☑️🟣 25.4. Understand Graph Search query results 

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

25.5. Create Graph Search query 

☑️🟣

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

📝🟢

note

For more information about the Query Library,
see Manage the Graph Search Query Library.

💡🔵

tip

After running the query, you can view the
complete query by hovering over the last
THAT... in the Graph Search query interface, and
the query is displayed in a tooltip.
If your query doesn't find any results or you
want to change your query for any reason, you
can always click anywhere in the Graph Search
query interface, where your existing query is...

25.6. Graph Search examples 

🔒🟣

cont…

This feature is included with a Cortex XSIAM
Premium license. It is also included with any
other Cortex XSIAM license that has the Cloud
Posture Security or Cloud Runtime Security
add-on.

☑️🟣

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

25.7. Manage the Graph Search Query Library 

☑️🟣

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

🧭➡️

navi…

From Graph Search in the Query Builder Select
Investigation & Response → Search → Query
Builder → Graph Search

Investigation & Response → Search → Query
Builder → Graph Search

Investigation & Response → Search → Query
Center

25.8. Edit and run queries in Query Center 

📝🟢

note

If query limits are applied to your tenant, the
number of concurrent running queries is limited
per user. If query usage is reaching the defined
limit, a system message warns you that a high
query load is impacting performance. If you
exceed the limit, new queries are blocked until
query usage drops. You can view all active
queries under Query Center...

📝🟢

note

You can cancel your own queries. To cancel
queries run by other users, you must have
View/Edit permissions for Configurations →
Query Management. By default, Instance
administrators have View/Edit permission.

📝🟢

note

Cancelled queries show a Canceled status. You
can see details of all canceled queries in the
Query History tab.
You cannot cancel correlation rule queries.
If you cancel a scheduled query, only the
current query is cancelled. Future recurrences
of the scheduled query are not affected.

🧭➡️

navi…

Investigation & Response → Search → Query
Center → Query History

📝🟢 25.8.1. Query Center reference information 

note

Certain fields are exposed and hidden by
default. An asterisk (*) is beside every field that
is exposed by default.

☑️🟣 25.9. Supported assets and findings 

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

☑️🟣 25.10. FAQ on Graph Search 

prer…

Graph Search requires View or View/Edit RBAC
permissions for Graph Search under
Investigation & Response → Search.

26. API specification inventory 

📝🟢

note

Even if you have already imported the
specification, you can edit the API specification
in Cortex Cloud and add or update the server
list.

📝🟢 26.1. Import API specification 

note

Even if you already imported the file, you can
edit the API asset and add or update the server
list.